

# GDG DevFest TAIPEI 2016



## Activity Recognition with TensorFlow

TensorFlow, Deep Learning, and Activity Recognition

Presented by: Aaron Lai

# | Who Am I



Data Scientist

Building System

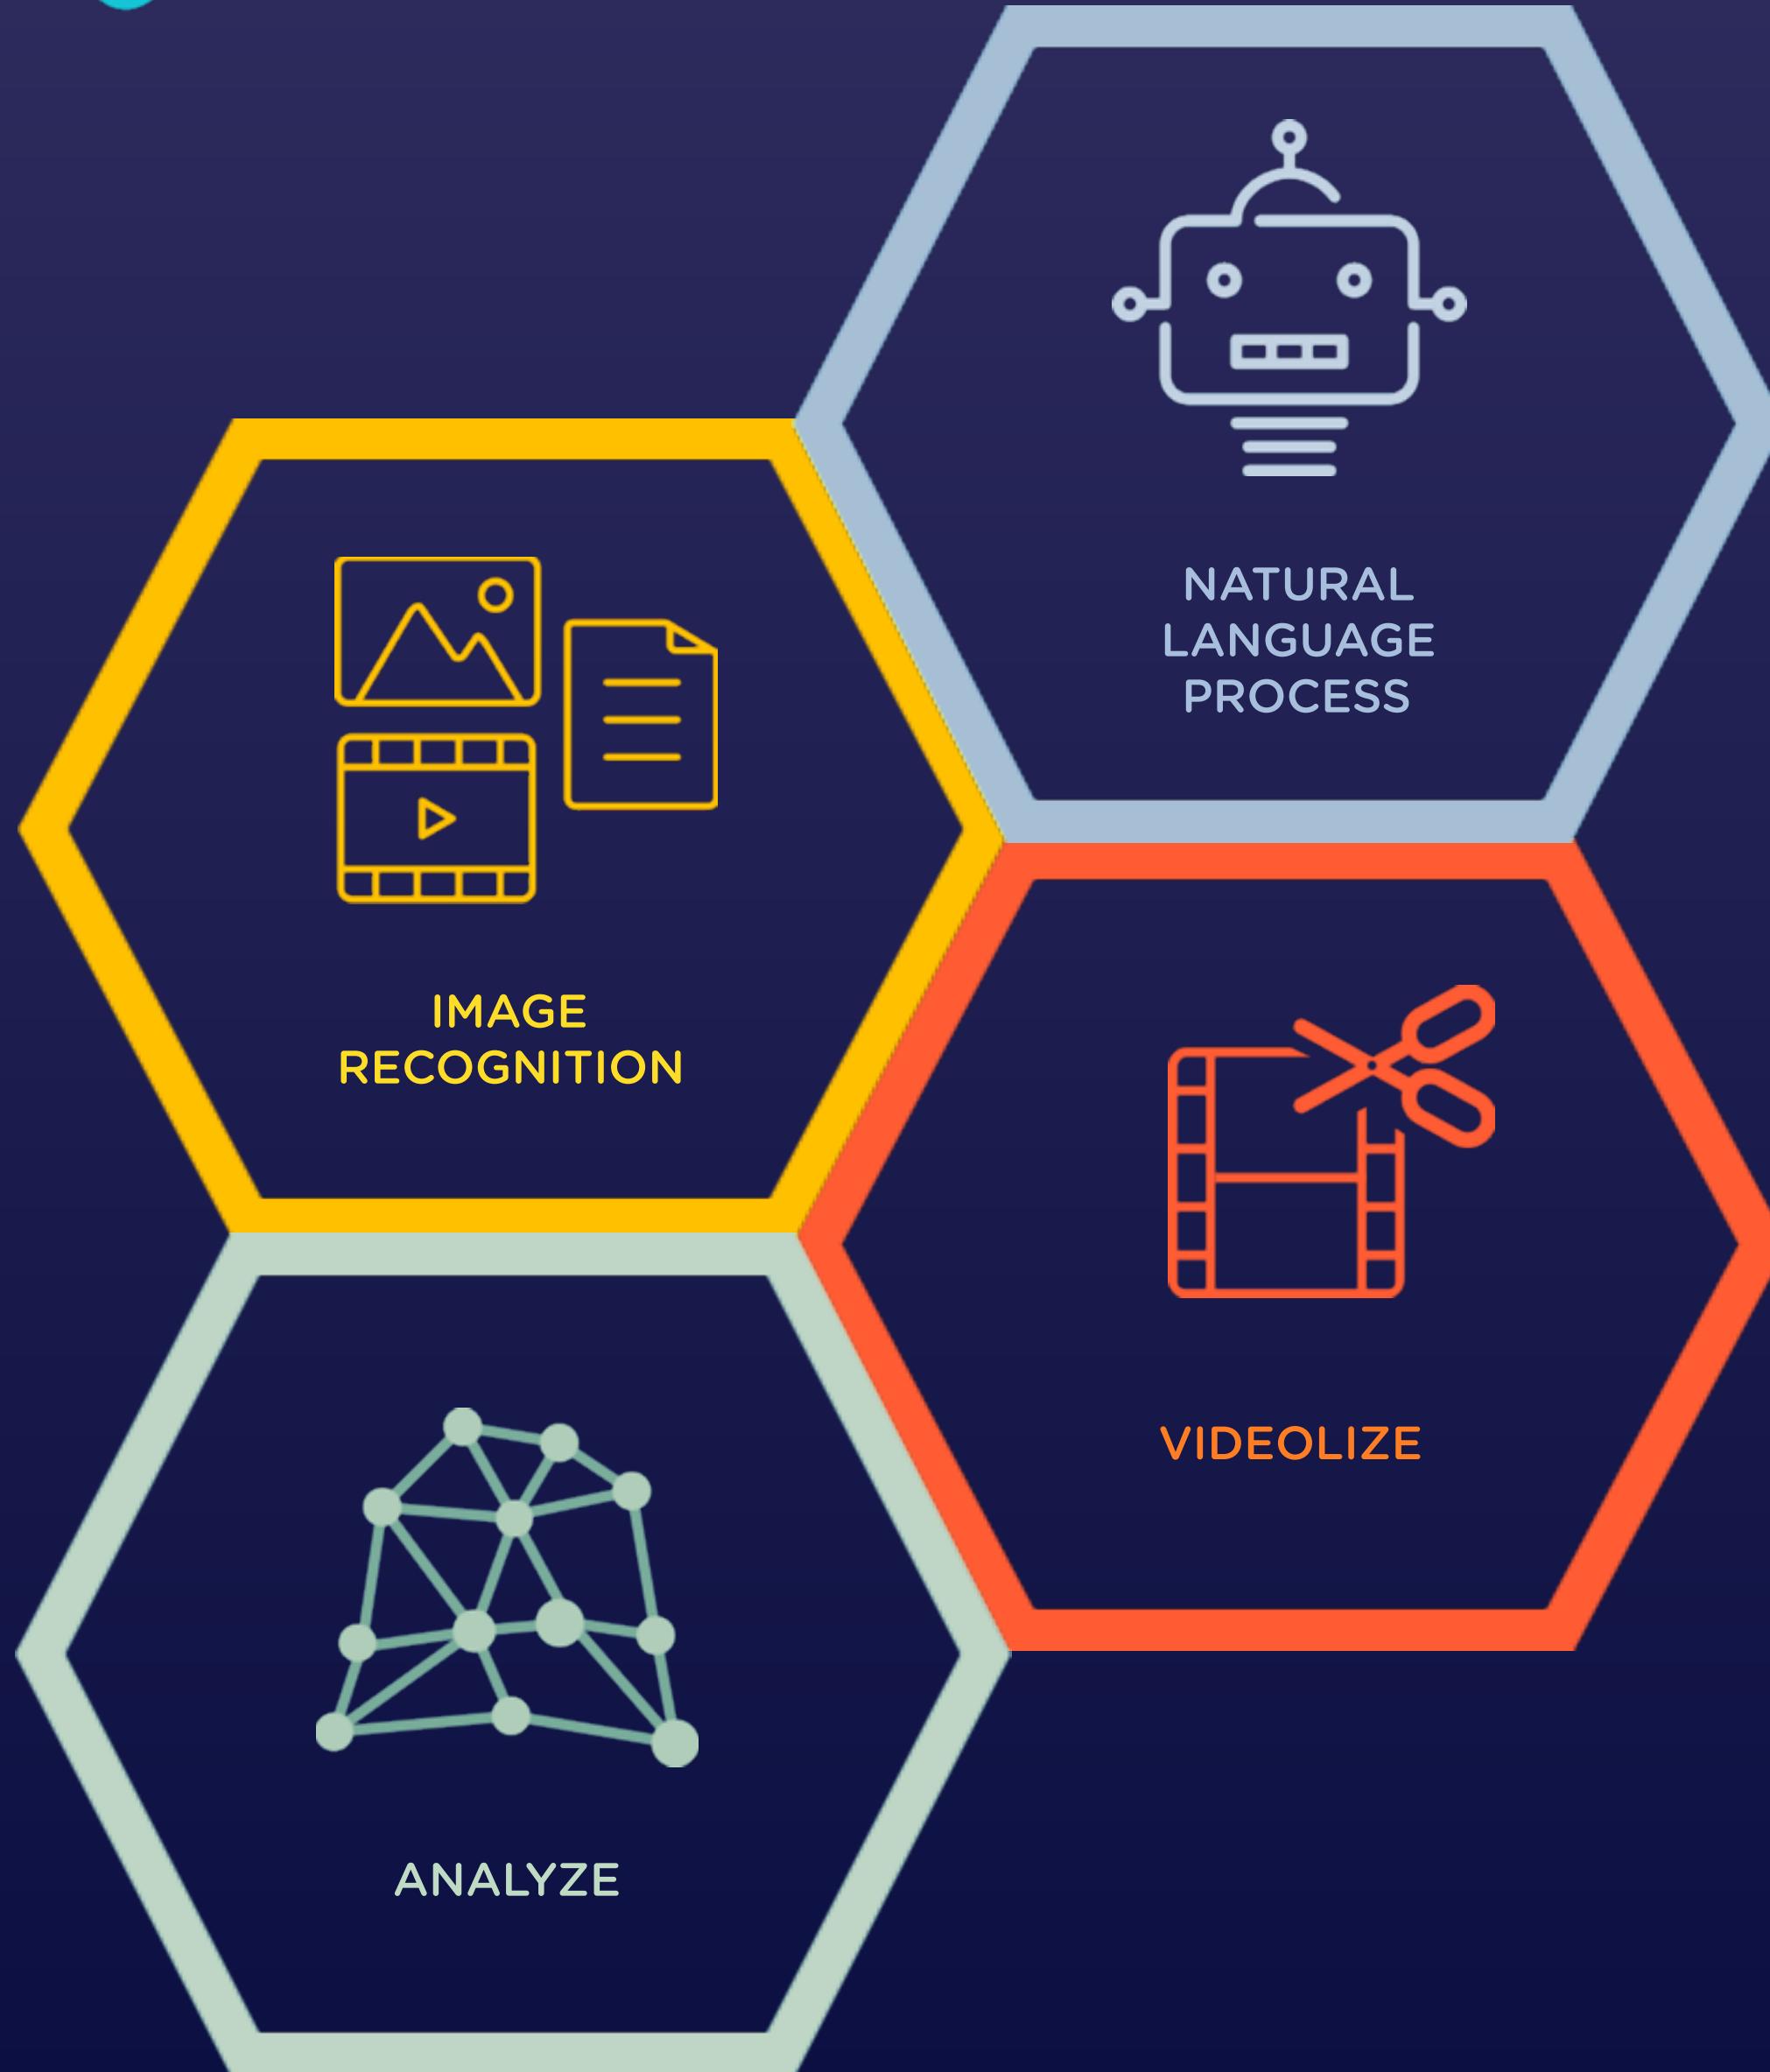
Playing New Tools

Data Analysis

Machine Learning



# GLIA STUDIO - AI Video Creation Platform



Create Video Analysis

All your videos

Text

The highly anticipated opening of the 2016-17 NBA season is just around the corner and it is already clear who are the payroll losers and winners. **NBA MVP Stephen Curry** tops the list of players with underwhelming salary, while **Finals MVP LeBron James** leads the biggest moneymakers.

Underpaid Players

According to Hoops Hype, **Curry** will be underpaid next season with his **salary of \$12,112,359**. The 28-year-old point guard has one more season left with the **Golden State Warriors** before turning free agent next summer, where **he is expected to get the highest salary raise**.

In the past two seasons, **Curry** achieved **two NBA MVP awards** and won **one championship title**, but he has been the least paid All-Star player in the team. This could be attributed to the **four-year \$44-million rookie scale contract extension** he signed with the Warriors in 2012. He battled several health woes during the early years of his career, making the franchise

Next season, per

Abstract

- NBA MVP Stephen Curry tops the list of players with underwhelming salary

- Finals MVP LeBron James leads the biggest moneymakers

- Curry will be underpaid next season with his salary of \$12,112,359

- Curry achieved two NBA MVP awards and won one championship title, but he has been the least paid All-Star player in the team

A

B

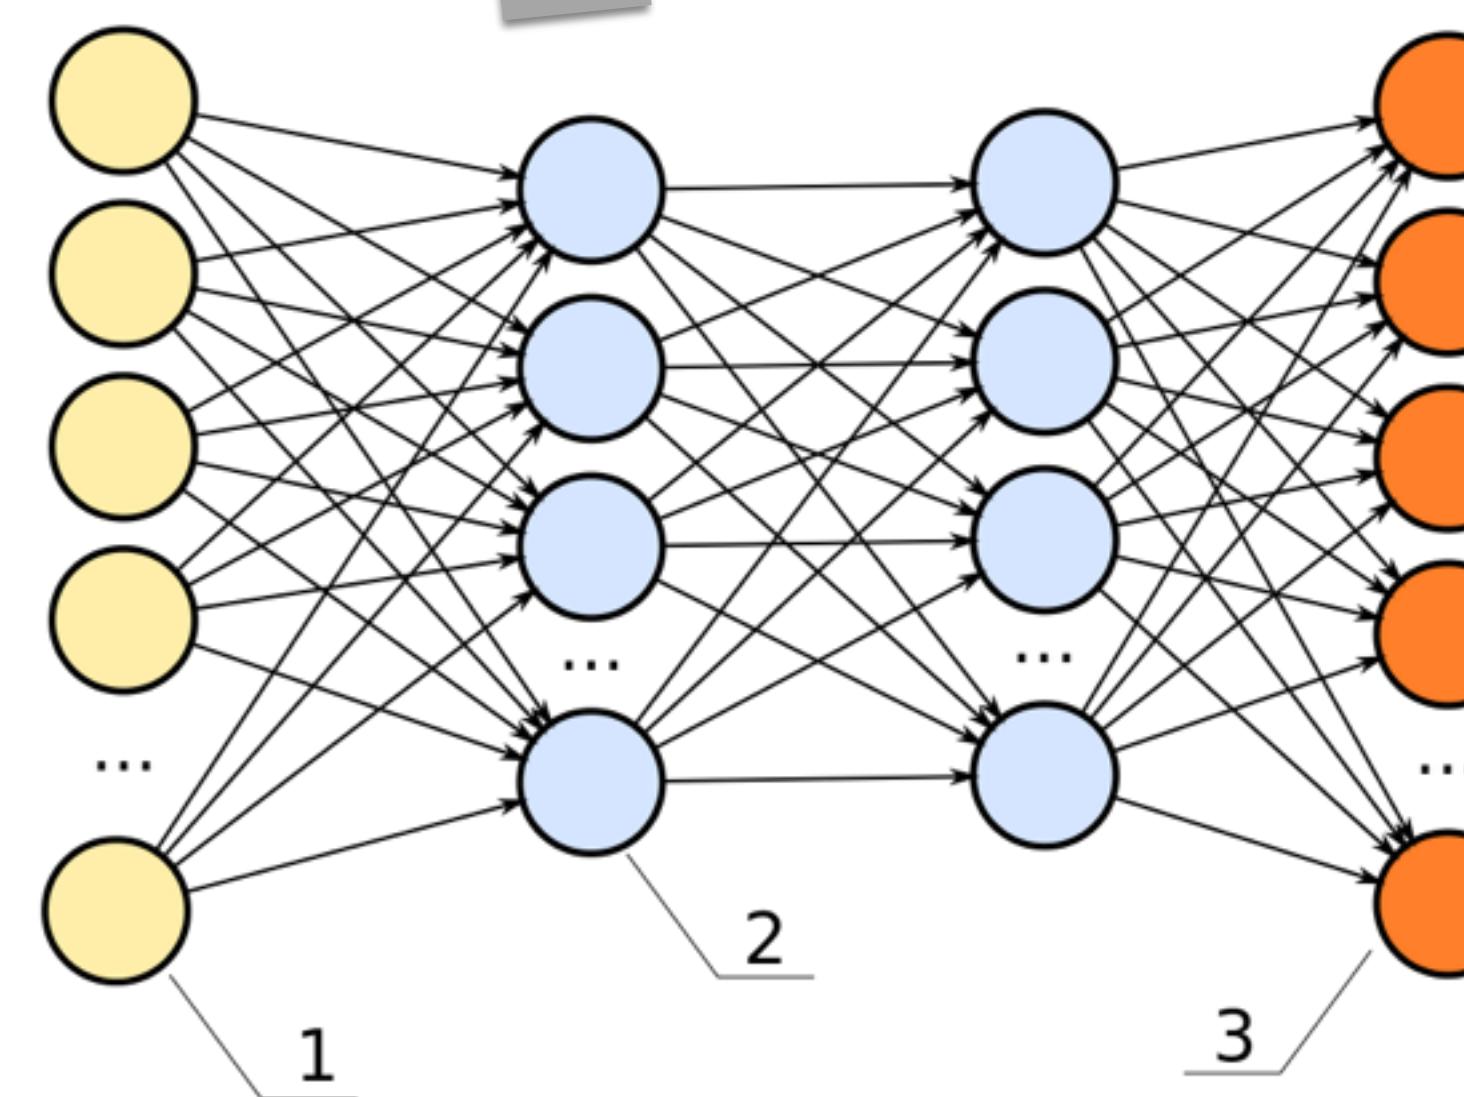
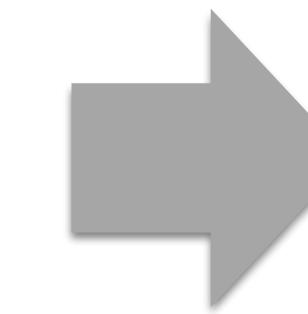
Engagement: 980 people

Engagement: 319 people

# Outline



**ACTIVITYNET**

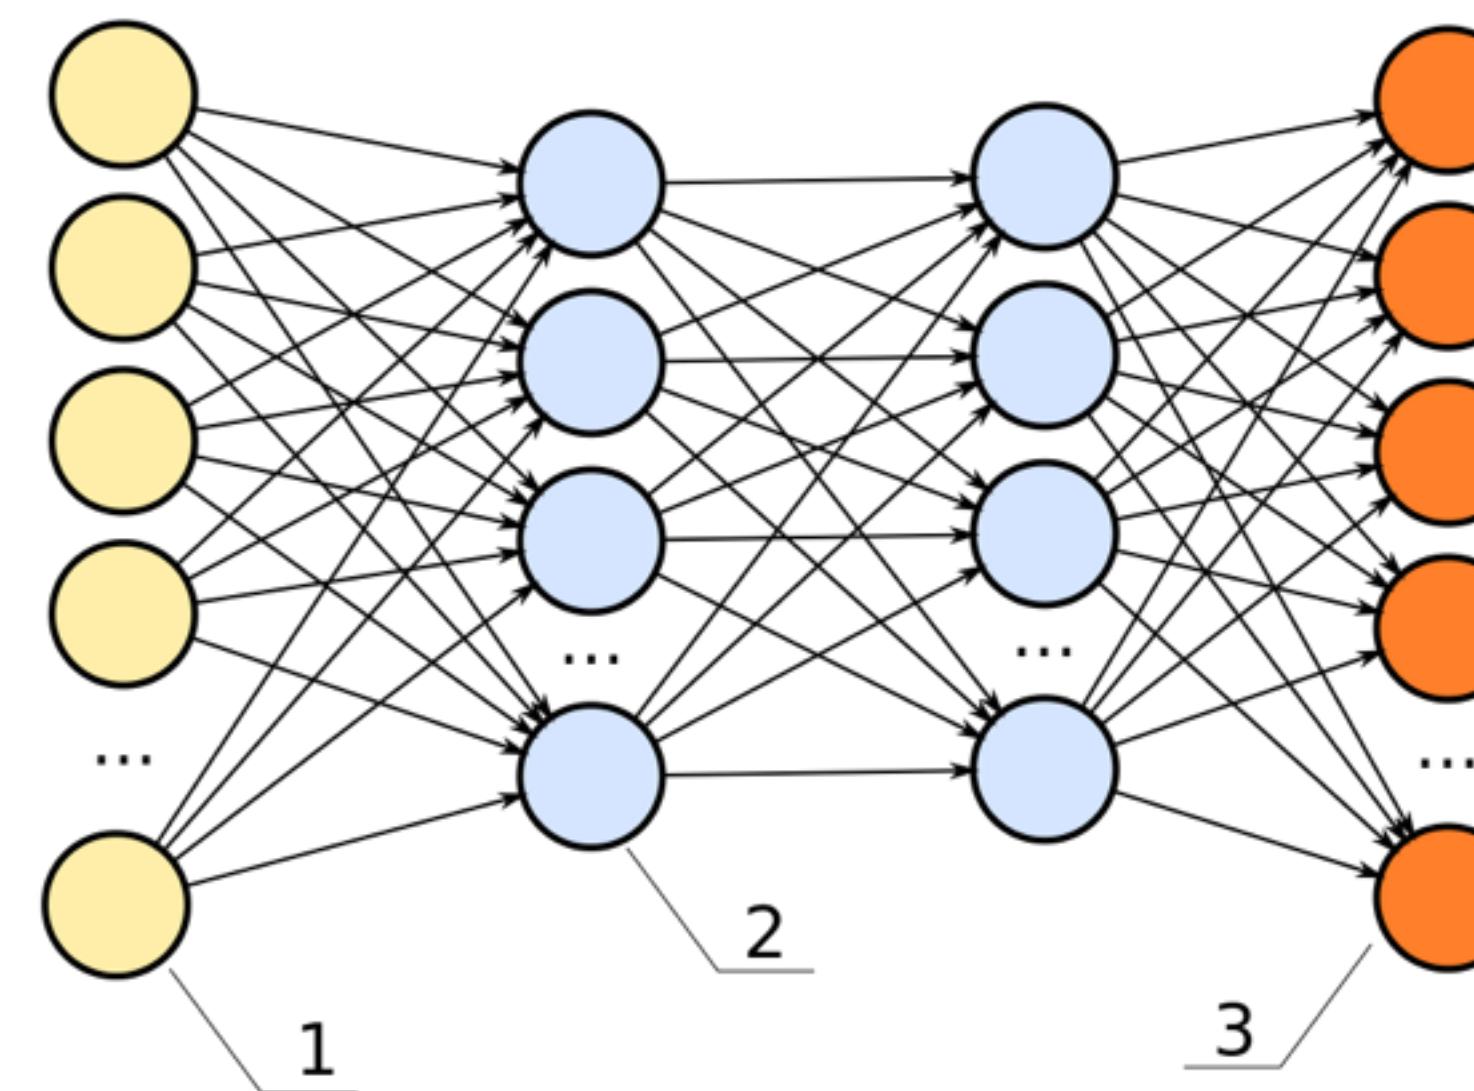


**TensorFlow**

# Outline



ACTIVITYNET



TensorFlow

# Activity Recognition



## Dataset

- more than 648 hours of untrimmed videos
- ~20K videos
- 200 activity categories
- Youtube

Reference : <http://activity-net.org/challenges/2016/index.html>

## Untrimmed Classification Challenge

Given a long video, predict the labels of the activities present in the video

## Detection Challenge

Given a long video, predict the labels and temporal extents of the activities present in the video

# Activity Recognition



Reference : <http://activity-net.org/challenges/2016/index.html>



GDG DevFest  
TAIPEI 2016

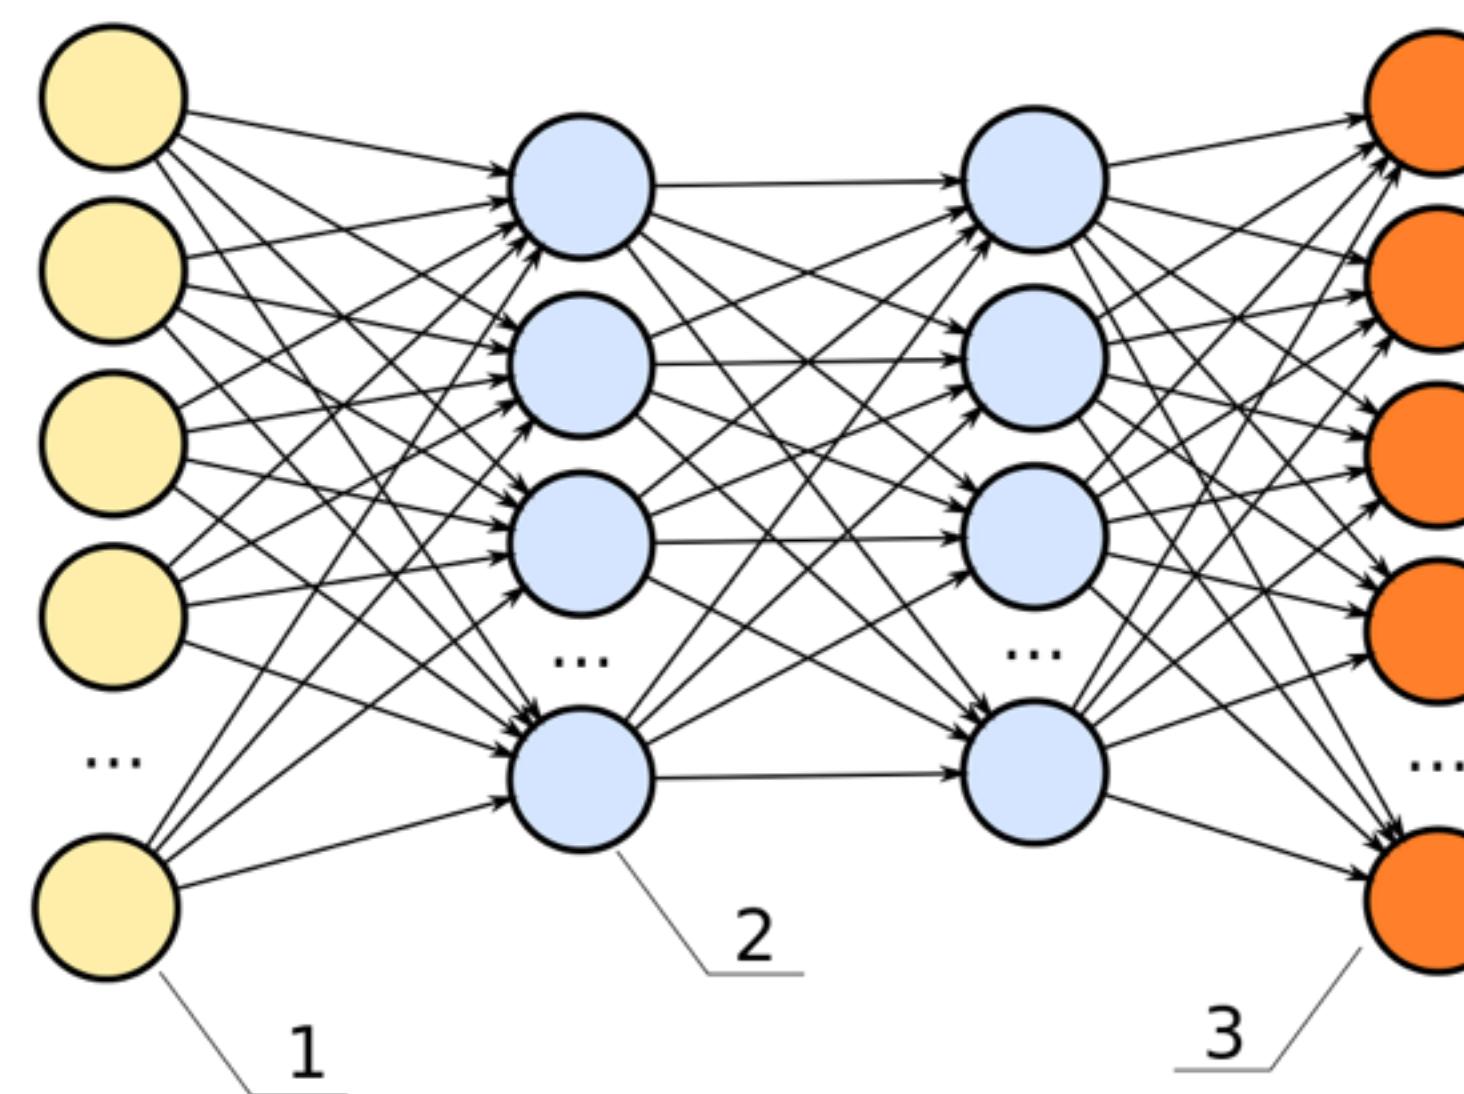
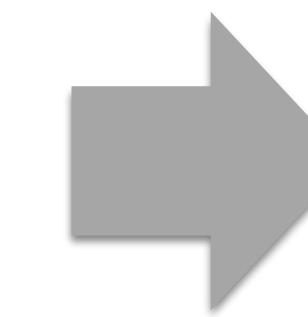


GLIA CLOUD

# Outline



**ACTIVITYNET**



GDG DevFest  
TAIPEI 2016

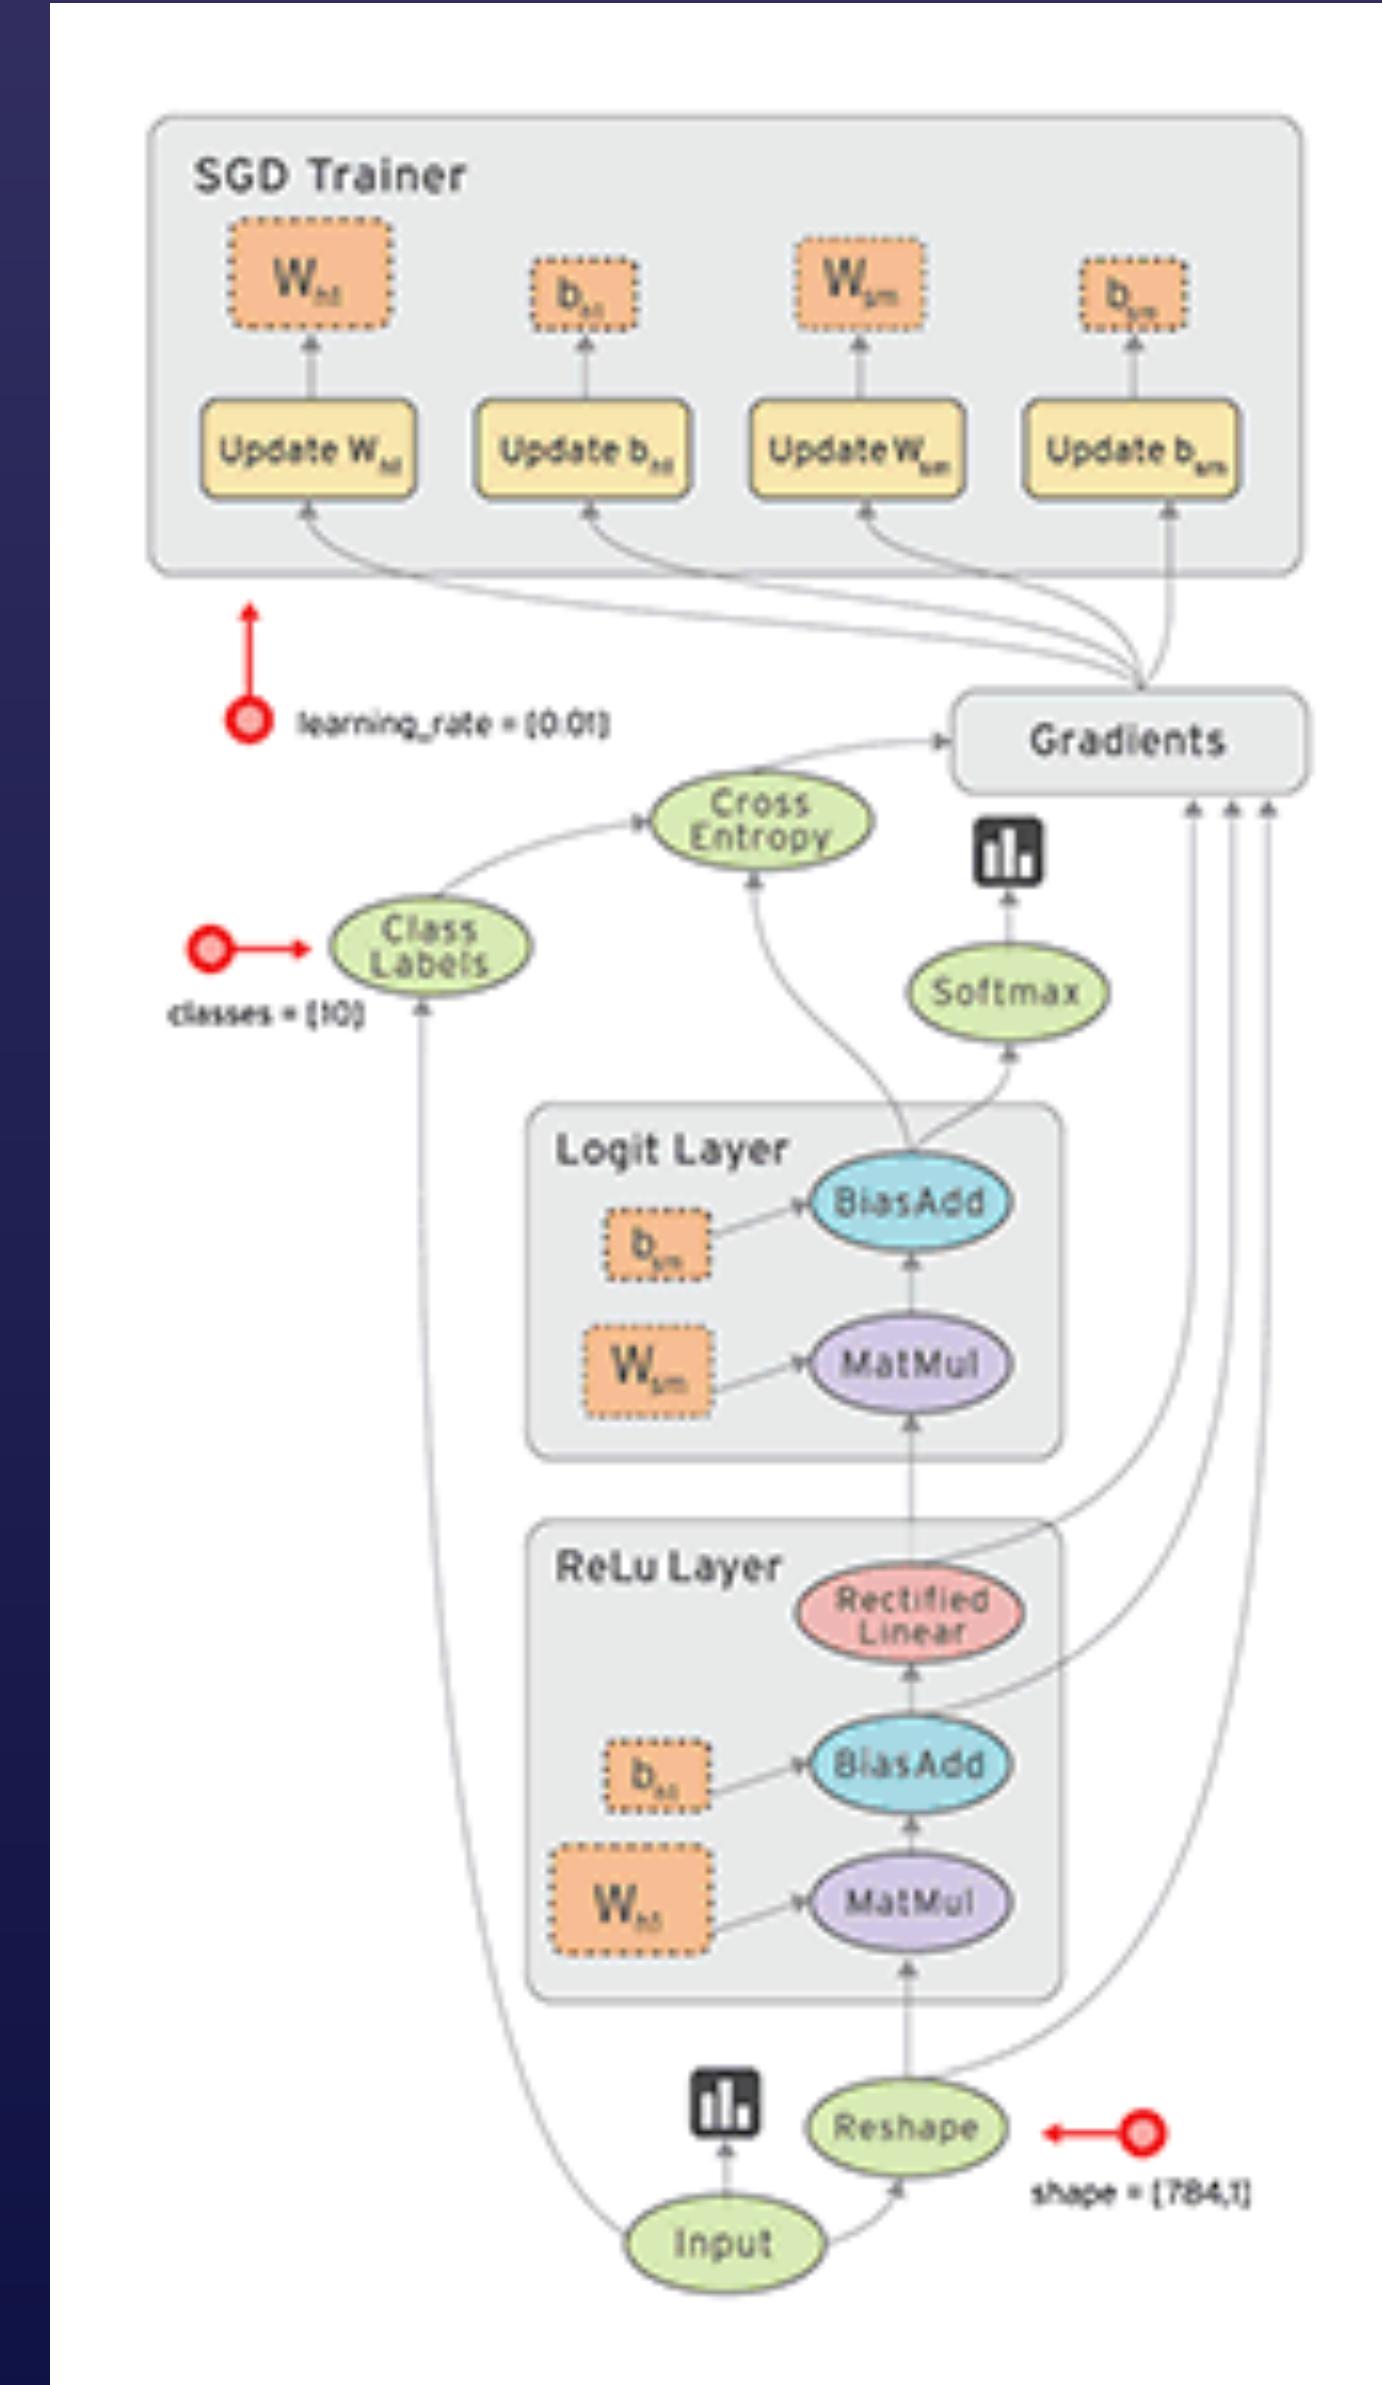


GLIA CLOUD

# TensorFlow

## Library for Machine Intelligence

- Open source at 2015.11
- Deep Learning
- Numerical computation : “data flow graphs” (node: ops, edge: tensors)
- One or more CPUs or GPUs
- **Auto-Differentiation!!** (Numerical Methods)



Reference : <https://www.tensorflow.org/>



GDG DevFest  
TAIPEI 2016



GLIA CLOUD

# Inputs, Variables and Operations



## Inputs - Placeholders

- Nodes for data inputs
- Static or dynamic shape

```
# tf Graph Input
x = tf.placeholder("float", shape=[None, 784])
y = tf.placeholder("float", shape=[None, 10])
```

## Variables

- Pass initial values
- Need to be initialized (later)

```
# Set model weights
w = tf.Variable(np.random.randn(784,10))
b = tf.Variable(np.random.randn(10))
```

## Basic Ops

- Matrix multiply, concatenate
- add (broadcast), multiply

```
Wx = tf.matmul(x, w)
pred = tf.add(Wx, b)
```

```
# higher order
w0 = tf.random_normal([784*2, 10])
ws = tf.Variable(w0)
x2 = tf.mul(x, x)
xs = tf.concat(1, [x, x2])
```

```
WXs = tf.matmul(xs, ws)
pred = tf.add(WXs, b)
```

# How to Run

## Session

- Connect the backend
- Run the graph
- `tf.initialize_all_variables()`
- `tf.Session()`
- `tf.run()`
- Feed data : `feed_dict={}`

```
init = tf.initialize_all_variables()
sess = tf.Session()
# initialize variables
with sess.as_default():
    sess.run(init)

with sess.as_default():
    print(sess.run(b))
    sess.run(init)
    print(b.eval())

[ 1.6032598 -1.3250742  0.77060711
 -0.12344709  2.22597957 -0.52275437
 [-0.09863706 -2.65096951 -1.14960706
  1.18818402 -1.63062012 -0.935911
```

# What is Scope - Shared Variables



## Variable Scope

- Variables within the same scope could be “reused”
- `tf.get_variable(name, shape)`
- `tf.variable_scope()`
- Manually set “reuse”:  
`scope.reuse_variables()`

```
import tensorflow as tf

with tf.variable_scope("foo") as foo_scope:
    v = tf.get_variable("v", [1])
with tf.variable_scope(foo_scope):
    w = tf.get_variable("w", [1])

# reuse variables : get them back!
with tf.variable_scope(foo_scope, reuse=True):
    v1 = tf.get_variable("v", [1])
    w1 = tf.get_variable("w", [1])

assert v1 == v
assert w1 == w
```

# Example I : Linear Regression



## Update Parameters

- tf.train.optimizer()
- minimize(cost)
- Auto-gradient &

Auto-update!!

```
train_X = np.asarray([3.3,4.4,5.5,6.71,6.93,4.168,9.779,6.182,7.59,2.167,
                     7.042,10.791,5.313,7.997,5.654,9.27,3.1])
train_Y = np.asarray([1.7,2.76,2.09,3.19,1.694,1.573,3.366,2.596,2.53,1.221,
                     2.827,3.465,1.65,2.904,2.42,2.94,1.3])
n_samples = train_X.shape[0]

X = tf.placeholder("float")
Y = tf.placeholder("float")

W = tf.Variable(rng.randn(), name="weight")
b = tf.Variable(rng.randn(), name="bias")

pred = tf.add(tf.mul(X, W), b)
cost = tf.reduce_sum(tf.pow(pred-Y, 2)) / (2*n_samples)
# Gradient descent
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

# Example I : Linear Regression



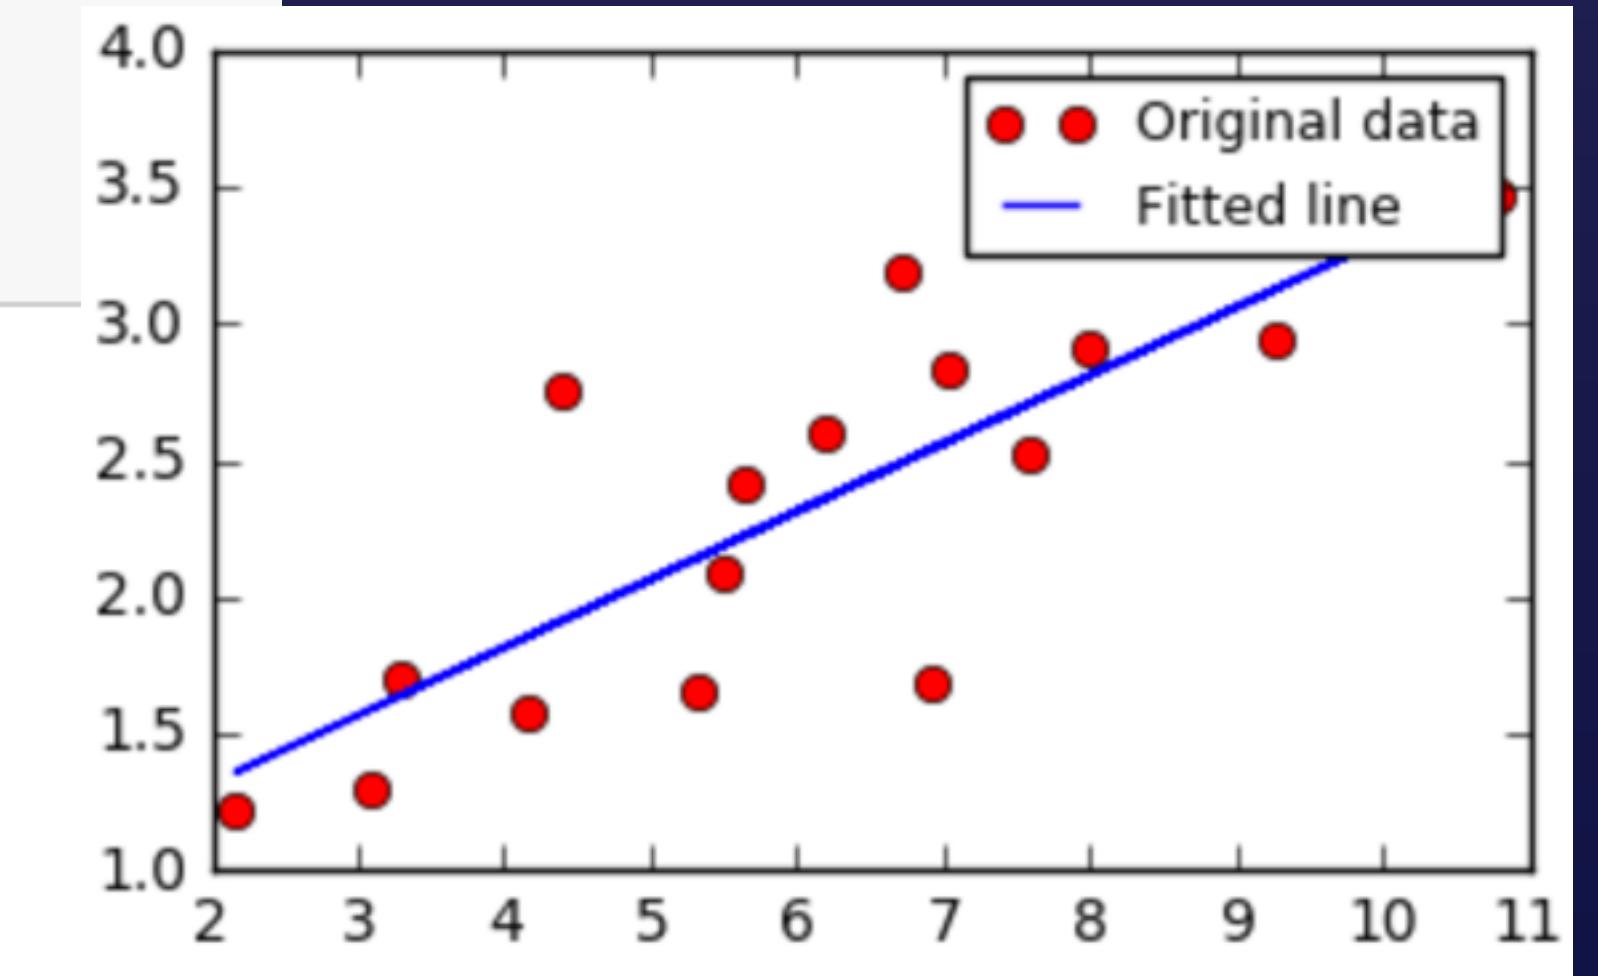
```
init = tf.initialize_all_variables()
with tf.Session() as sess:
    sess.run(init)

    for epoch in range(training_epochs):
        sess.run(optimizer, feed_dict={X: train_X, Y: train_Y})
        log(epoch, display_step, sess, cost, train_X, train_Y, w, b)

    training_cost = sess.run(cost, feed_dict={X: train_X, Y: train_Y})
    print("Training cost=", training_cost, "W=",
          sess.run(w), "b=", sess.run(b), '\n')

    plot_graph(train_X, train_Y, sess, w, b)

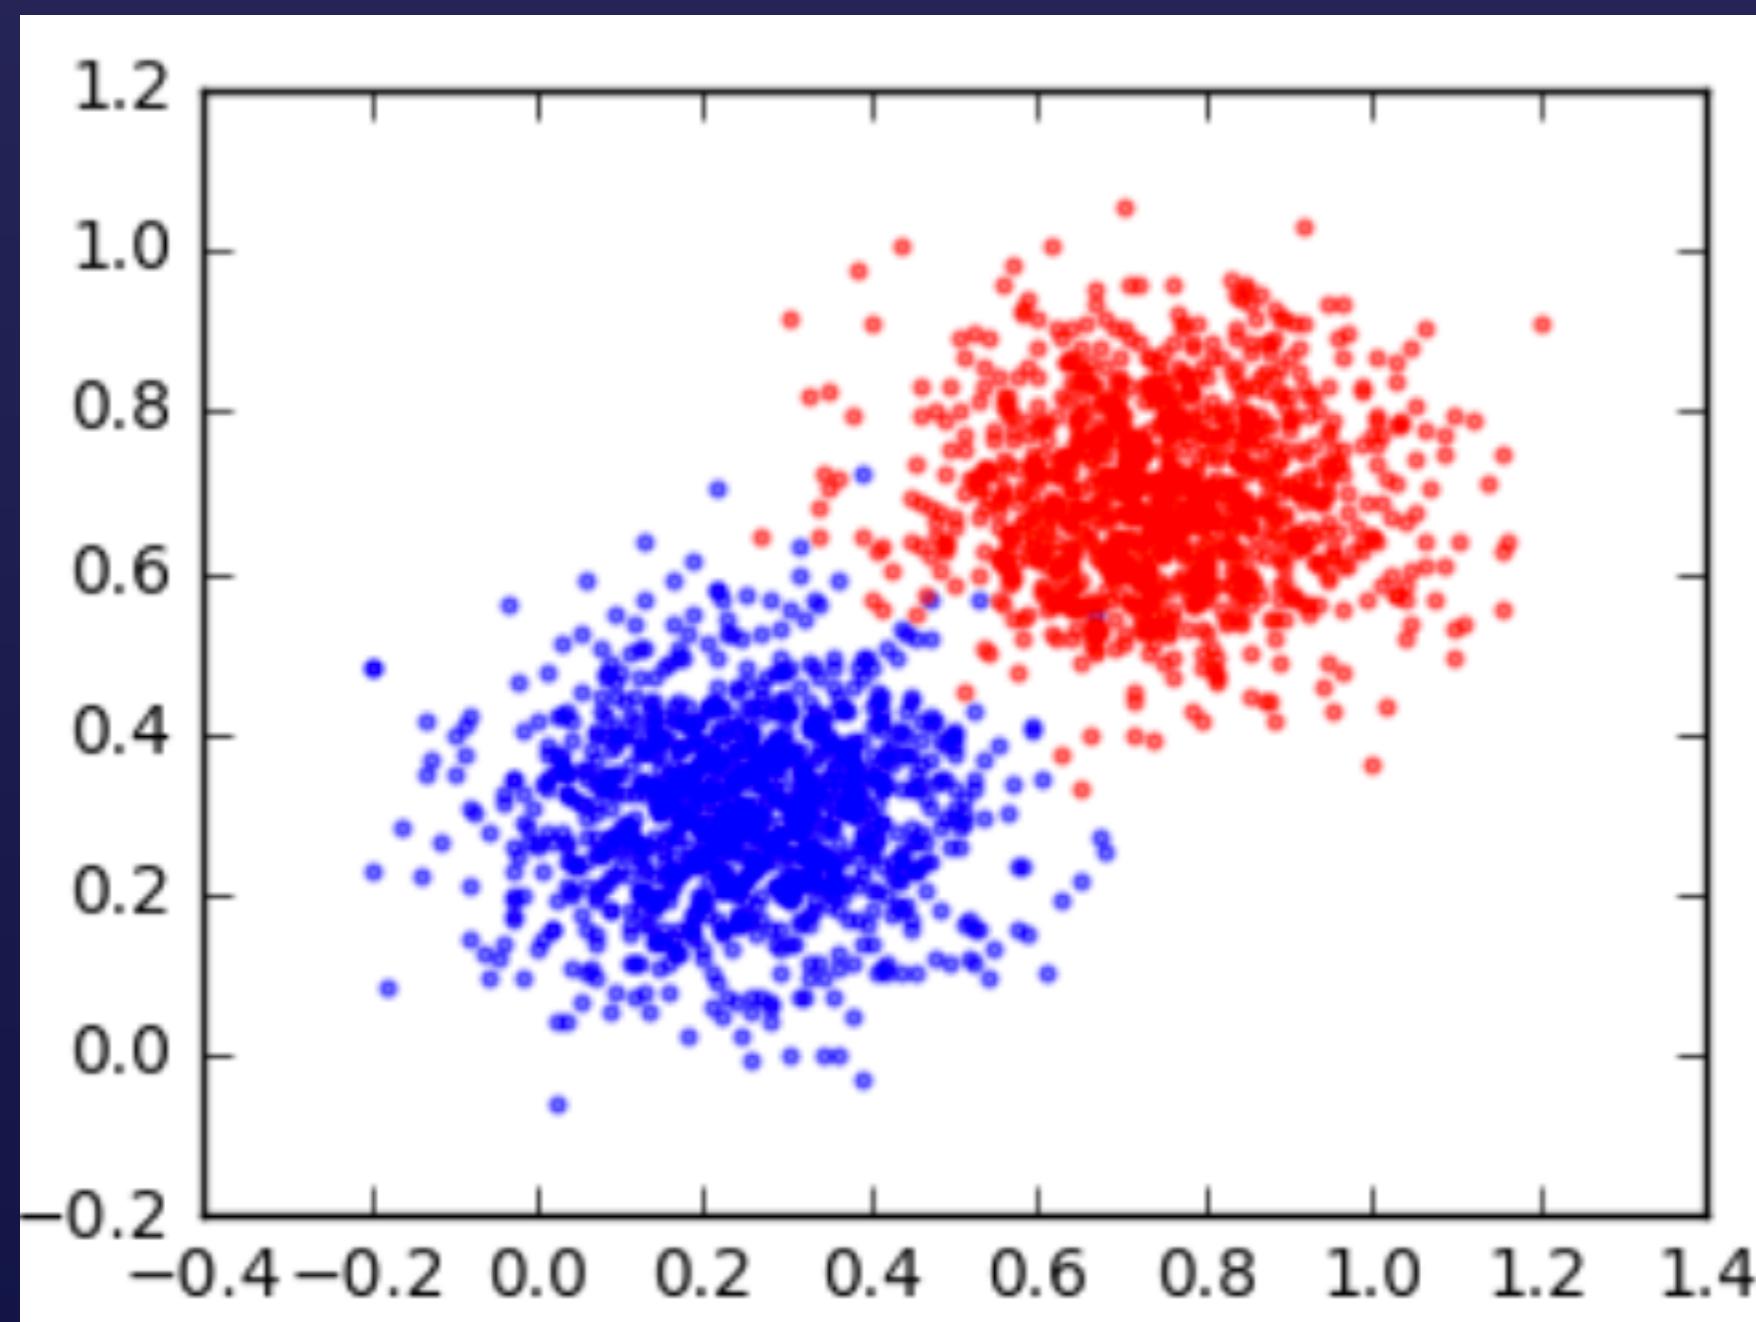
Epoch: 2000 cost= 0.088393934 W= 0.190929 b= 1.22918
Epoch: 4000 cost= 0.079599813 W= 0.222334 b= 1.00653
Epoch: 6000 cost= 0.077551104 W= 0.237493 b= 0.899064
Epoch: 8000 cost= 0.077073790 W= 0.244809 b= 0.847195
Epoch: 10000 cost= 0.076962613 W= 0.24834 b= 0.82216
Optimization Finished!
Training cost= 0.0769626 W= 0.24834 b= 0.82216
```



# Example II : Binary Classification



## Data



## Logistic Regression

- Loss function :

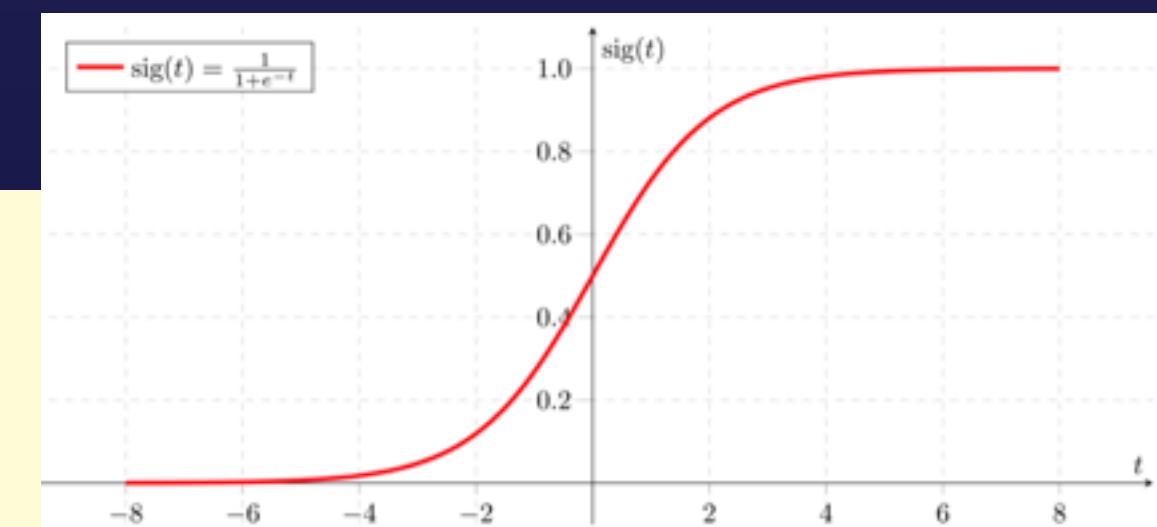
$$\min_{\mathbf{w}} E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \ln \left( 1 + \exp(-y_n \mathbf{w}^T \mathbf{x}_n) \right)$$

- Decision :

logistic regression: use

$$h(\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

to approximate target function  $f(\mathbf{x}) = P(+1|\mathbf{x})$



# Example II : Binary Classification



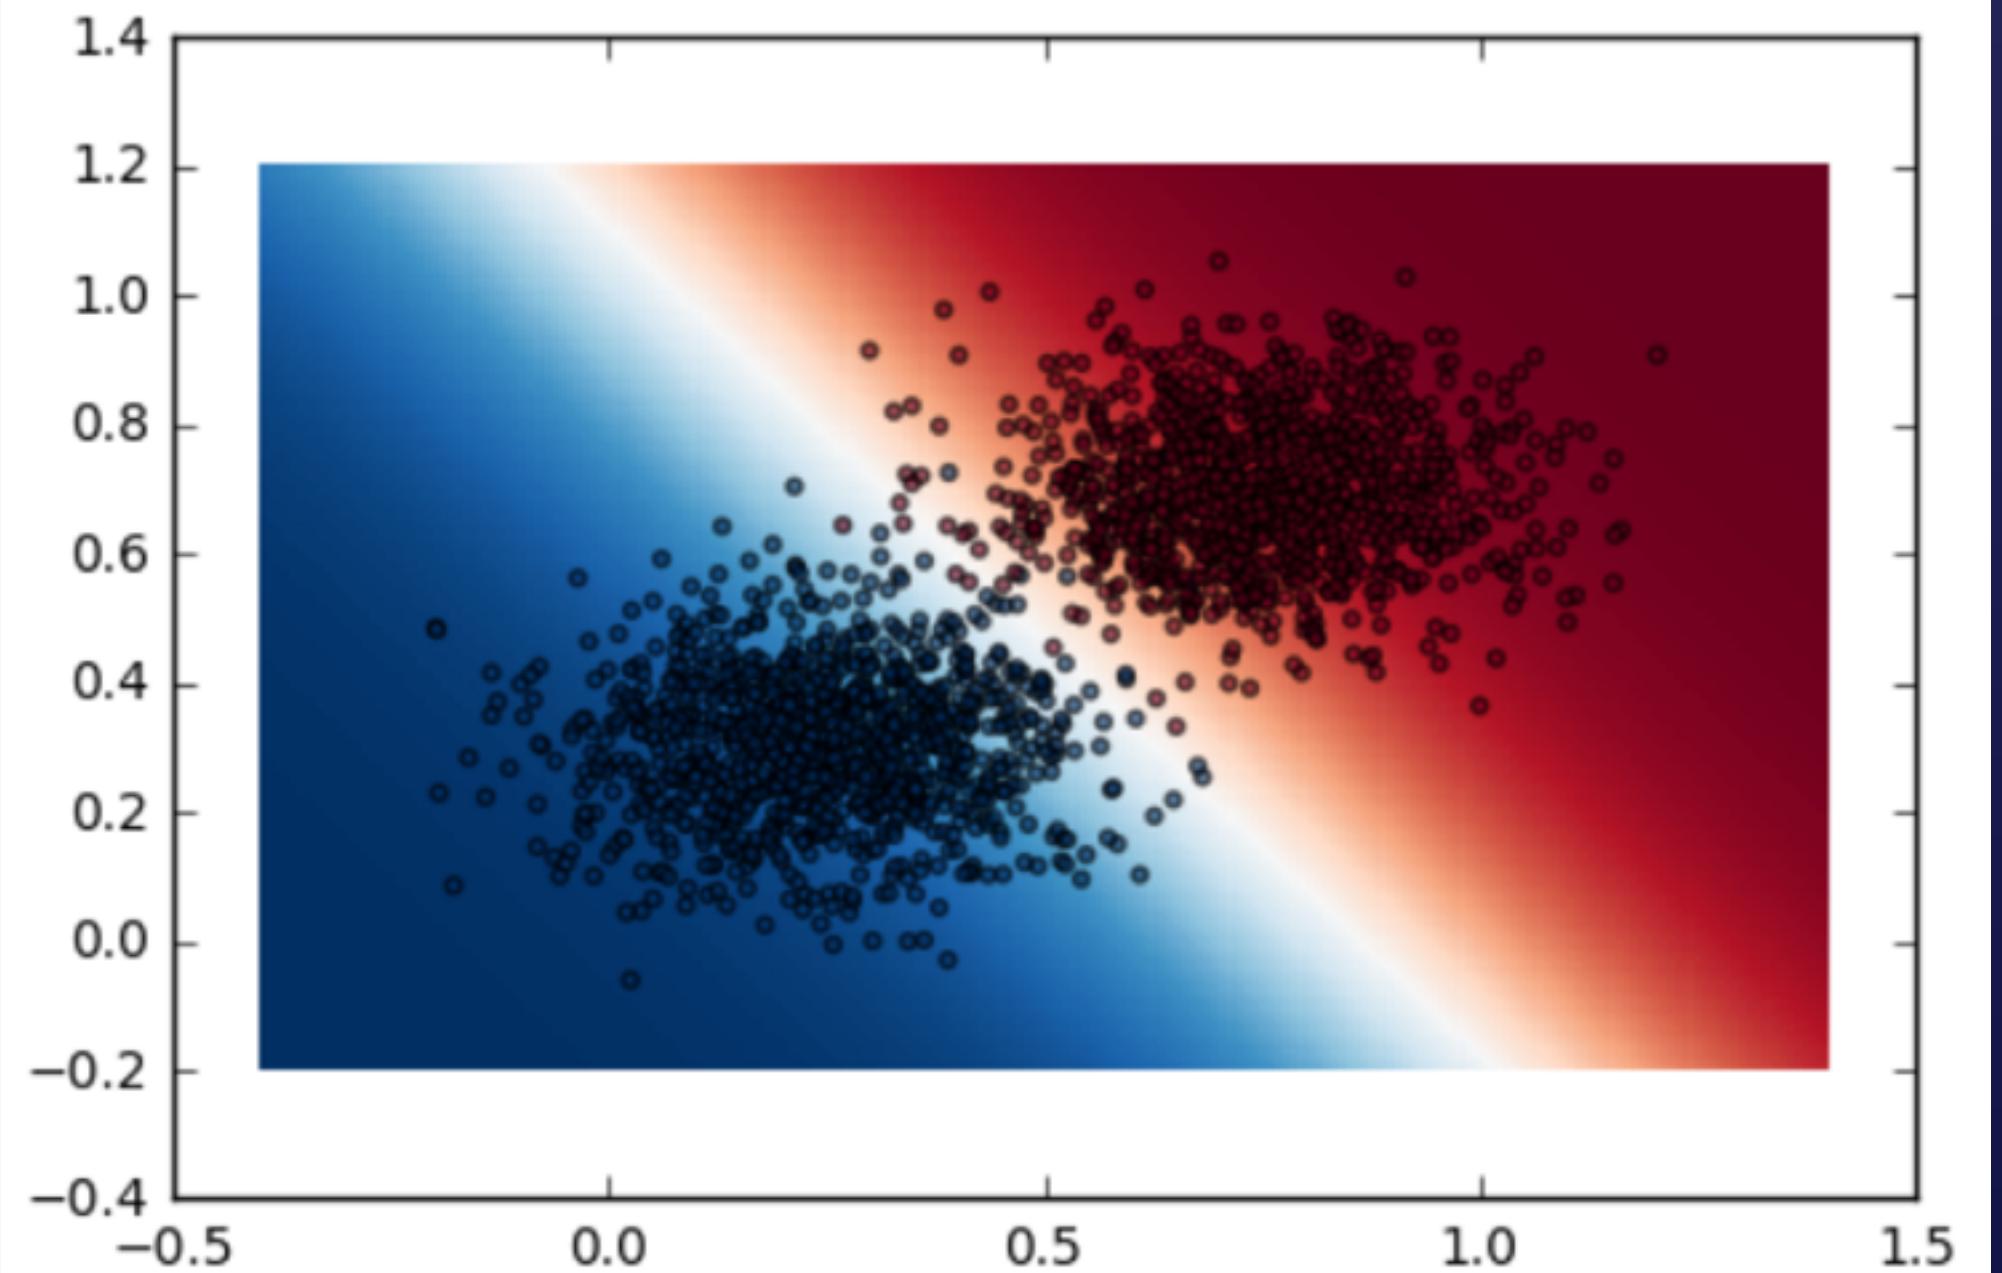
```
# tf Graph Input - dynamic shape
x = tf.placeholder(tf.float32, [None, 2])
y = tf.placeholder(tf.float32, [None, 1])

# Set model weights
W = tf.Variable(tf.zeros([2, 1]))
b = tf.Variable(tf.zeros([1]))

# Construct model
exp_a1 = tf.exp(tf.mul((tf.matmul(x, W) + b), -y)) + 1
pred = tf.nn.sigmoid(tf.add(tf.matmul(x, W), b))

cost = tf.reduce_mean(tf.reduce_sum(tf.log(exp_a1),
                                  reduction_indices=1))

# Gradient Descent
optimizer = tf.train.GradientDescentOptimizer(learning_rate)
optimizer = optimizer.minimize(cost)
```



Reference : <https://github.com/aymericdamien/TensorFlow-Examples>

GDG DevFest  
TAIPEI 2016

GLIA CLOUD

# Using Multicores



## Assign device

- tf.device()
- Matrix multiplication
- (500, 1000) \* (1000, 500)

```
Using cpu: 0:00:00.239567
```

```
Using gpu: 0:00:00.030540
```

```
# CPU
with tf.device('/cpu:0'):
    a = tf.constant(A)
    b = tf.constant(B)
    c = tf.matmul(a, b)

start = datetime.now()
with tf.Session() as sess:
    sess.run(c)
end = datetime.now()
print('Using cpu: %s' % str(end - start))
```

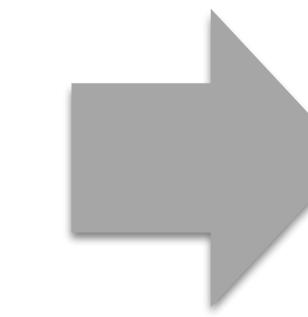
```
# GPU
with tf.device('/gpu:0'):
    a = tf.constant(A)
    b = tf.constant(B)
    c = tf.matmul(a, b)

start = datetime.now()
with tf.Session() as sess:
    sess.run(c)
end = datetime.now()
print('Using gpu: %s' % str(end - start))
```

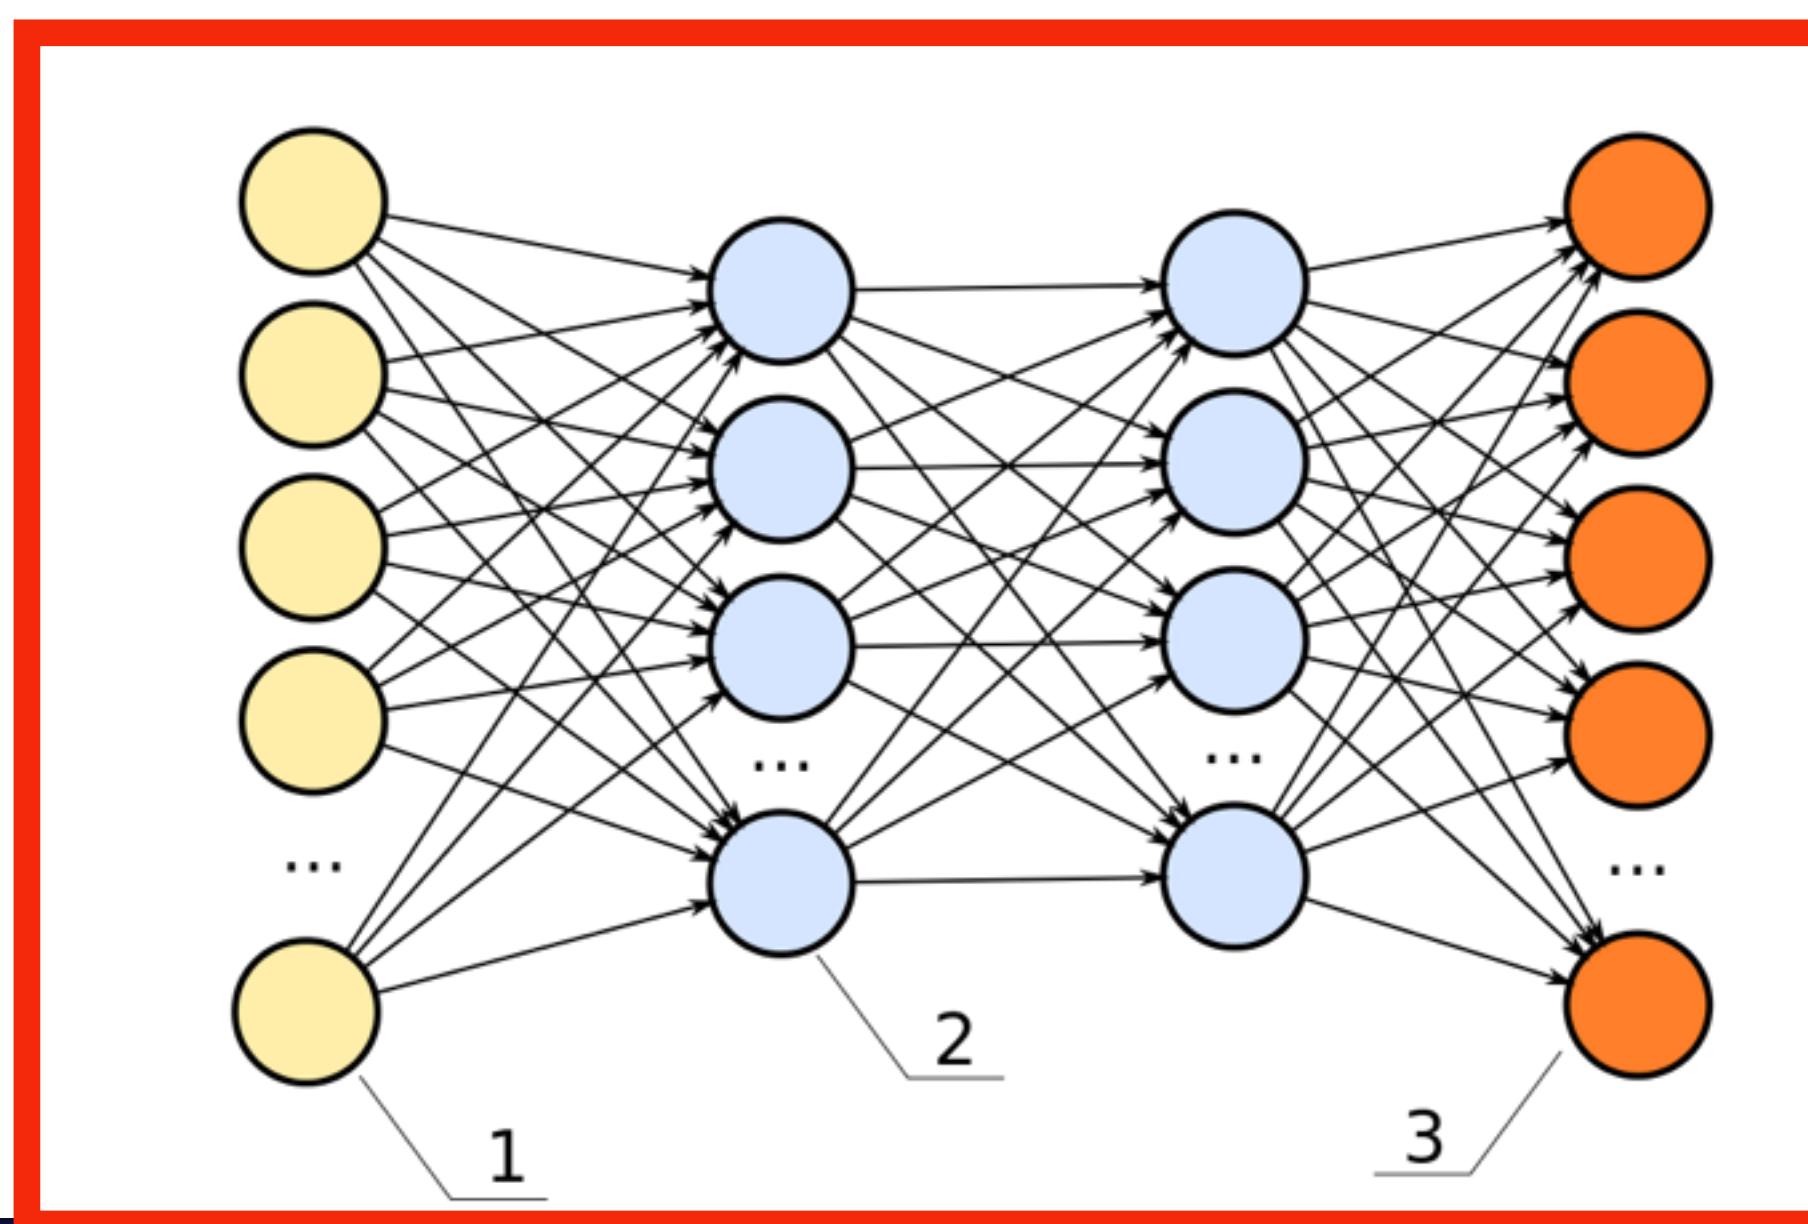
# Outline



ACTIVITYNET



TensorFlow



# Machine Learning Framework

## Hypothesis

$$F(x) \rightarrow y$$

Function Approximation

Define Hypothesis  
Function Set

## Inference

$$x \rightarrow F(x)$$

Given input, predict y

## Testing

## Training

Data  $\rightarrow$  “Best” F

Define and Minimize  
“Loss”

“Best” F:  
over “new” data

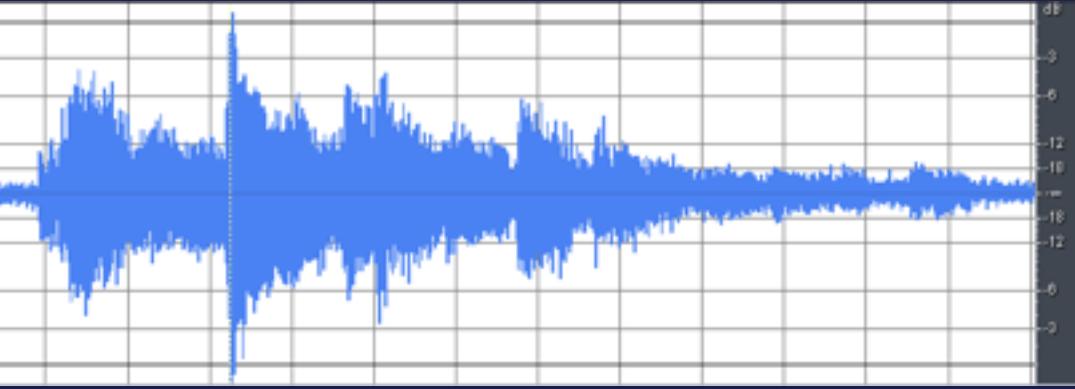
# Deep Learning

A set of very complex functions

- Based on “neuron” and being “deep”
- Architecture & parameters

Course by  
Hung-Yi Lee

[http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_MLSD15\\_2.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html)

$$f(\text{}) = \text{"5"} \quad f(\text{}) = \text{"Hi good day"}$$
$$f(\text{}) = \text{"dog run in the water"}$$

# Neural Network

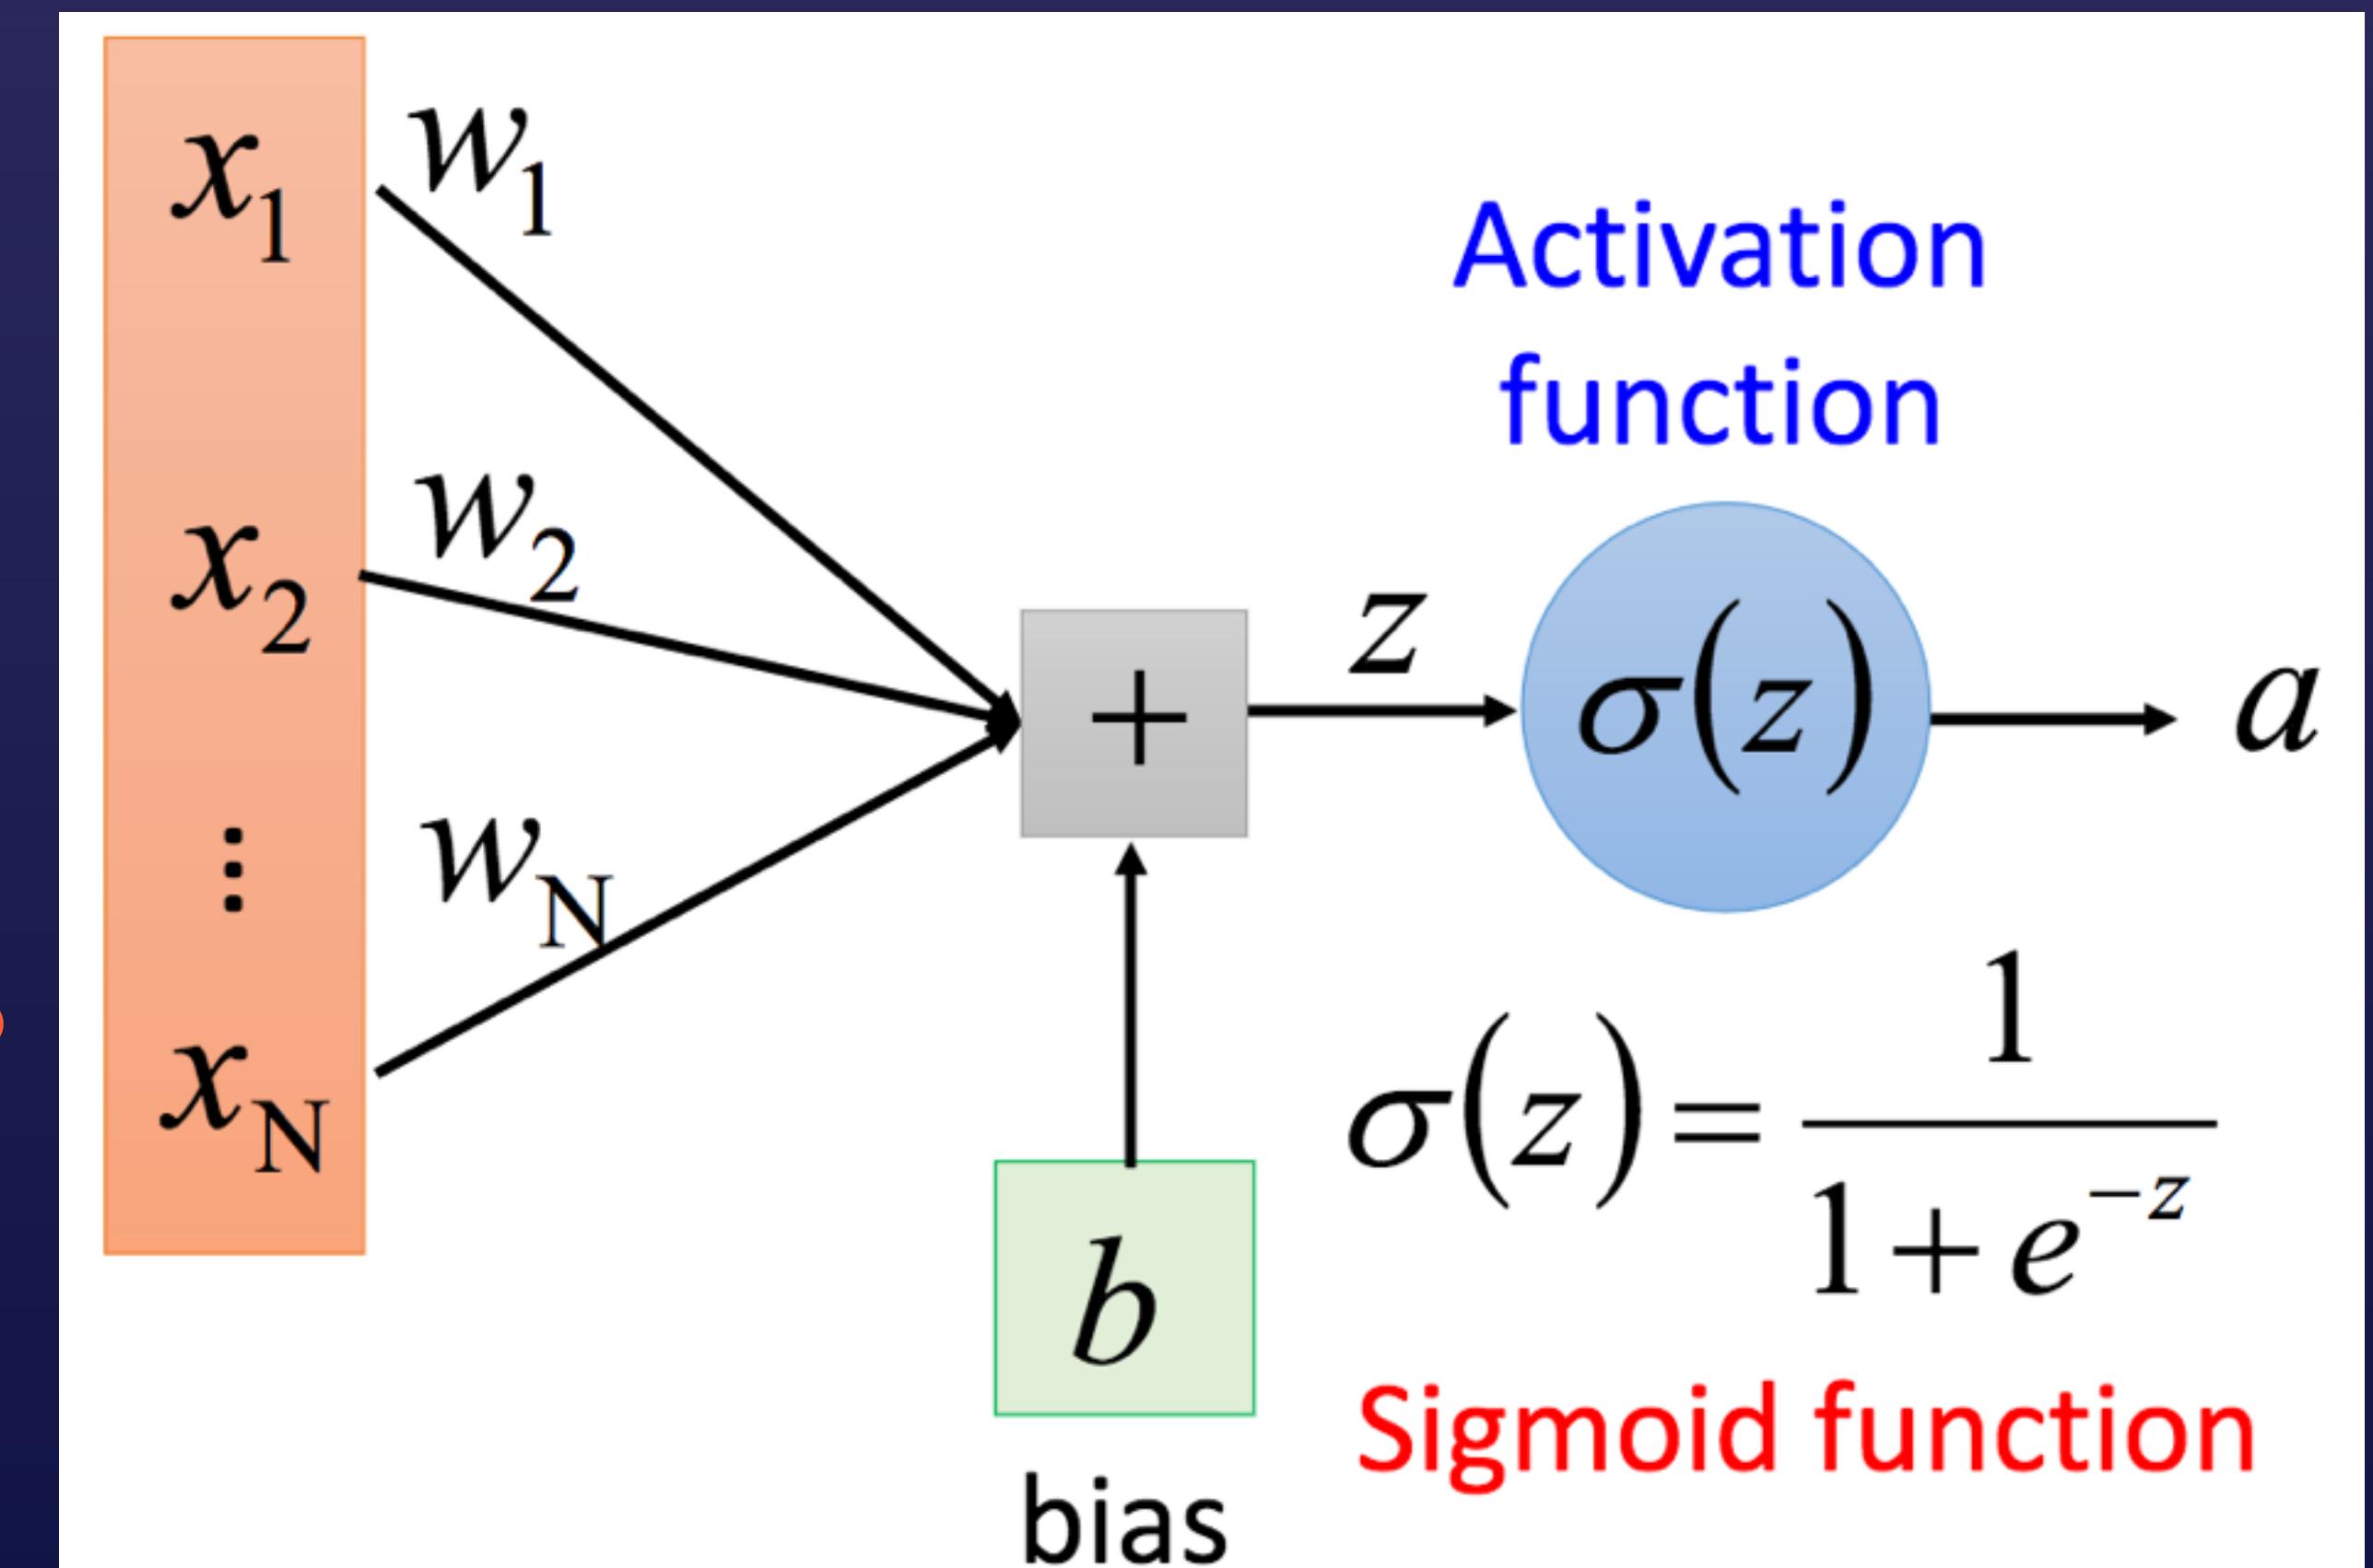
## Logistic Regression

- Just a single neuron with “sigmoid” activation function

Many neurons ?

Different activation functions ?

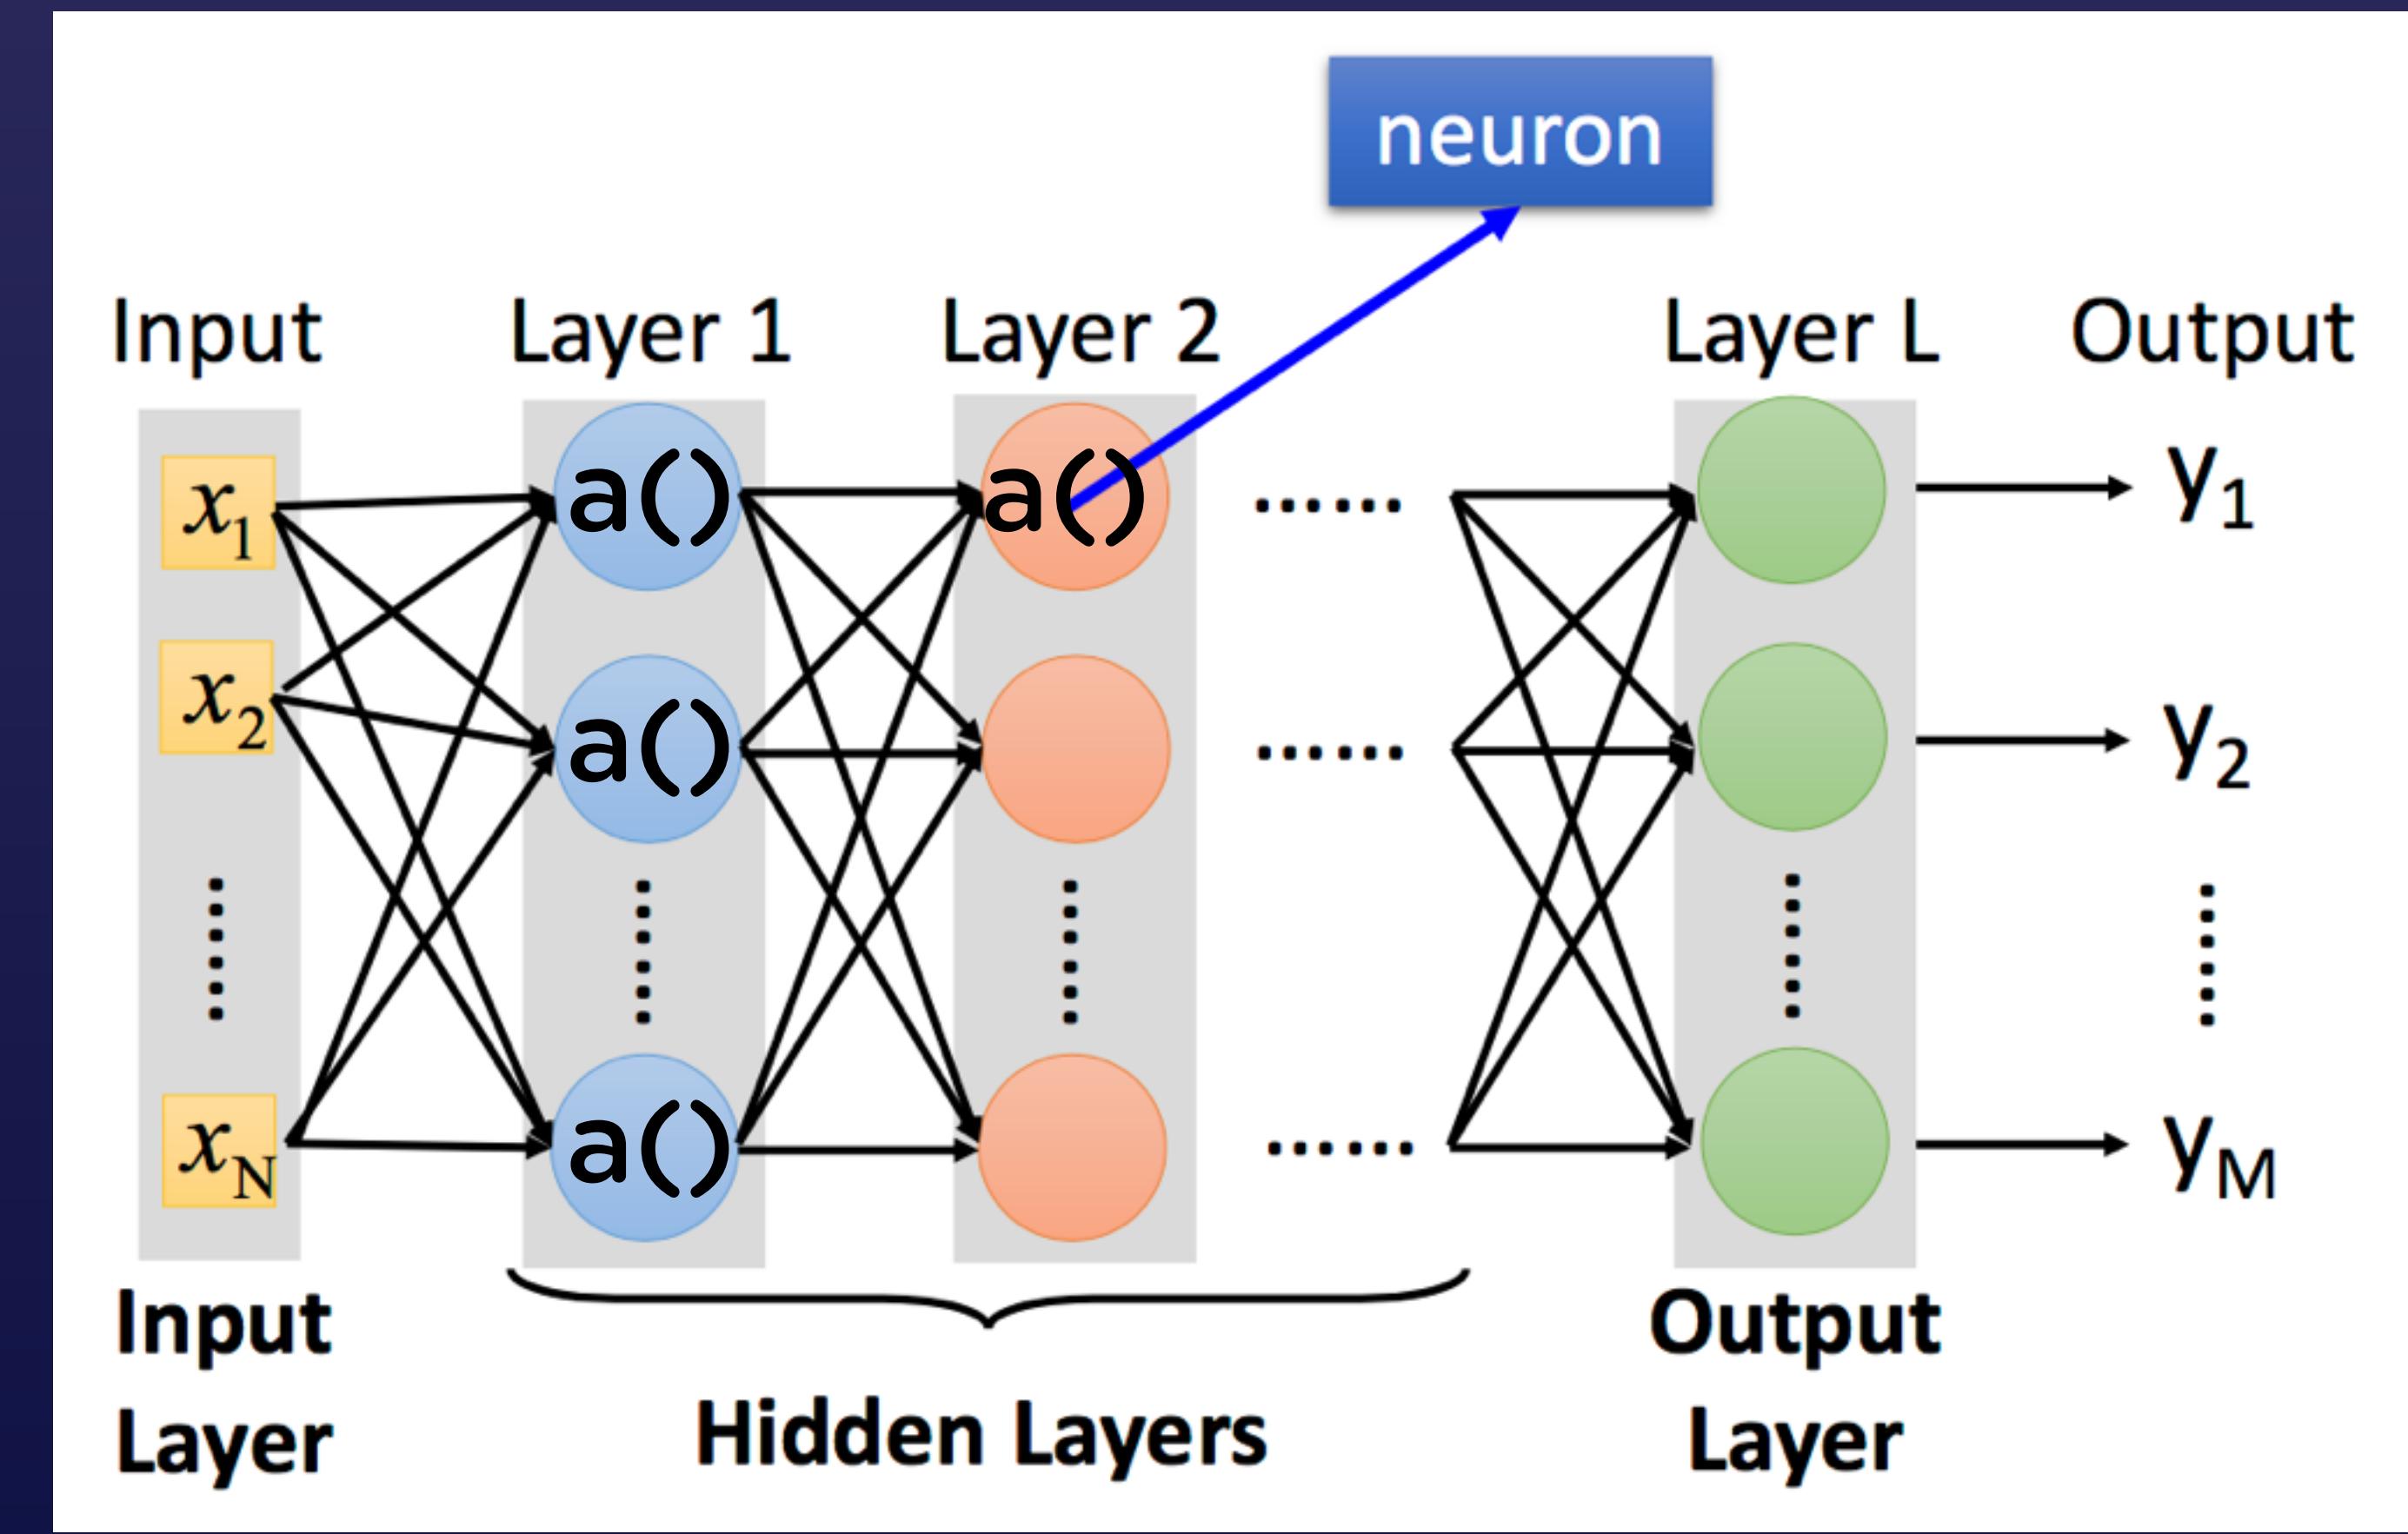
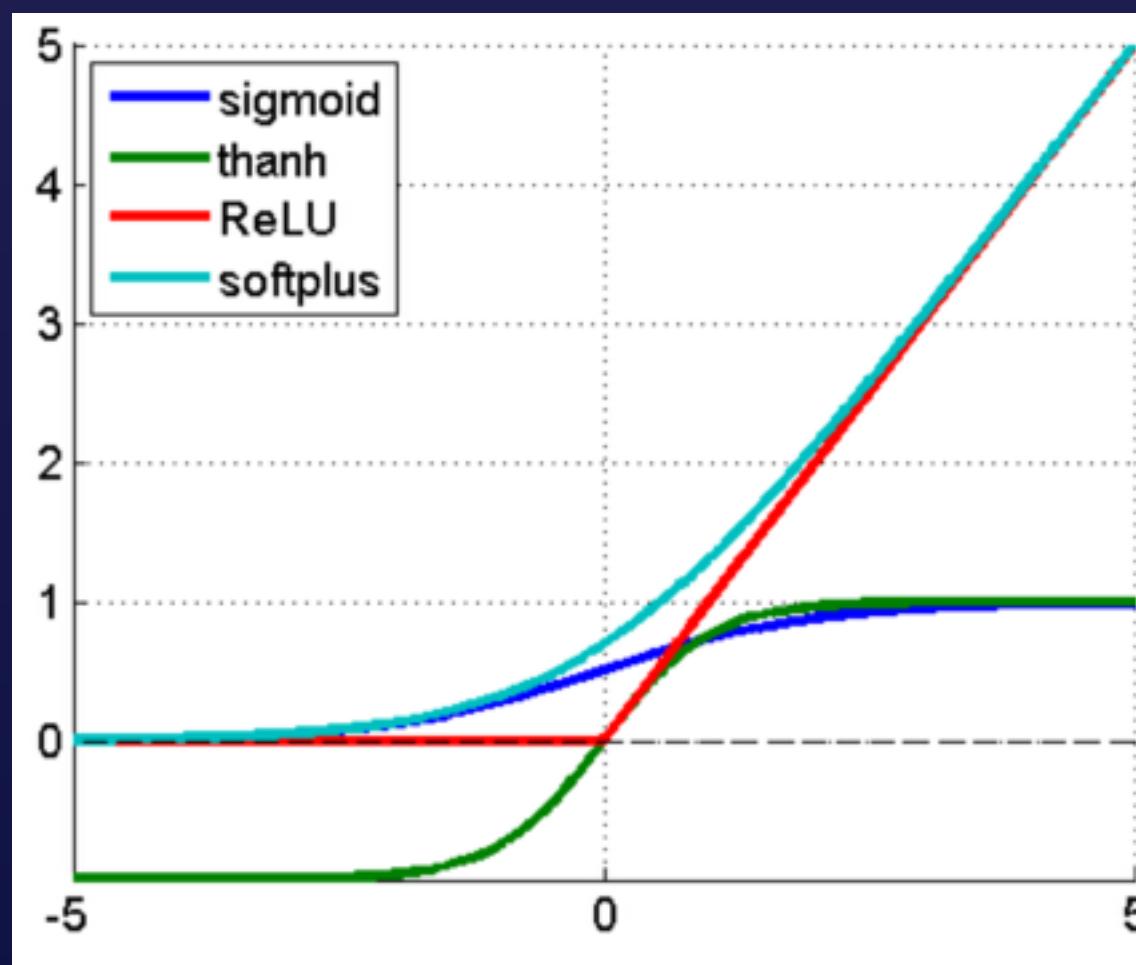
More layers (Cascade) ?



# Deep Neural Network

## Fully Connected

- Deep: many hidden layers
- Hidden layer: many neurons
- Neuron: non-linear transform of weighted sum of inputs
- Transform:



Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_ML16.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_ML16.html)  
<https://imiloainf.wordpress.com/2013/11/06/rectifier-nonlinearities/>

Very Rich Set  
of Hypothesis !



GDG DevFest  
TAIPEI 2016



GLIA CLOUD

# Deep Neural Network

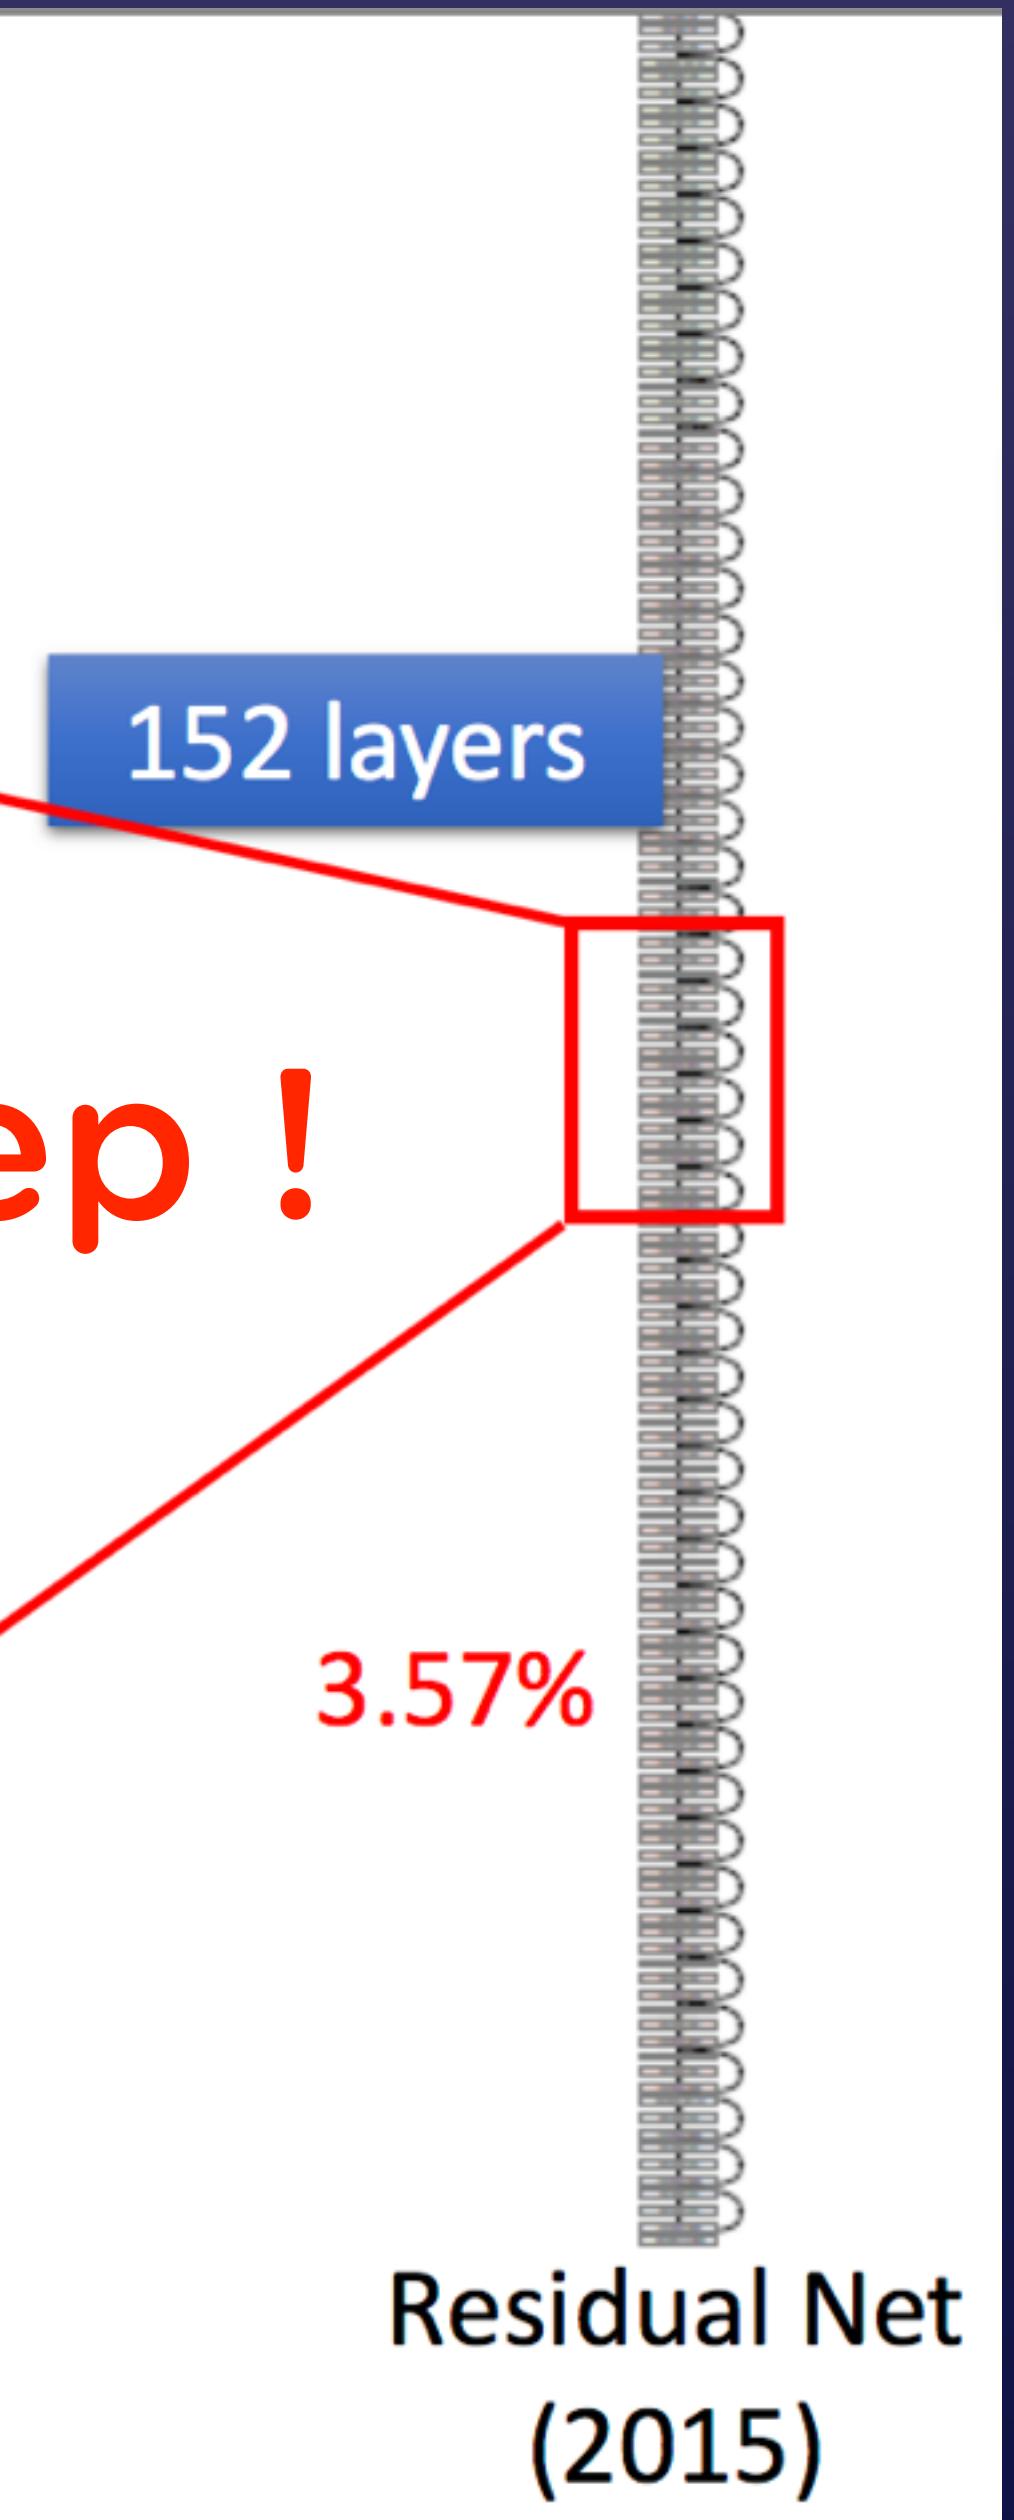
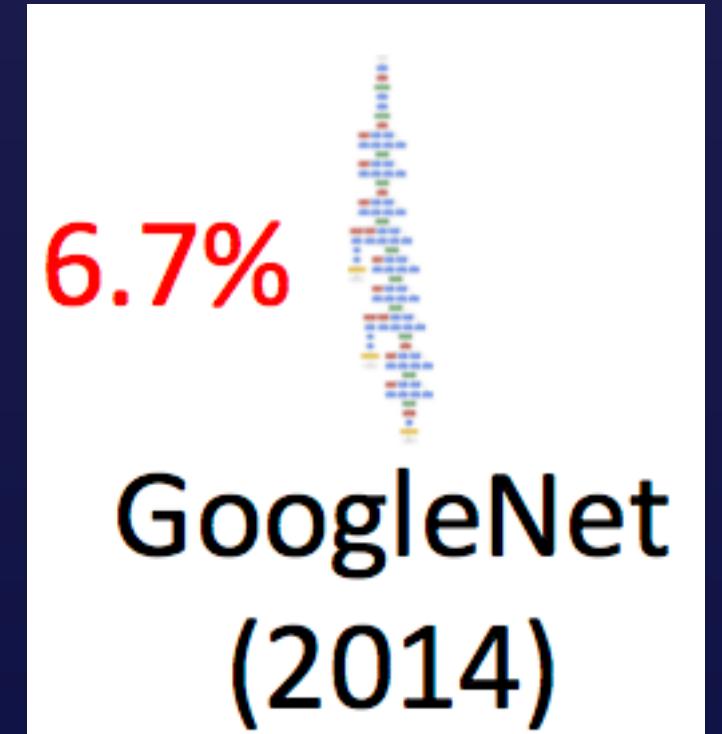
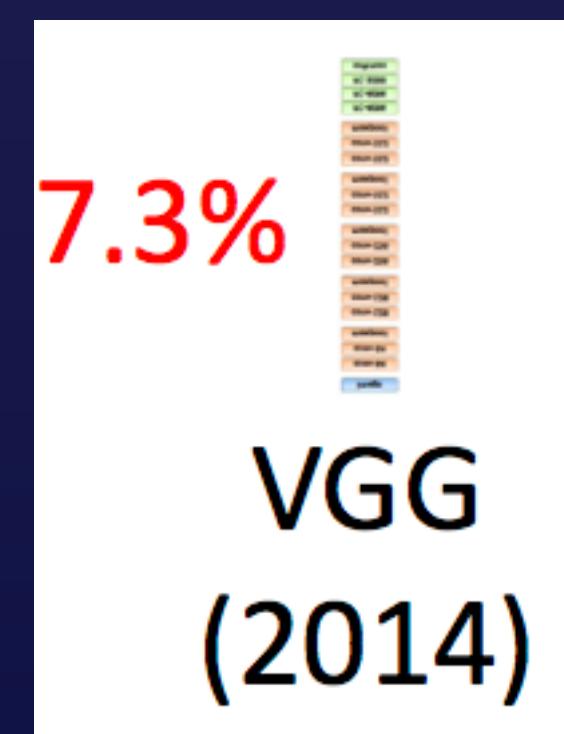
## ImageNet

### Large Scale Visual Recognition

#### Challenge (ILSVRC)

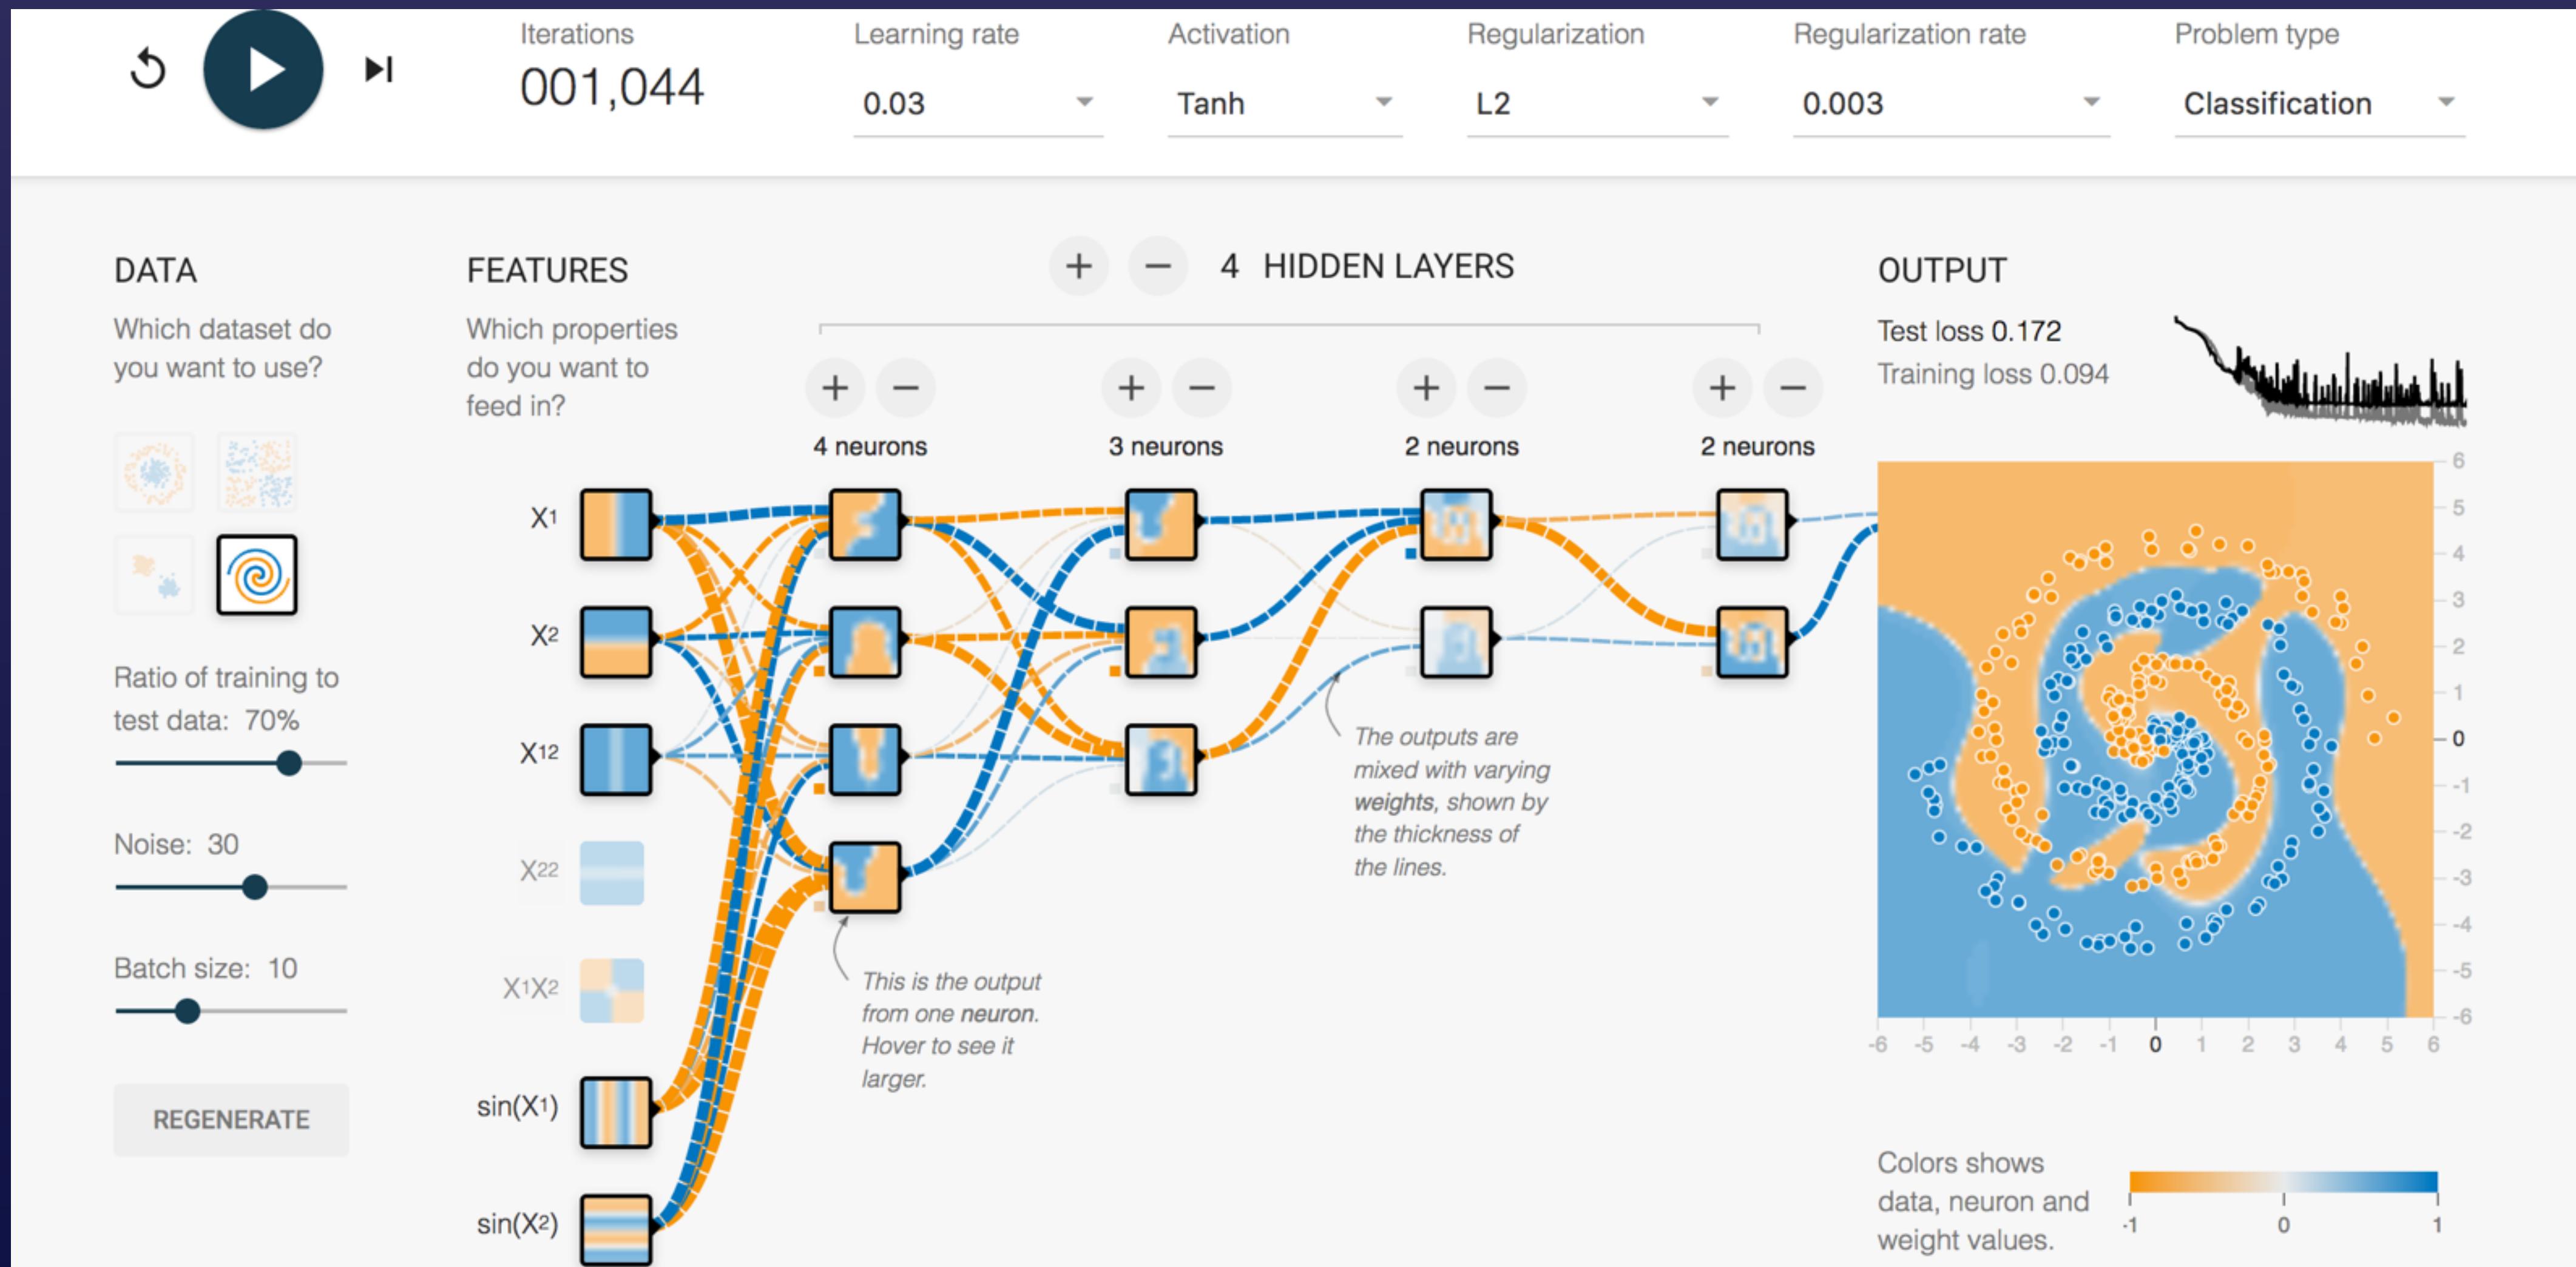
- Classification: 1000 categories
- Human error rate: 5.1%

Go Deep !



Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_ML16.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_ML16.html)

# TensorFlow Playground



Reference : <http://playground.tensorflow.org>



GDG DevFest  
TAIPEI 2016



GLIA CLOUD

# DNN Example : Autoencoder

## MNIST

- For hand-written digits recognition
- Each digit is a 28 \* 28 BW images (0.0 - 1.0 array)
- 10 categories: 0 - 9
- Demo: [https://www.youtube.com/watch?v=FwFduRA\\_L6Q](https://www.youtube.com/watch?v=FwFduRA_L6Q)

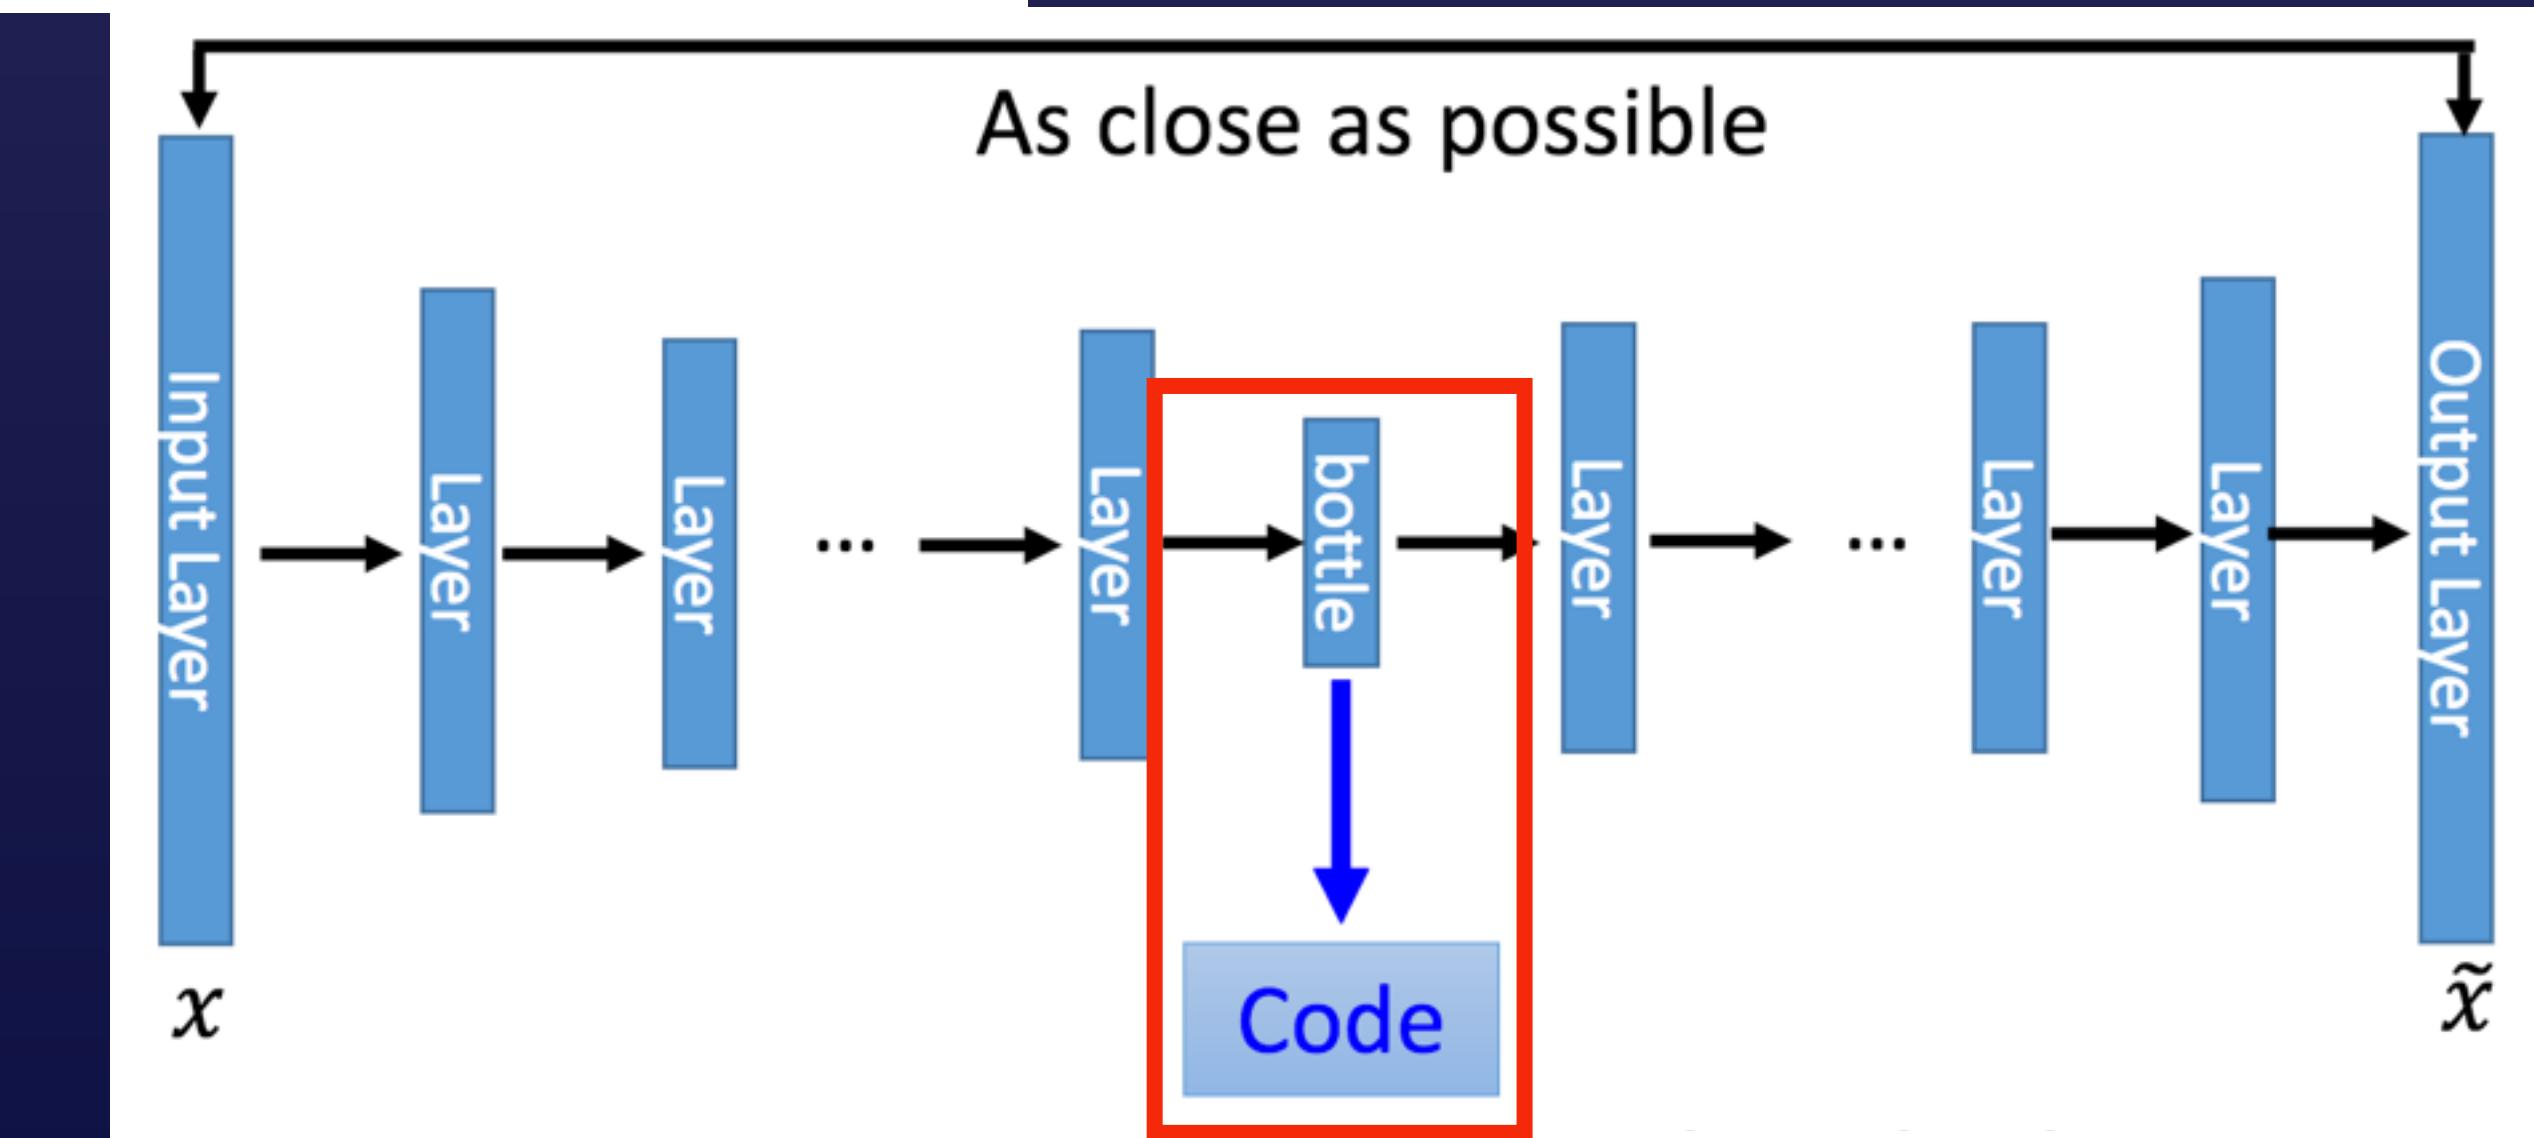
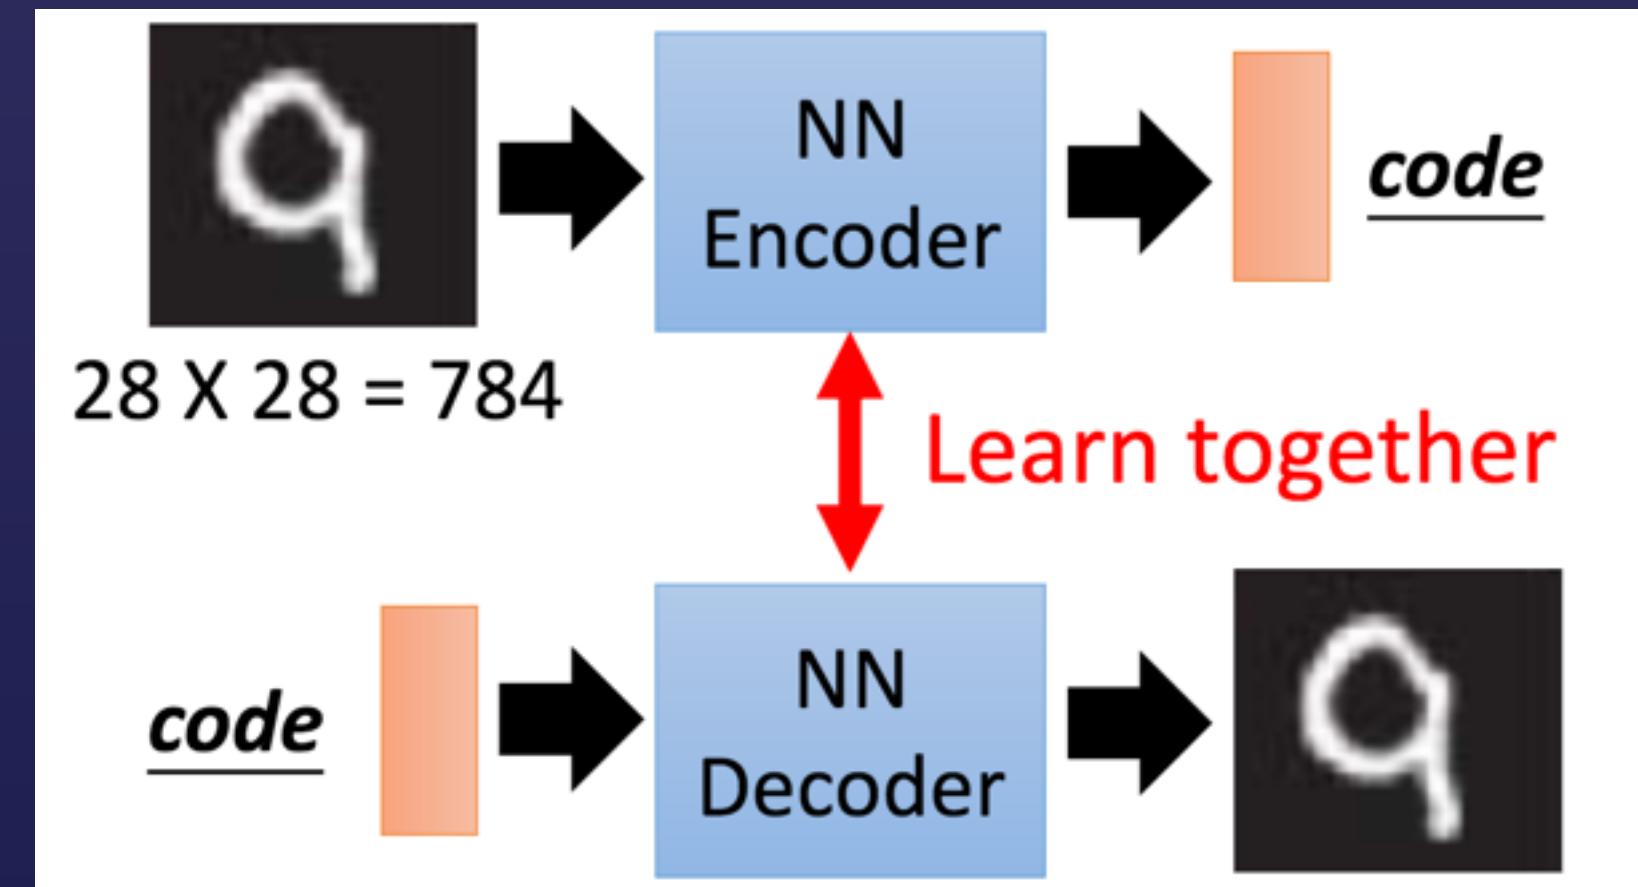
Reference : <https://www.tensorflow.org/>



# DNN Example : Autoencoder

## Auto-encoder

- Deep Neural Network as “Encoder” & “Decoder”
- Unsupervised (Dimension reduction)
- “Representation” of the image



Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_MLSD15\\_2.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html)

# DNN Example : Autoencoder

## Encoder

```
def encoder(x, scope, n_input, n_hidden, n_code, n_layer):
    with tf.variable_scope(scope):
        W = tf.get_variable("W_in", [n_input, n_hidden])
        b = tf.get_variable("b_in", [n_hidden])
        layer_i = tf.nn.tanh(tf.add(tf.matmul(x, W), b))

        for i in range(n_layer - 1):
            W = tf.get_variable("W_en_%d" % i, [n_hidden, n_hidden])
            b = tf.get_variable("b_en_%d" % i, [n_hidden])
            layer_i = tf.nn.tanh(tf.add(tf.matmul(layer_i, W), b))

        W = tf.get_variable("W_encode", [n_hidden, n_code])
        b = tf.get_variable("b_encode", [n_code])
        code = tf.nn.tanh(tf.add(tf.matmul(layer_i, W), b))

    return code
```

# DNN Example : Autoencoder

## Decoder

```
def decoder(code, scope, n_output, n_hidden, n_code, n_layer):
    with tf.variable_scope(scope):
        W = tf.get_variable("W_decode", [n_code, n_hidden])
        b = tf.get_variable("b_decode", [n_hidden])
        layer_i = tf.nn.tanh(tf.add(tf.matmul(code, W), b))

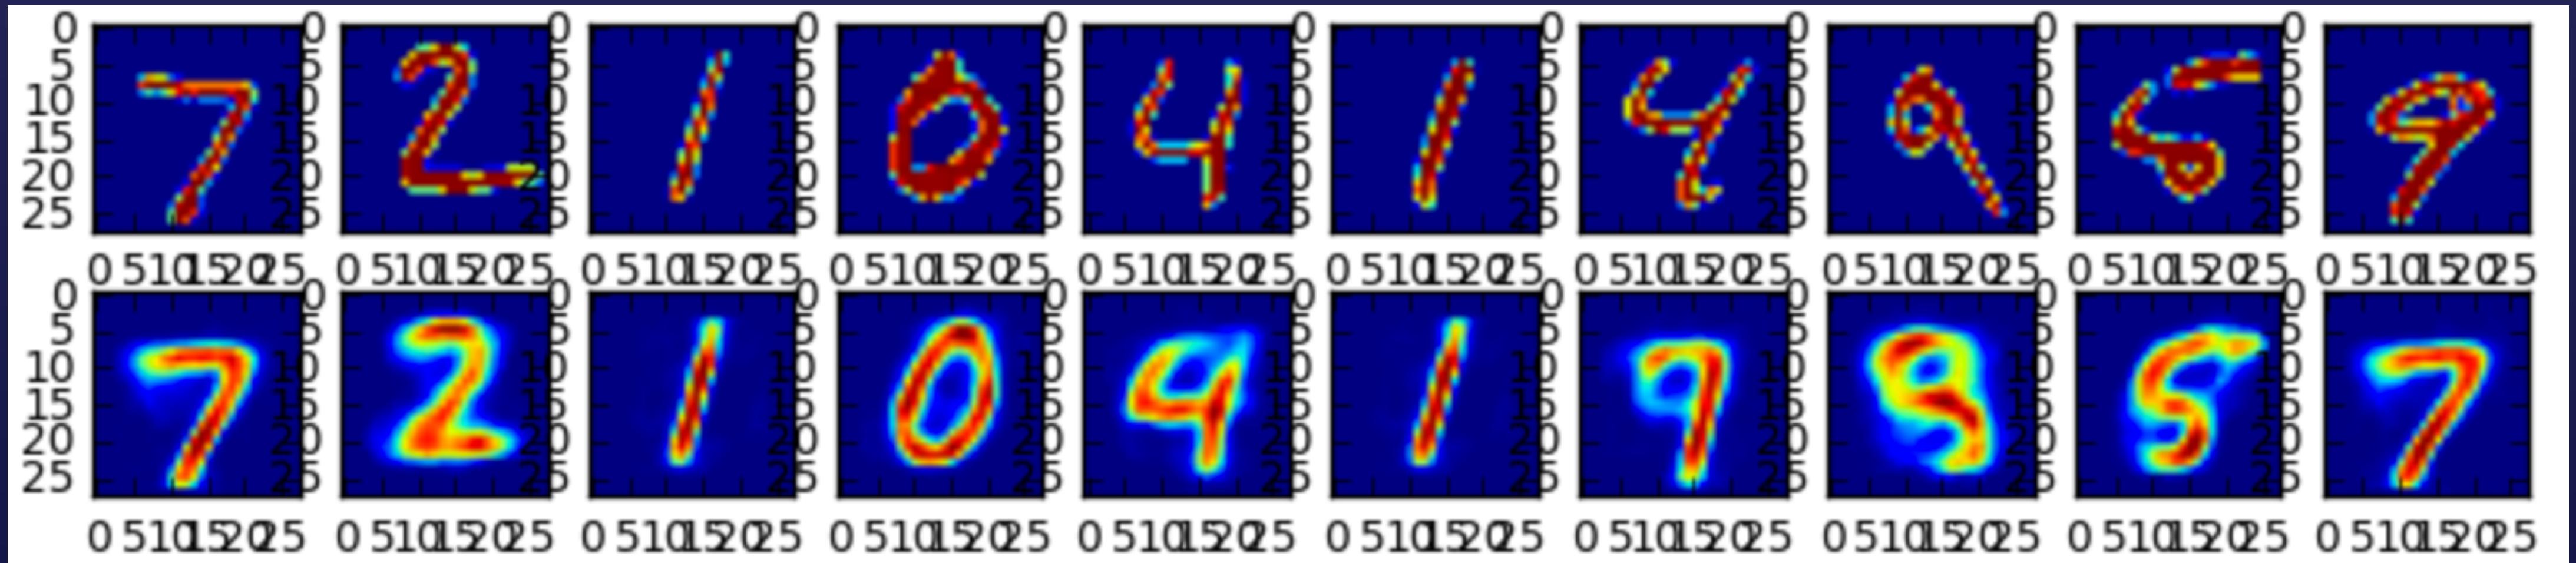
        for i in range(n_layer - 1):
            W = tf.get_variable("W_de_%d" % i, [n_hidden, n_hidden])
            b = tf.get_variable("b_de_%d" % i, [n_hidden])
            layer_i = tf.nn.tanh(tf.add(tf.matmul(layer_i, W), b))

        W = tf.get_variable("W_out", [n_hidden, n_output])
        b = tf.get_variable("b_out", [n_output])
        output = tf.nn.sigmoid(tf.add(tf.matmul(layer_i, W), b))

    return output
```

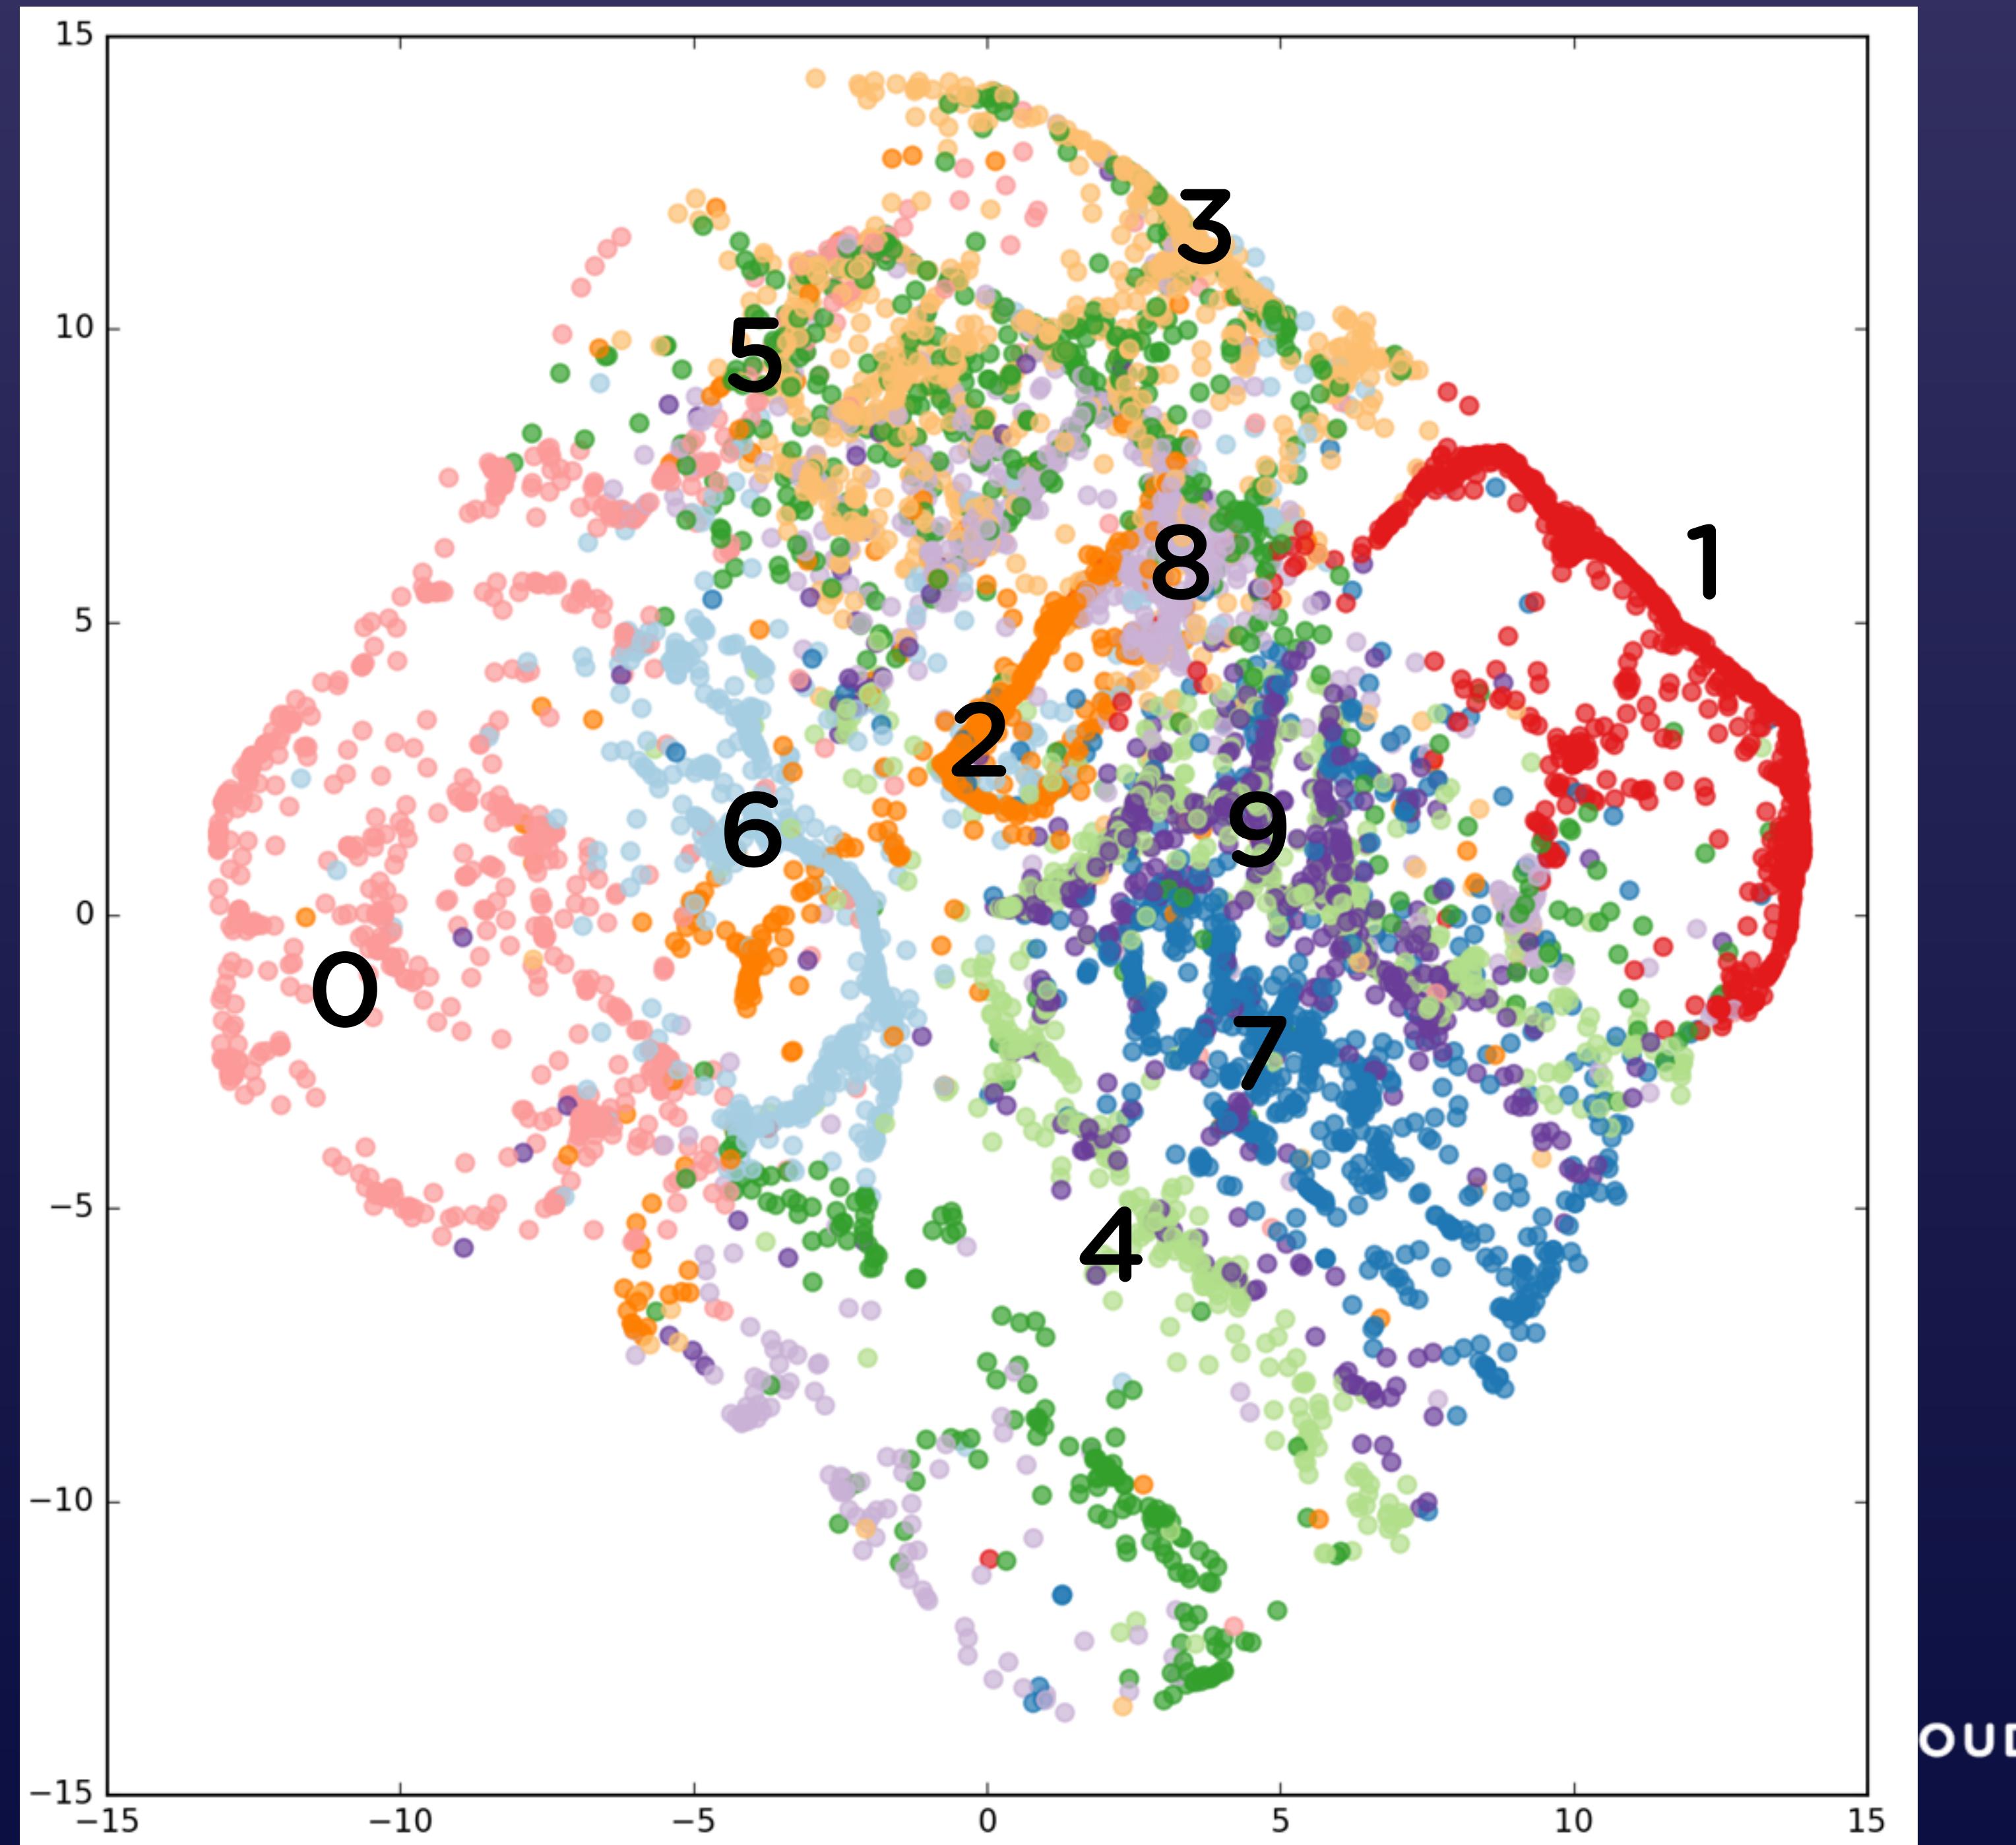
# Autoencoder

## Results



Reference : [https://github.com/AaronYALai/tensorflow\\_activity\\_recognition](https://github.com/AaronYALai/tensorflow_activity_recognition)

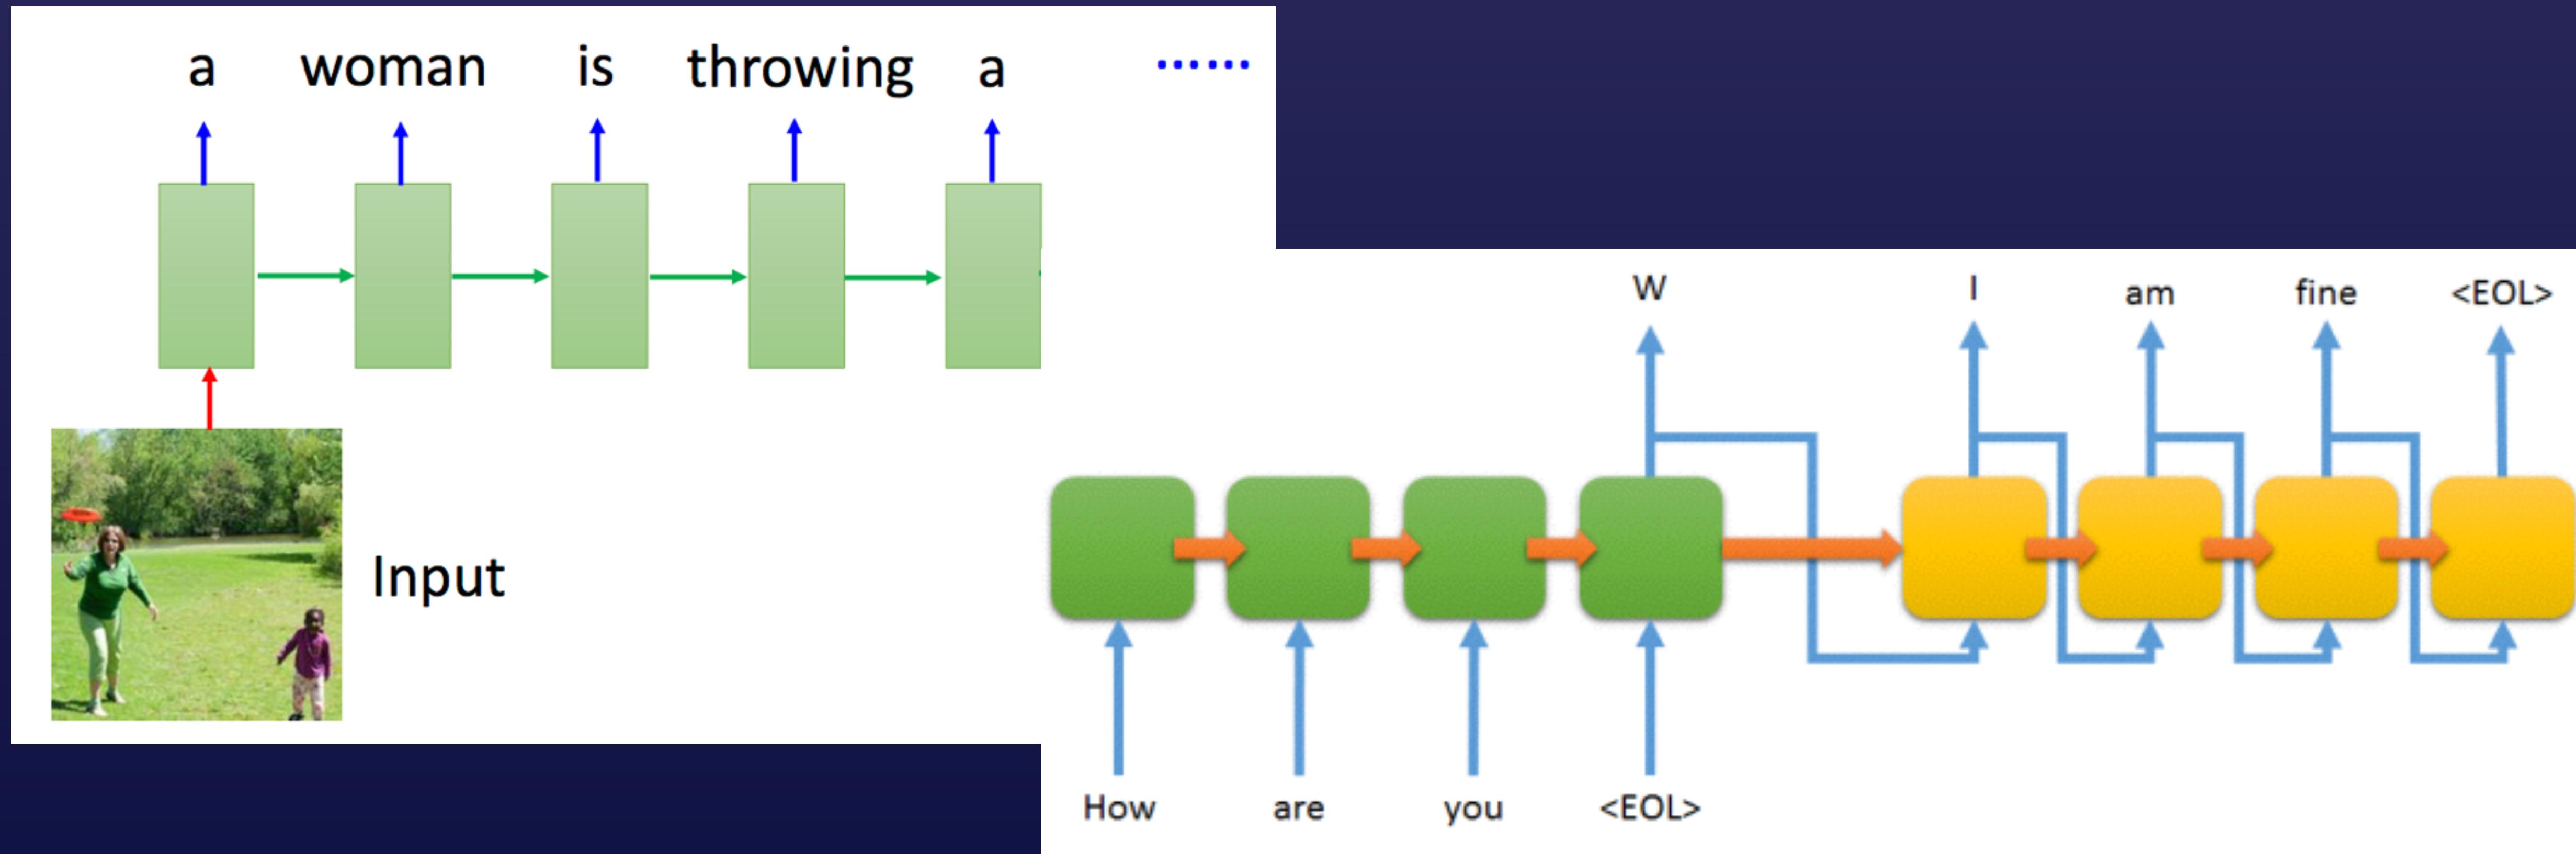
# Autoencoder



Reference : [https://github.com/AaronYALai/tensorflow\\_activity\\_recognition](https://github.com/AaronYALai/tensorflow_activity_recognition)

# Neural Network with Memory

Memory is important for sequence tasks



Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS\\_2015\\_2](http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2)  
<https://github.com/farizrahman4u/seq2seq>

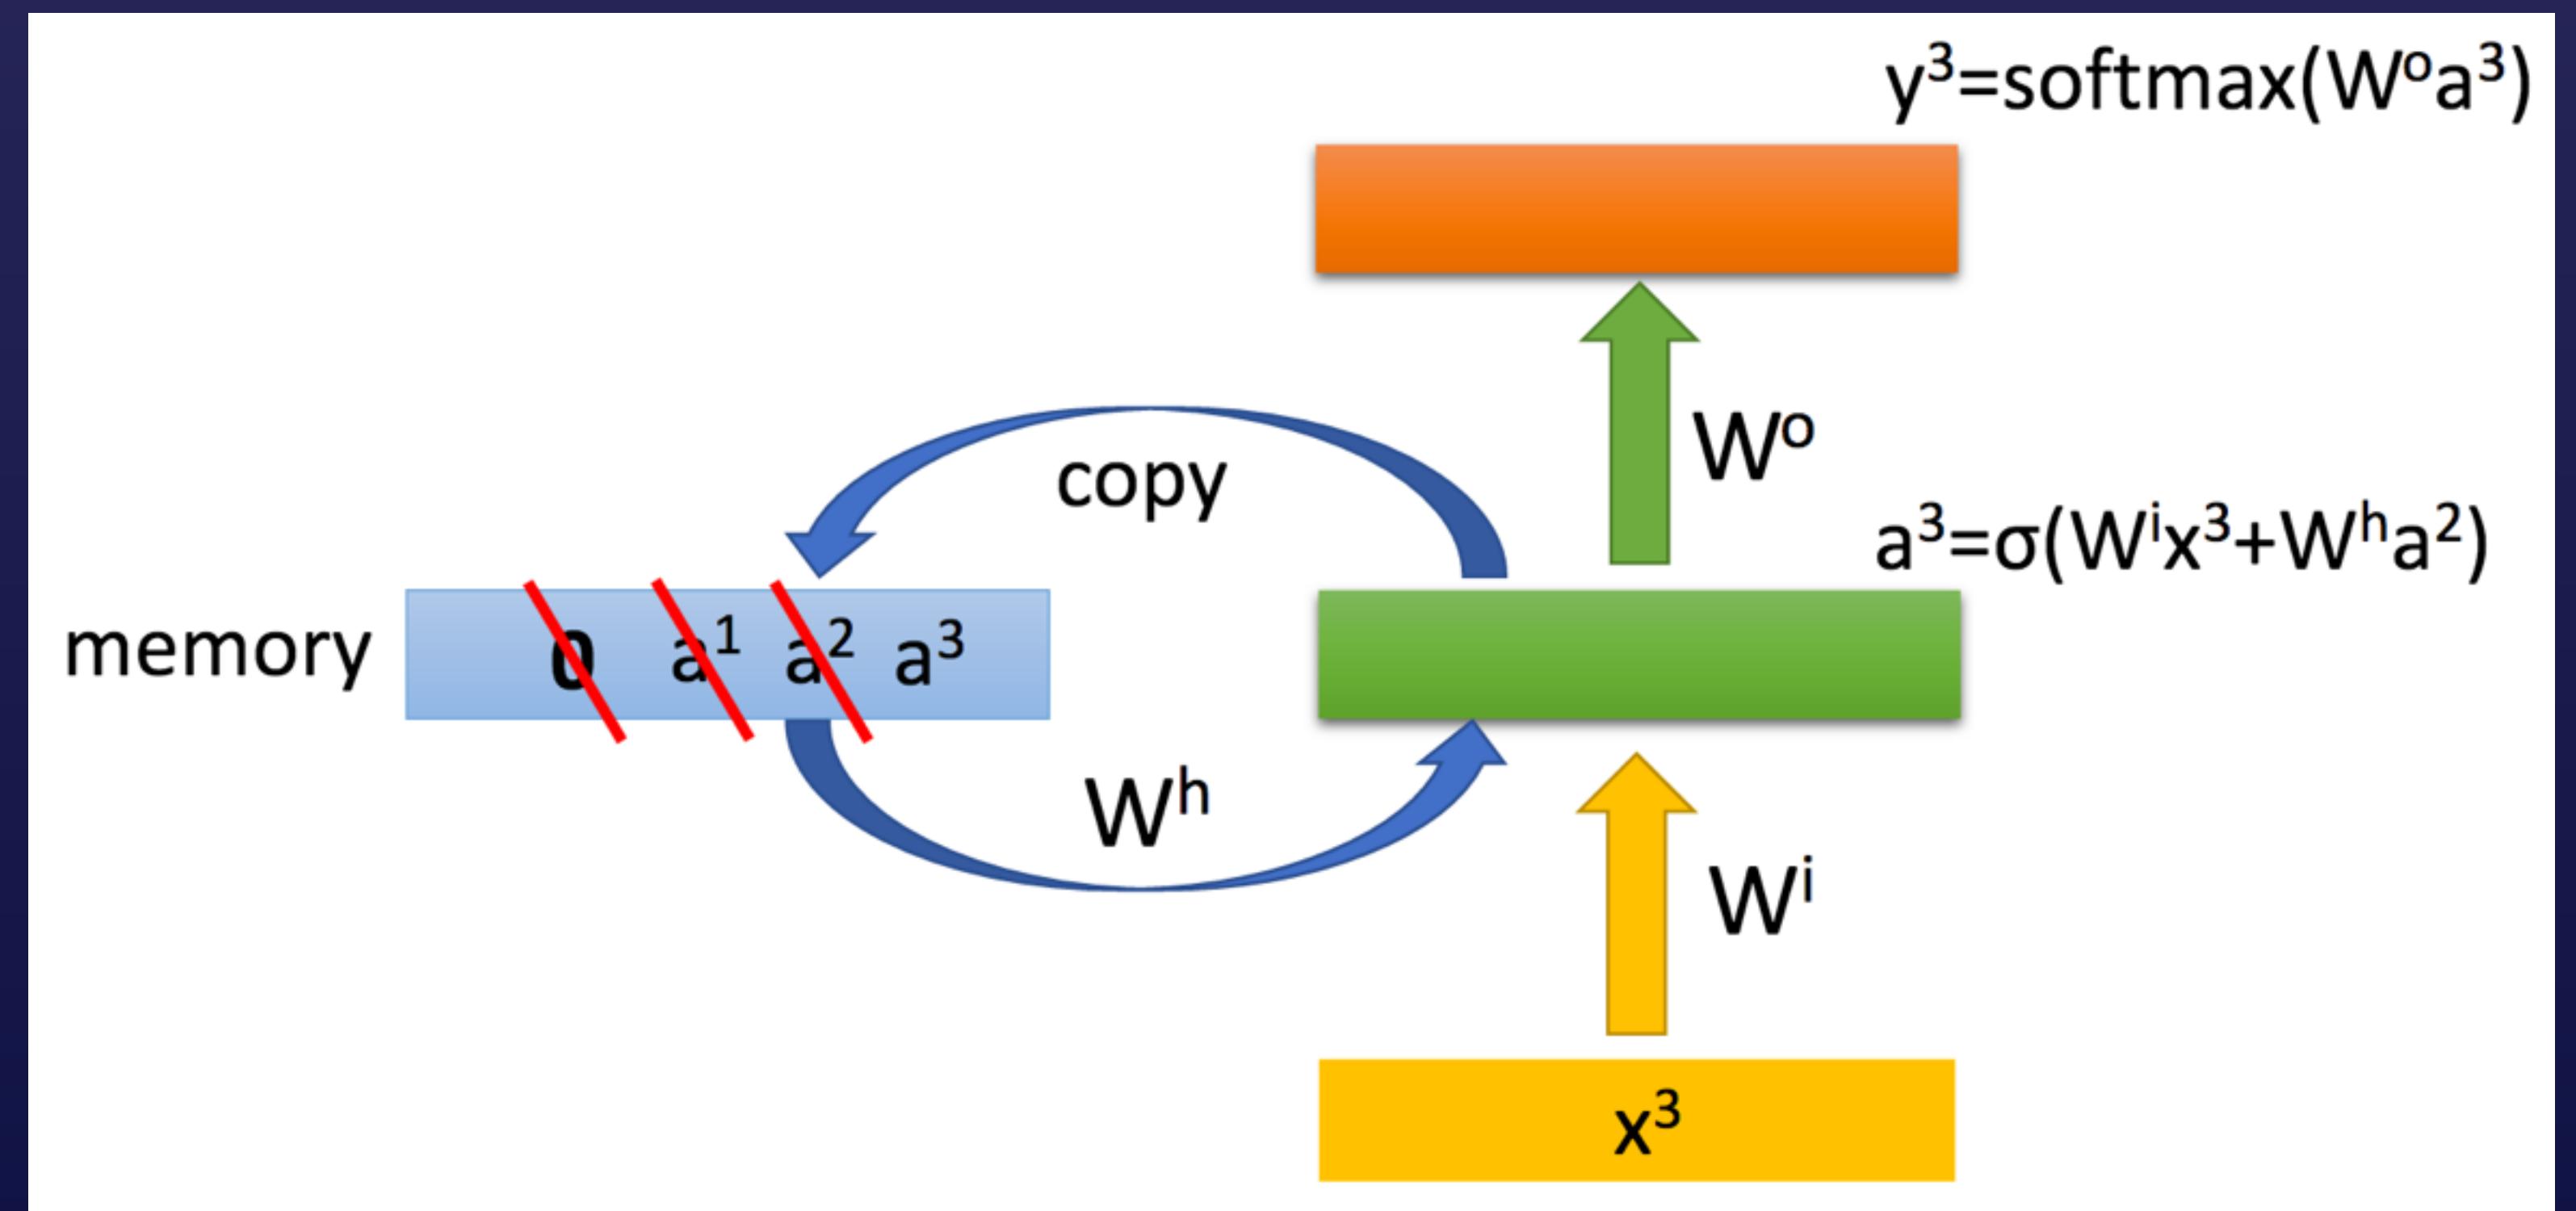
GDG DevFest  
TAIPEI 2016

GLIA CLOUD

# Neural Network with Memory

## Recurrent Neural Network

- Memory cell: pass and update at every step
- “Reuse”:  $W_o$ ,  $W_i$ , and  $W_h$  are shared
- “Sequence”: feature vectors are input “one by one”
- Hard to Train:  $df/dW$ ? Autogradient!

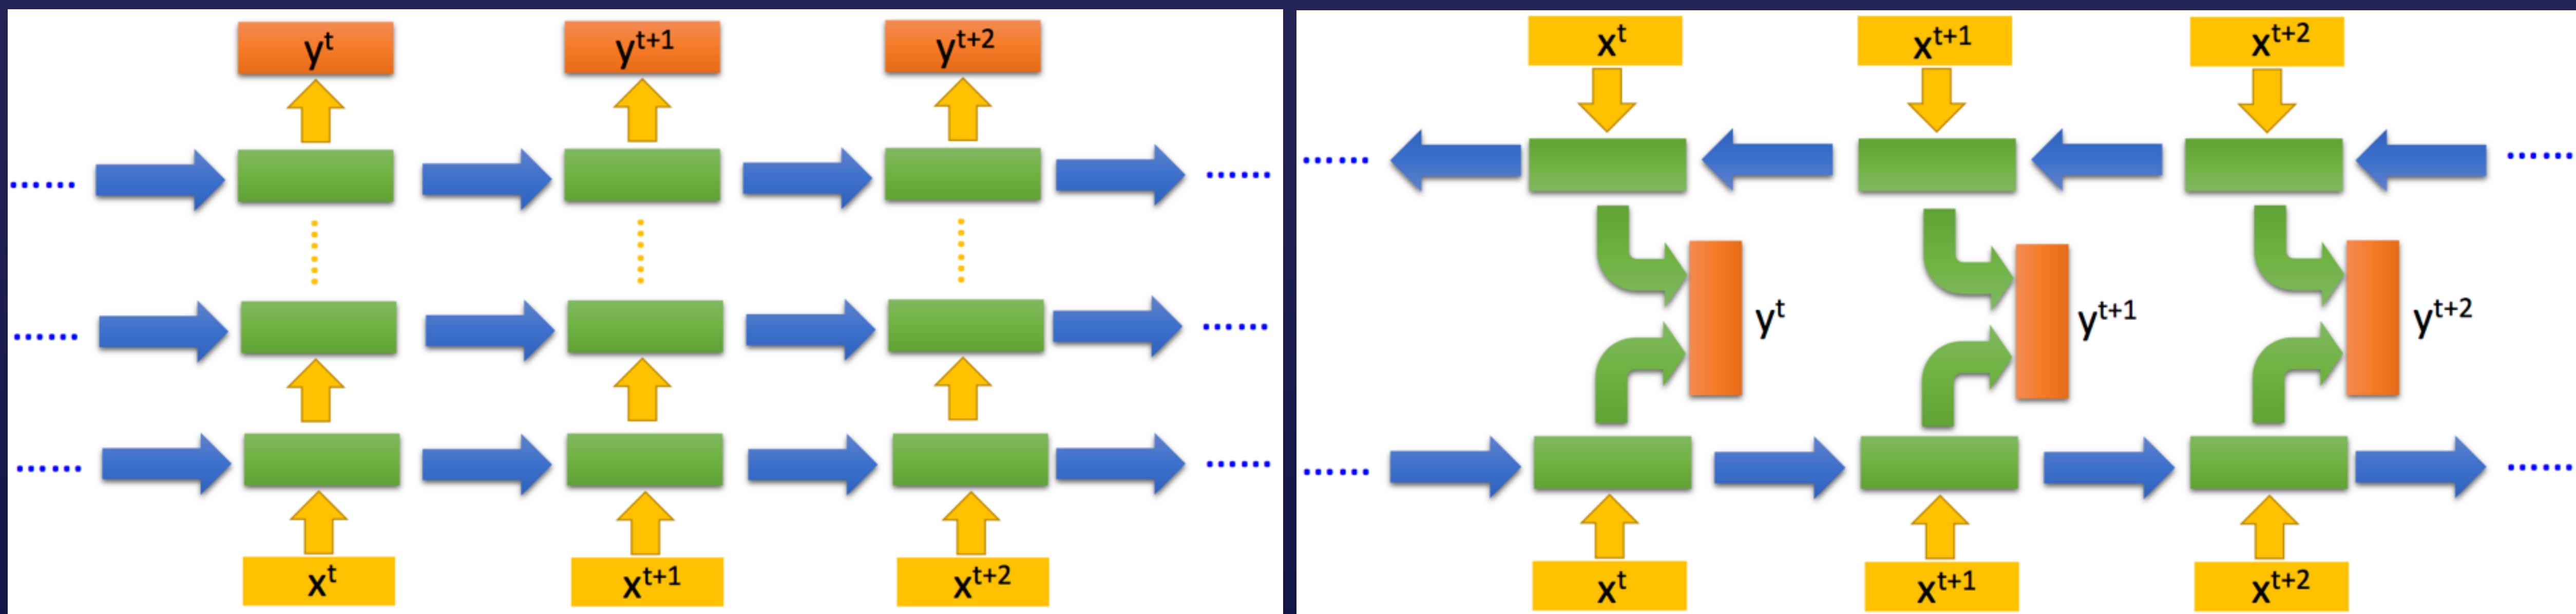


Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_MLSD15\\_2.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html)

# Recurrent Neural Network

Go Deep!

Bidirectional



Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_MLSD15\\_2.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html)

GDG DevFest  
TAIPEI 2016



# Sample Code

- RNN on MNIST
- Break 28\*28 image into 28 vectors (seq)
- Output the last one

```
Start
Iter 256, Loss= 2.118008, Accuracy= 0.31641
Iter 12800, Loss= 0.369384, Accuracy= 0.89453
Iter 25600, Loss= 0.383506, Accuracy= 0.87109
Iter 38400, Loss= 0.284901, Accuracy= 0.92578
Iter 51200, Loss= 0.220428, Accuracy= 0.92578
Iter 64000, Loss= 0.216666, Accuracy= 0.94141
Iter 76800, Loss= 0.246528, Accuracy= 0.92969
Iter 89600, Loss= 0.224747, Accuracy= 0.92578
Optimization Finished!
Testing Accuracy: 0.96875
```

```
def RNN(x, n_input, n_hidden, n_classes, n_steps, scope=None):
    with tf.variable_scope(scope or "RNN"):
        # x:[batchsize, n_steps, n_input]; Permuting batch_size and n_steps
        x = tf.transpose(x, [1, 0, 2])
        # Reshape to (n_steps*batch_size, n_input)
        x = tf.reshape(x, [-1, n_input])
        # Split: list of 'n_steps' tensors of shape (batch_size, n_input)
        x = tf.split(0, n_steps, x)

        outputs = []
        memory = tf.get_variable("memory", [1, n_hidden])
        for i in range(len(x)):
            W_o = tf.get_variable("W_out", [n_hidden, n_classes])
            W_h = tf.get_variable("W_mem", [n_hidden, n_hidden])
            W_i = tf.get_variable("W_in", [n_input, n_hidden])

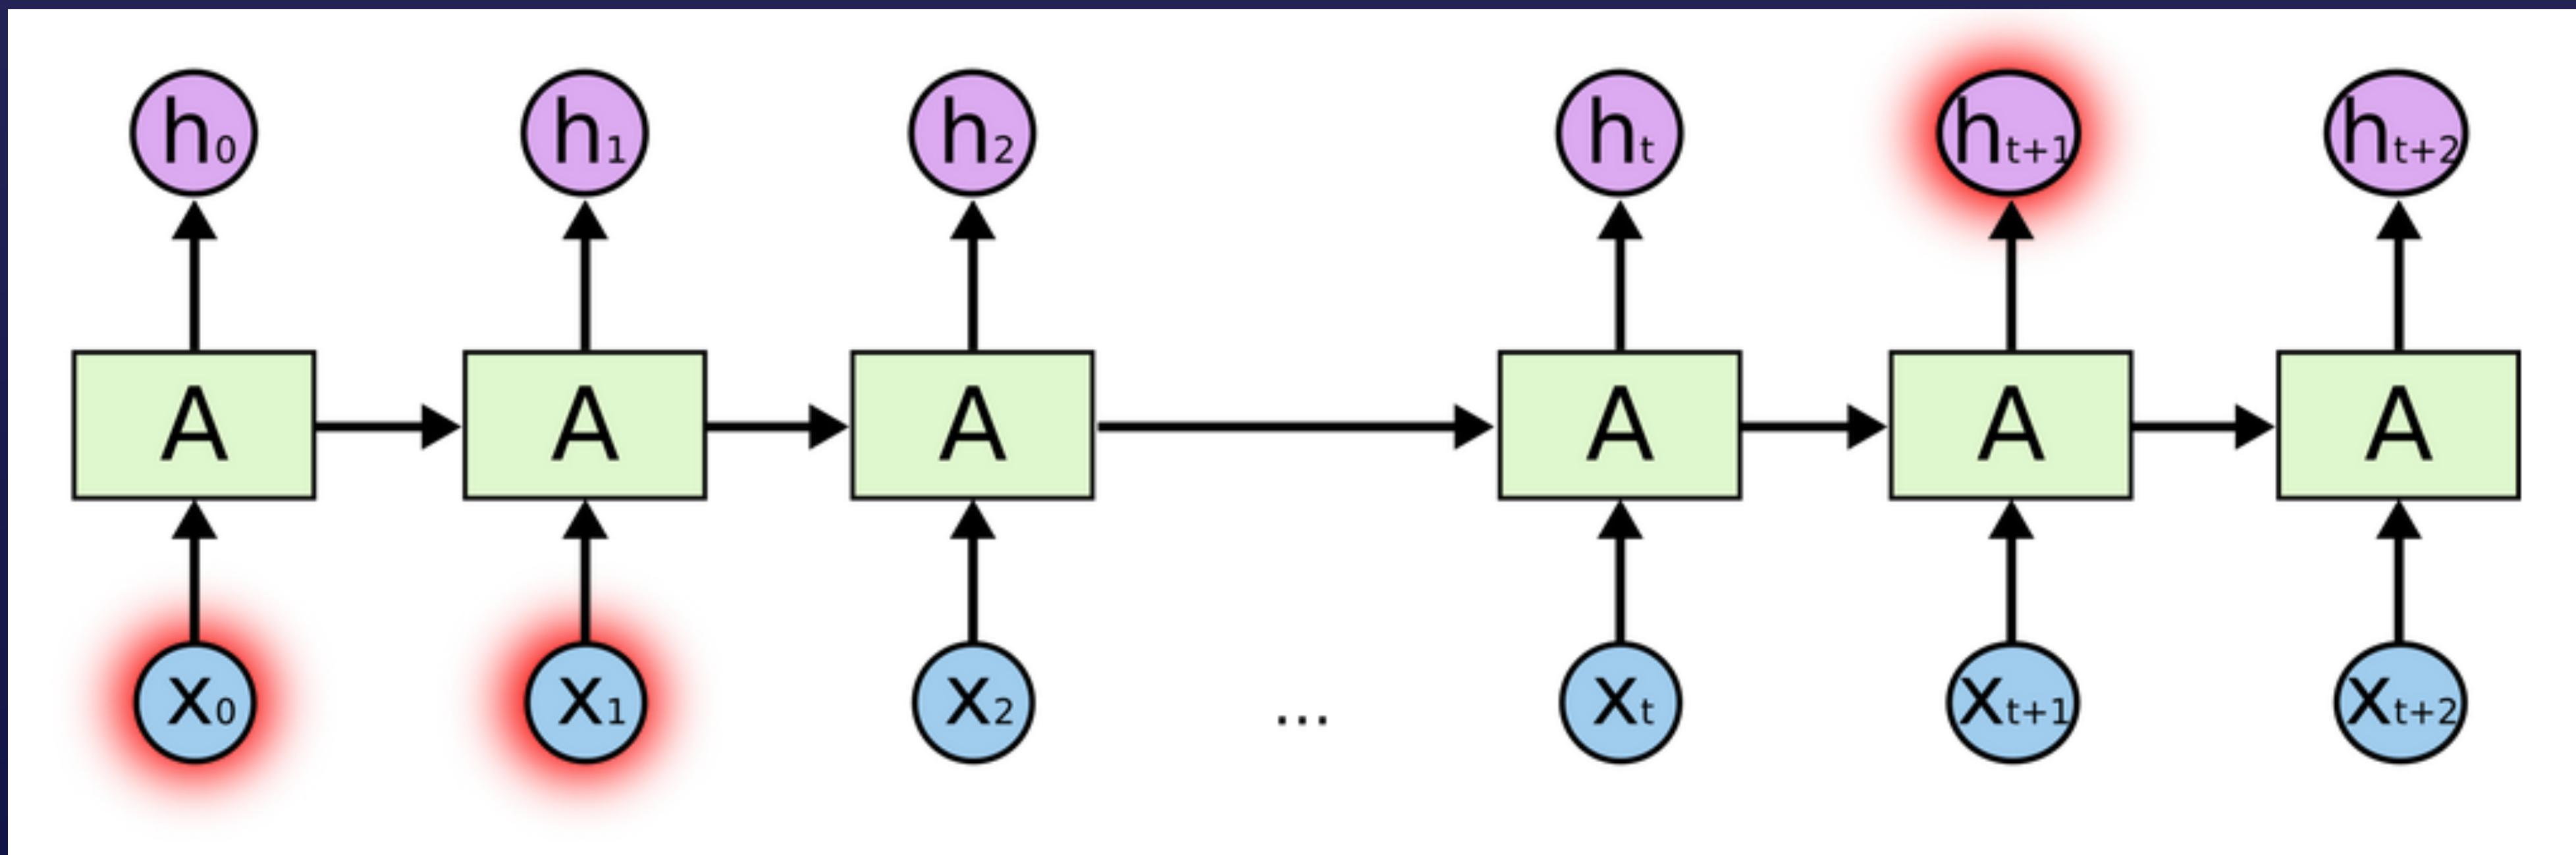
            memory_read = tf.matmul(memory, W_h)
            memory = tf.nn.tanh(tf.add(tf.matmul(x[i], W_i), memory_read))
            outputs.append(tf.matmul(memory, W_o))

            # set the variables to be shared !
            if i == 0:
                tf.get_variable_scope().reuse_variables()

    return outputs[-1]
```

# Recurrent Neural Network

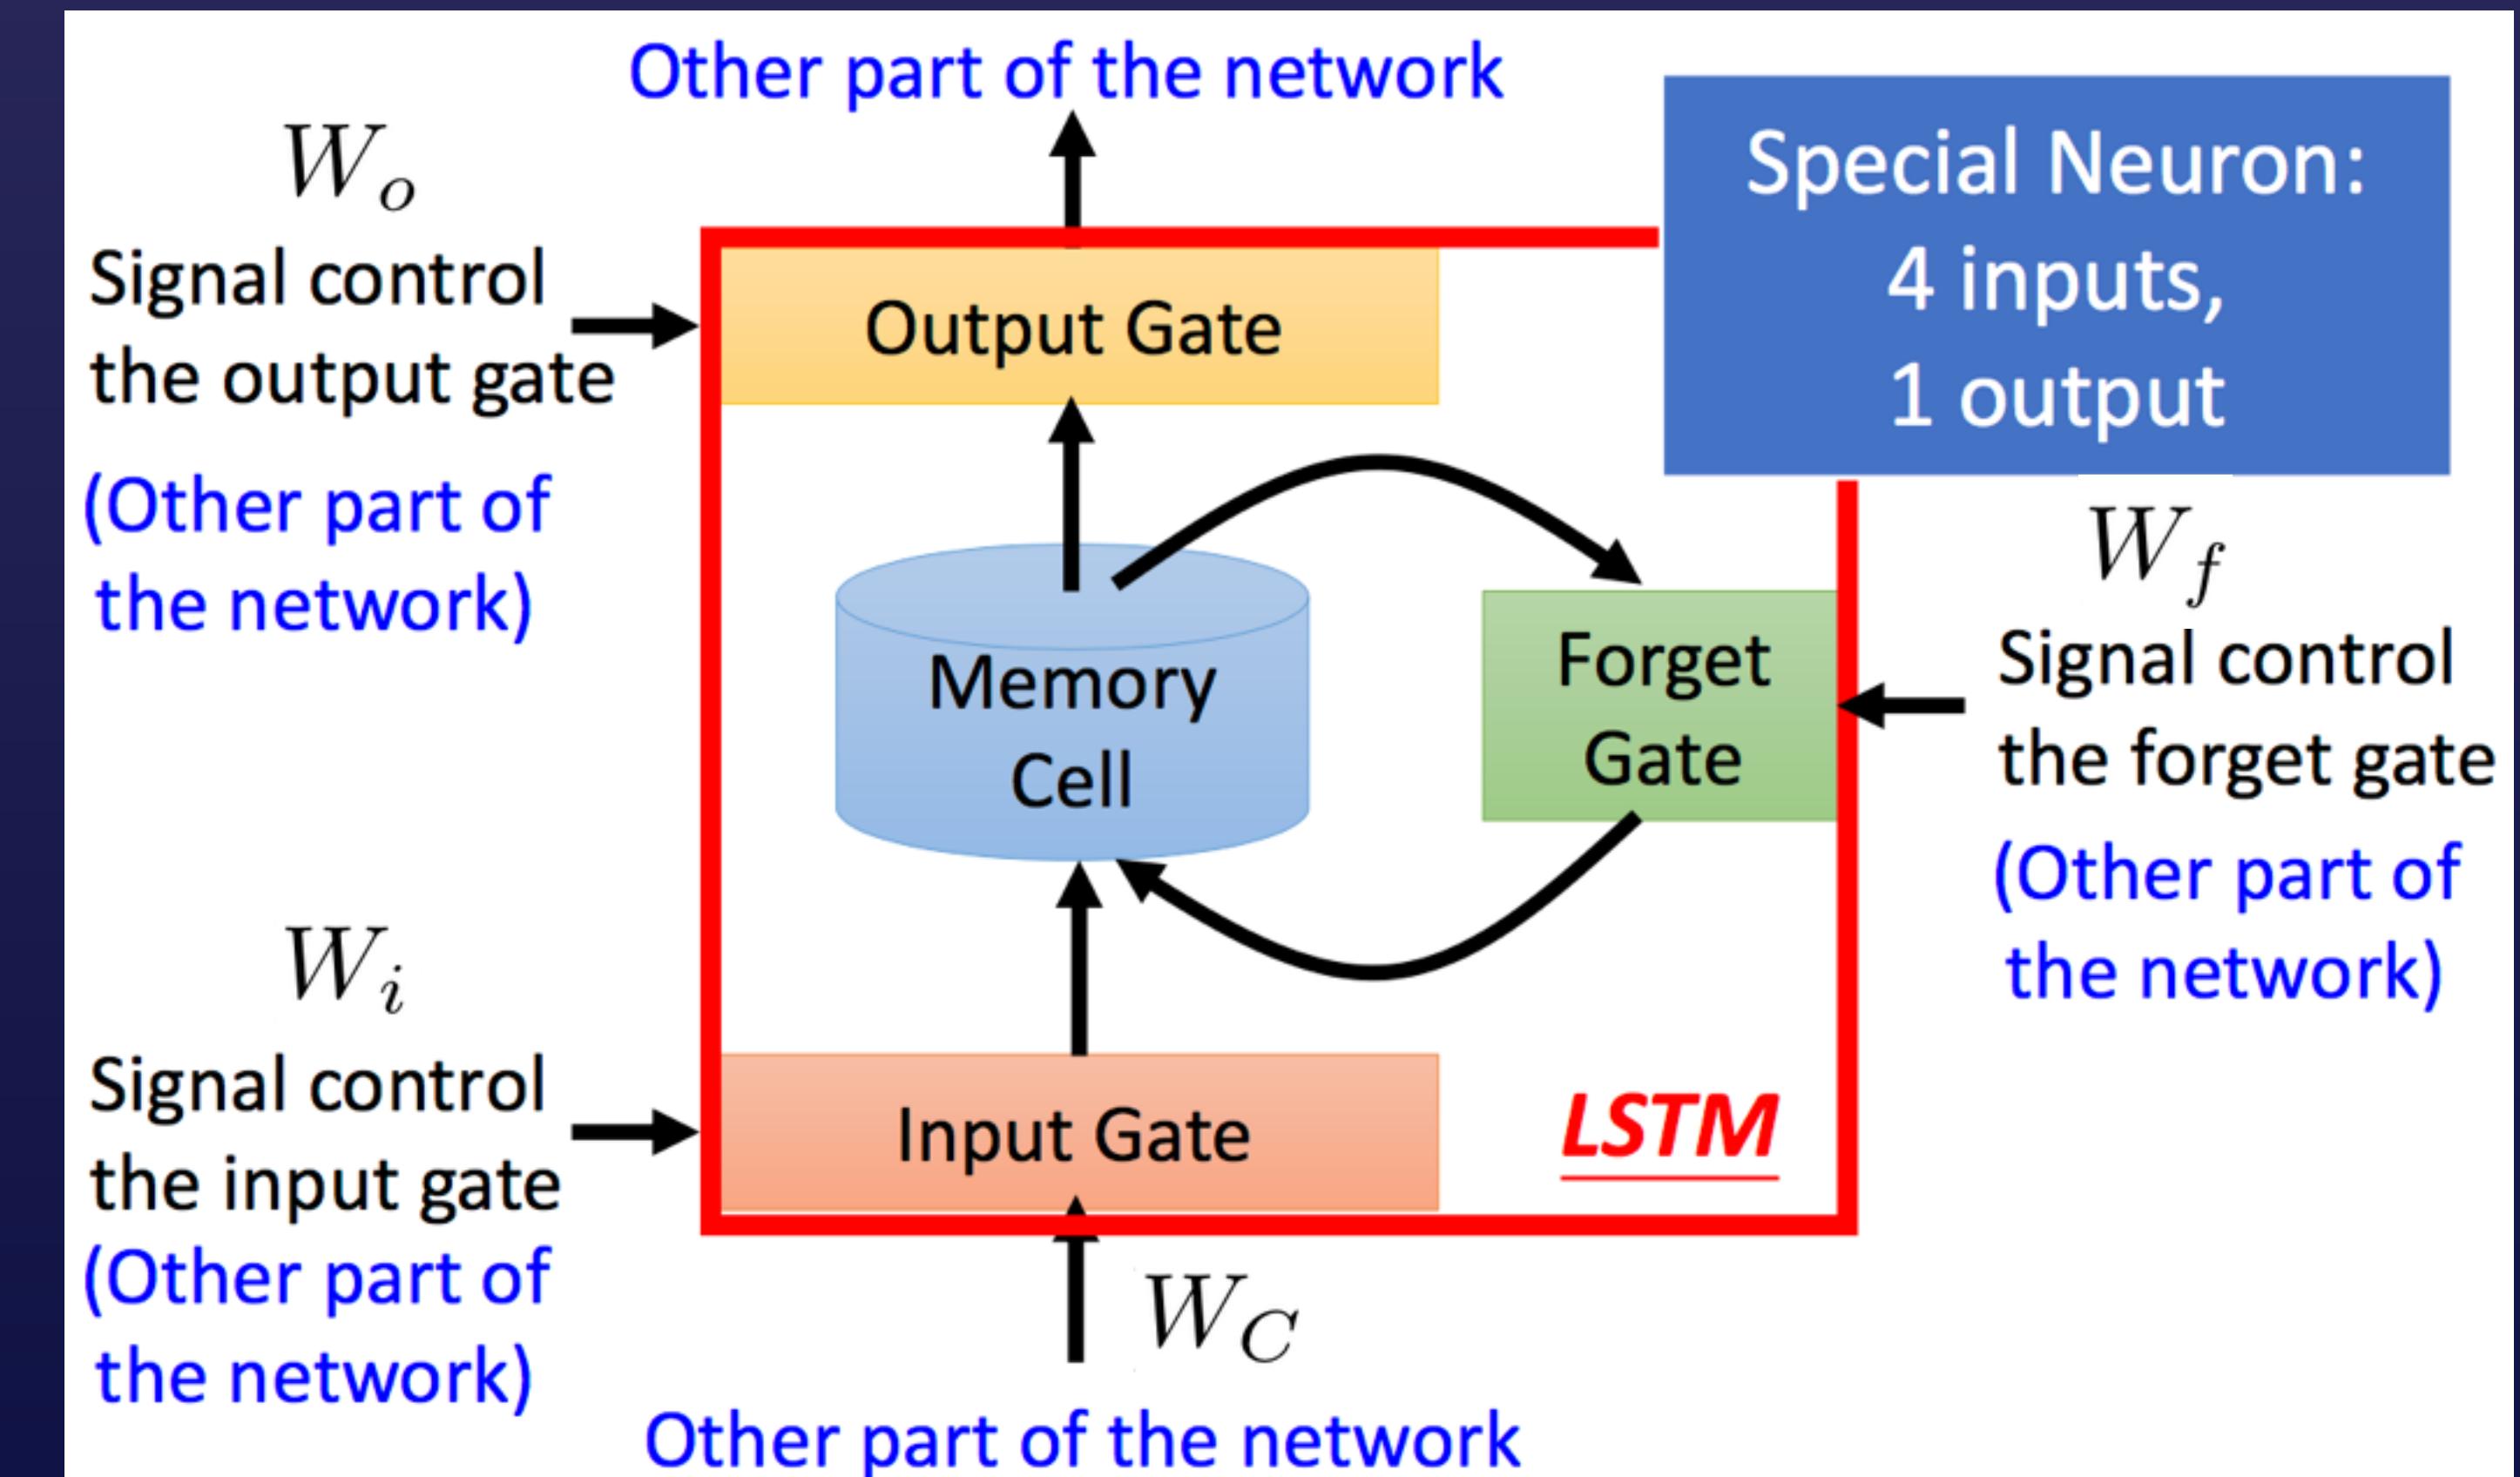
However, RNN update memory at each step  
=> Cannot “connect the dots”



# Neural Network with Memory

## Long Short-Term Memory (LSTM)

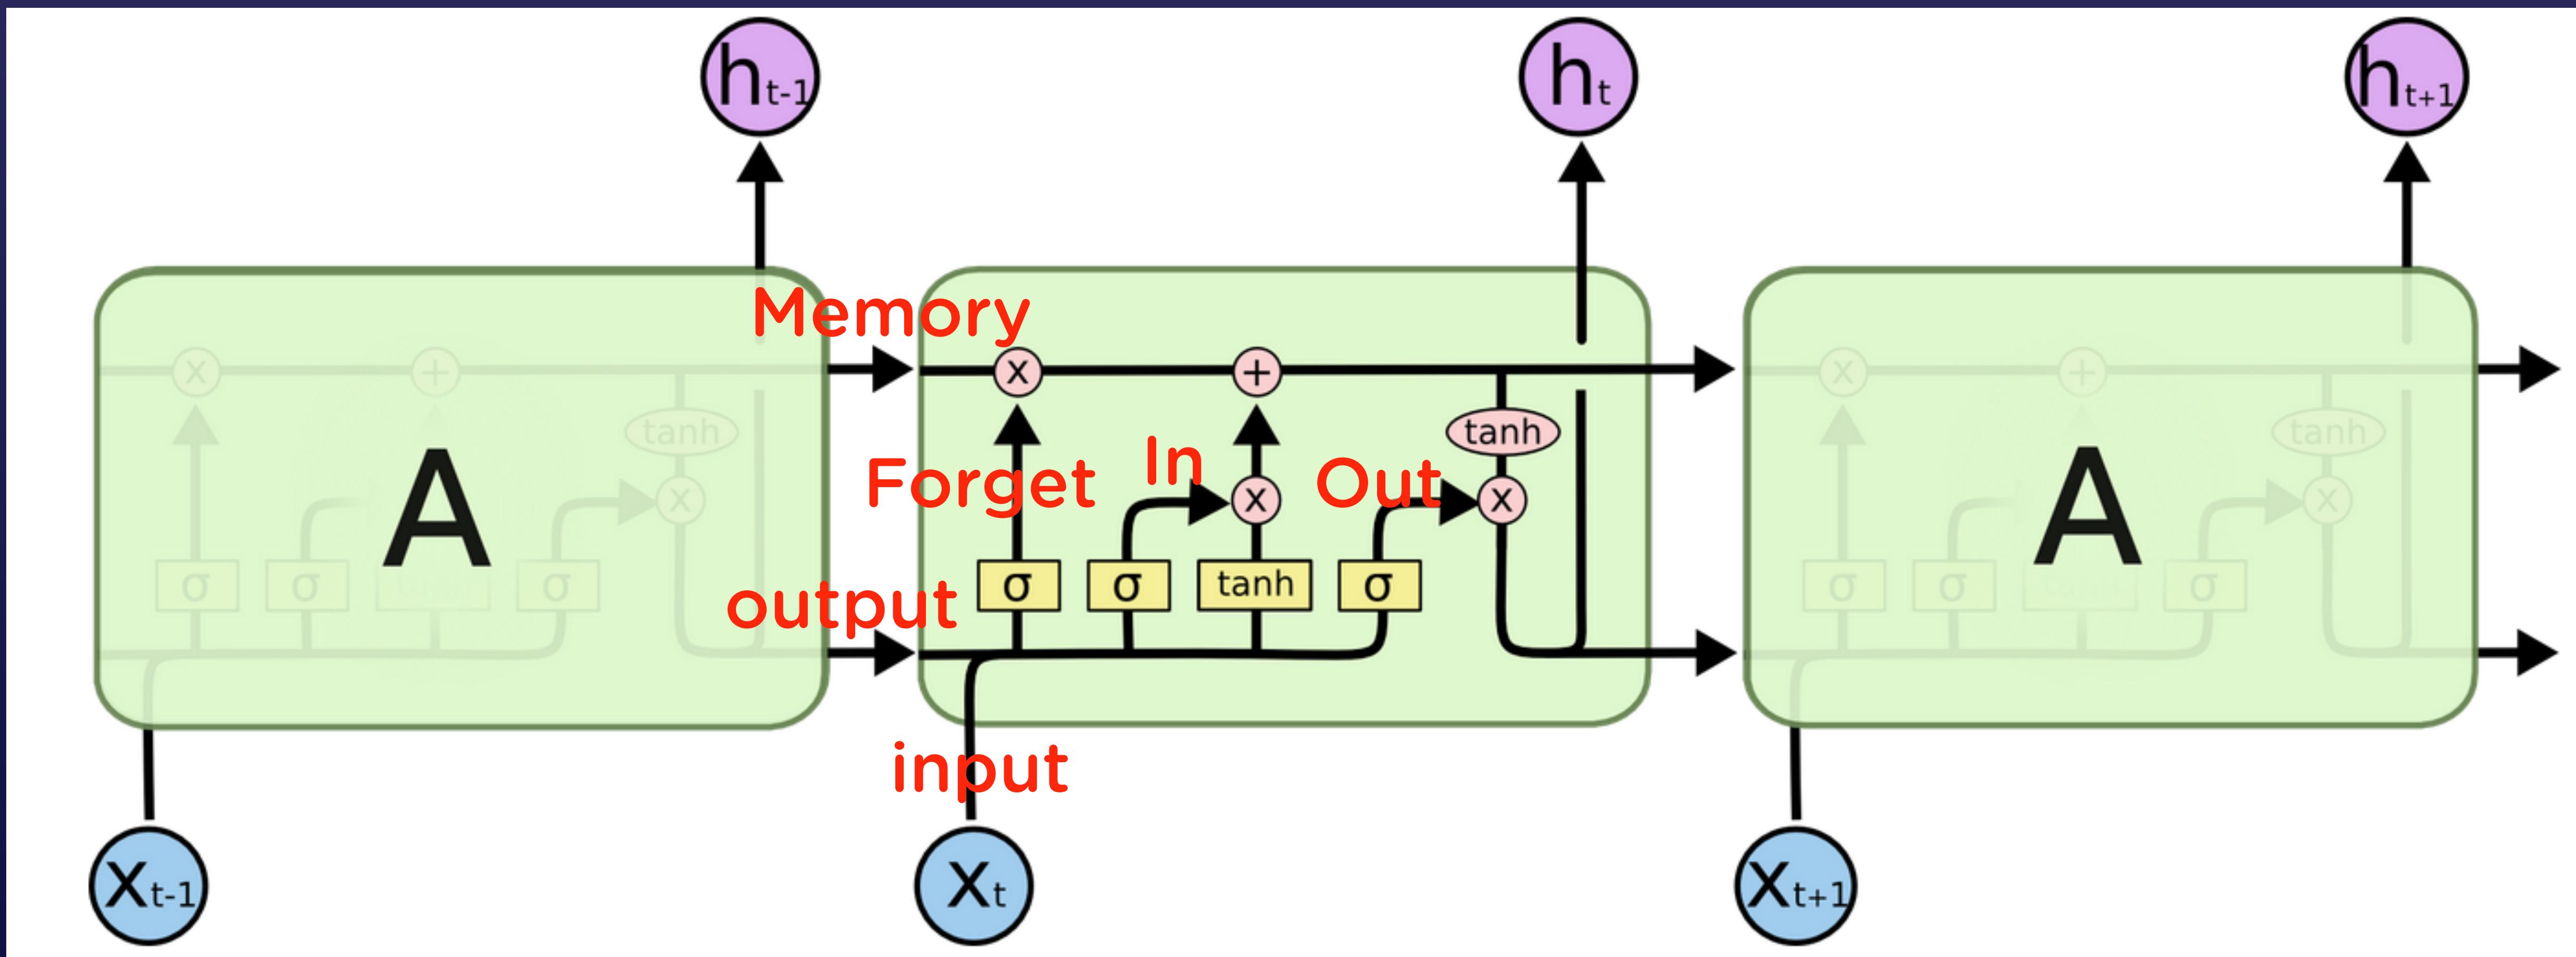
- Gate Mechanism: Control the output, update or forget of the memory
- Memory cell
- “Reuse”:  $W_o$ ,  $W_i$ ,  $W_f$  and  $W_c$  are shared
- “Sequence”: feature vectors are input “one by one”



Reference : [http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_MLSD15\\_2.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html)

# Long Short-Term Memory

## Another View

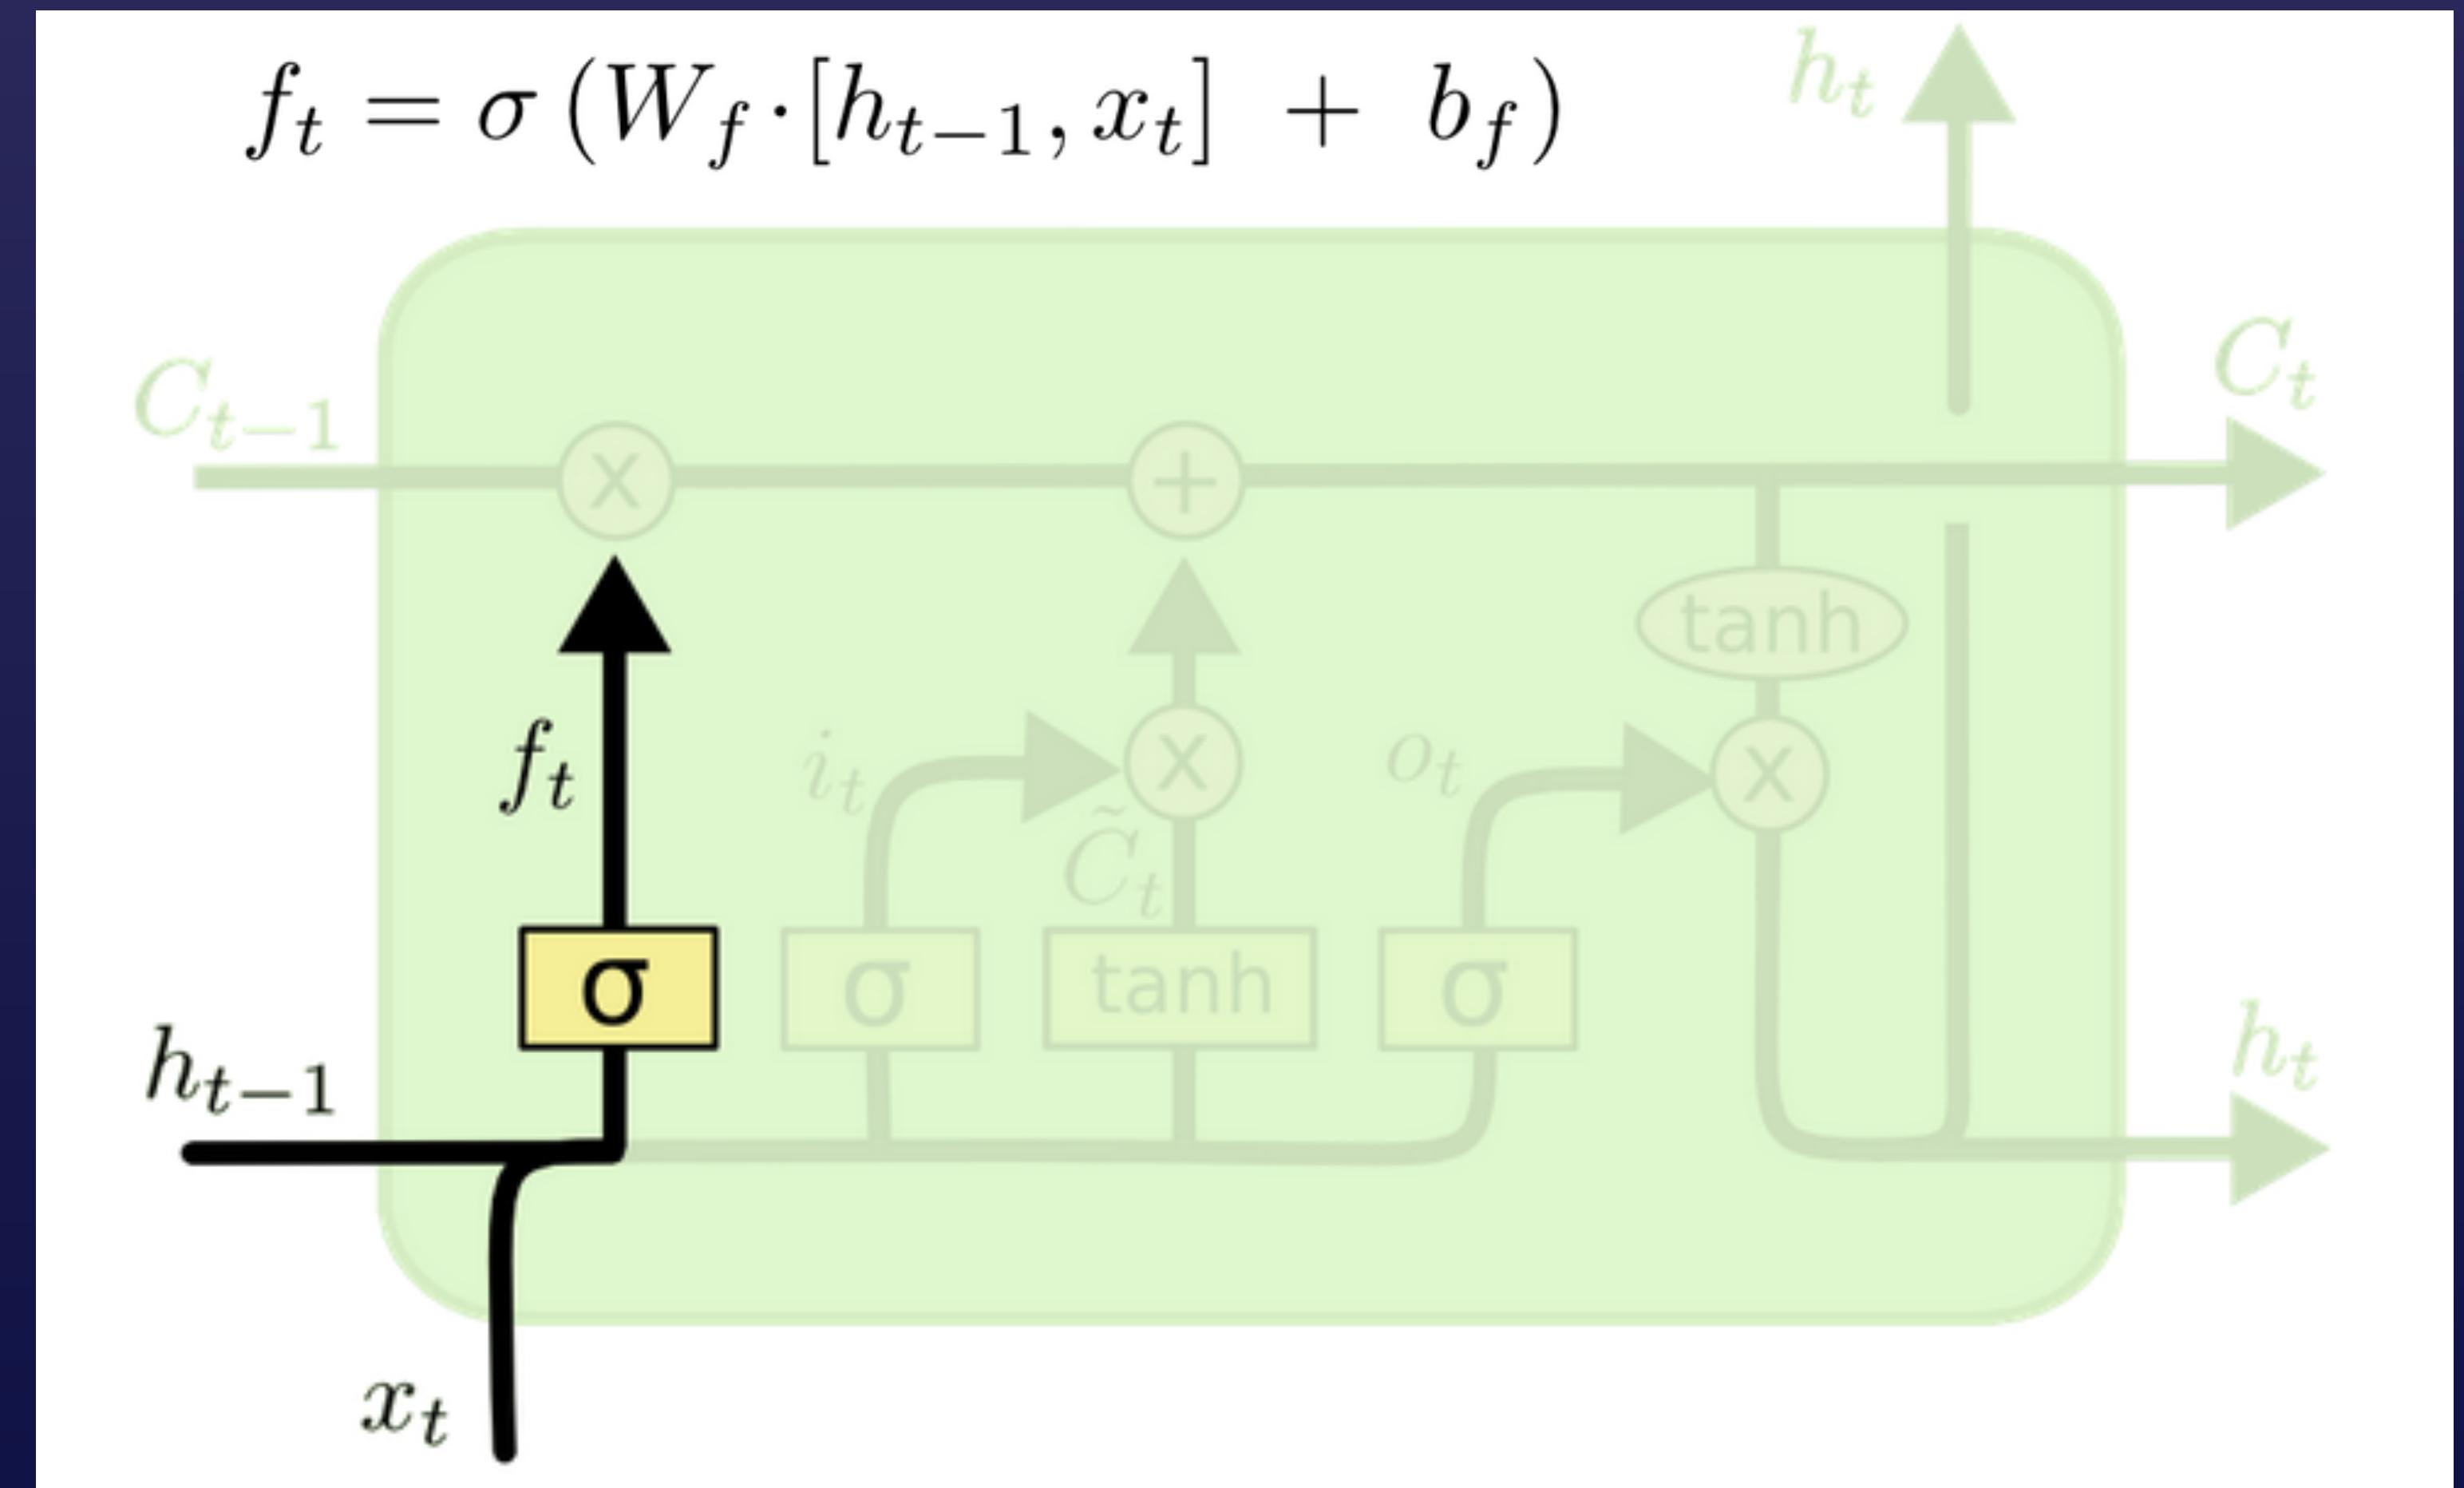


Reference : <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# Long Short-Term Memory

## Forget Step

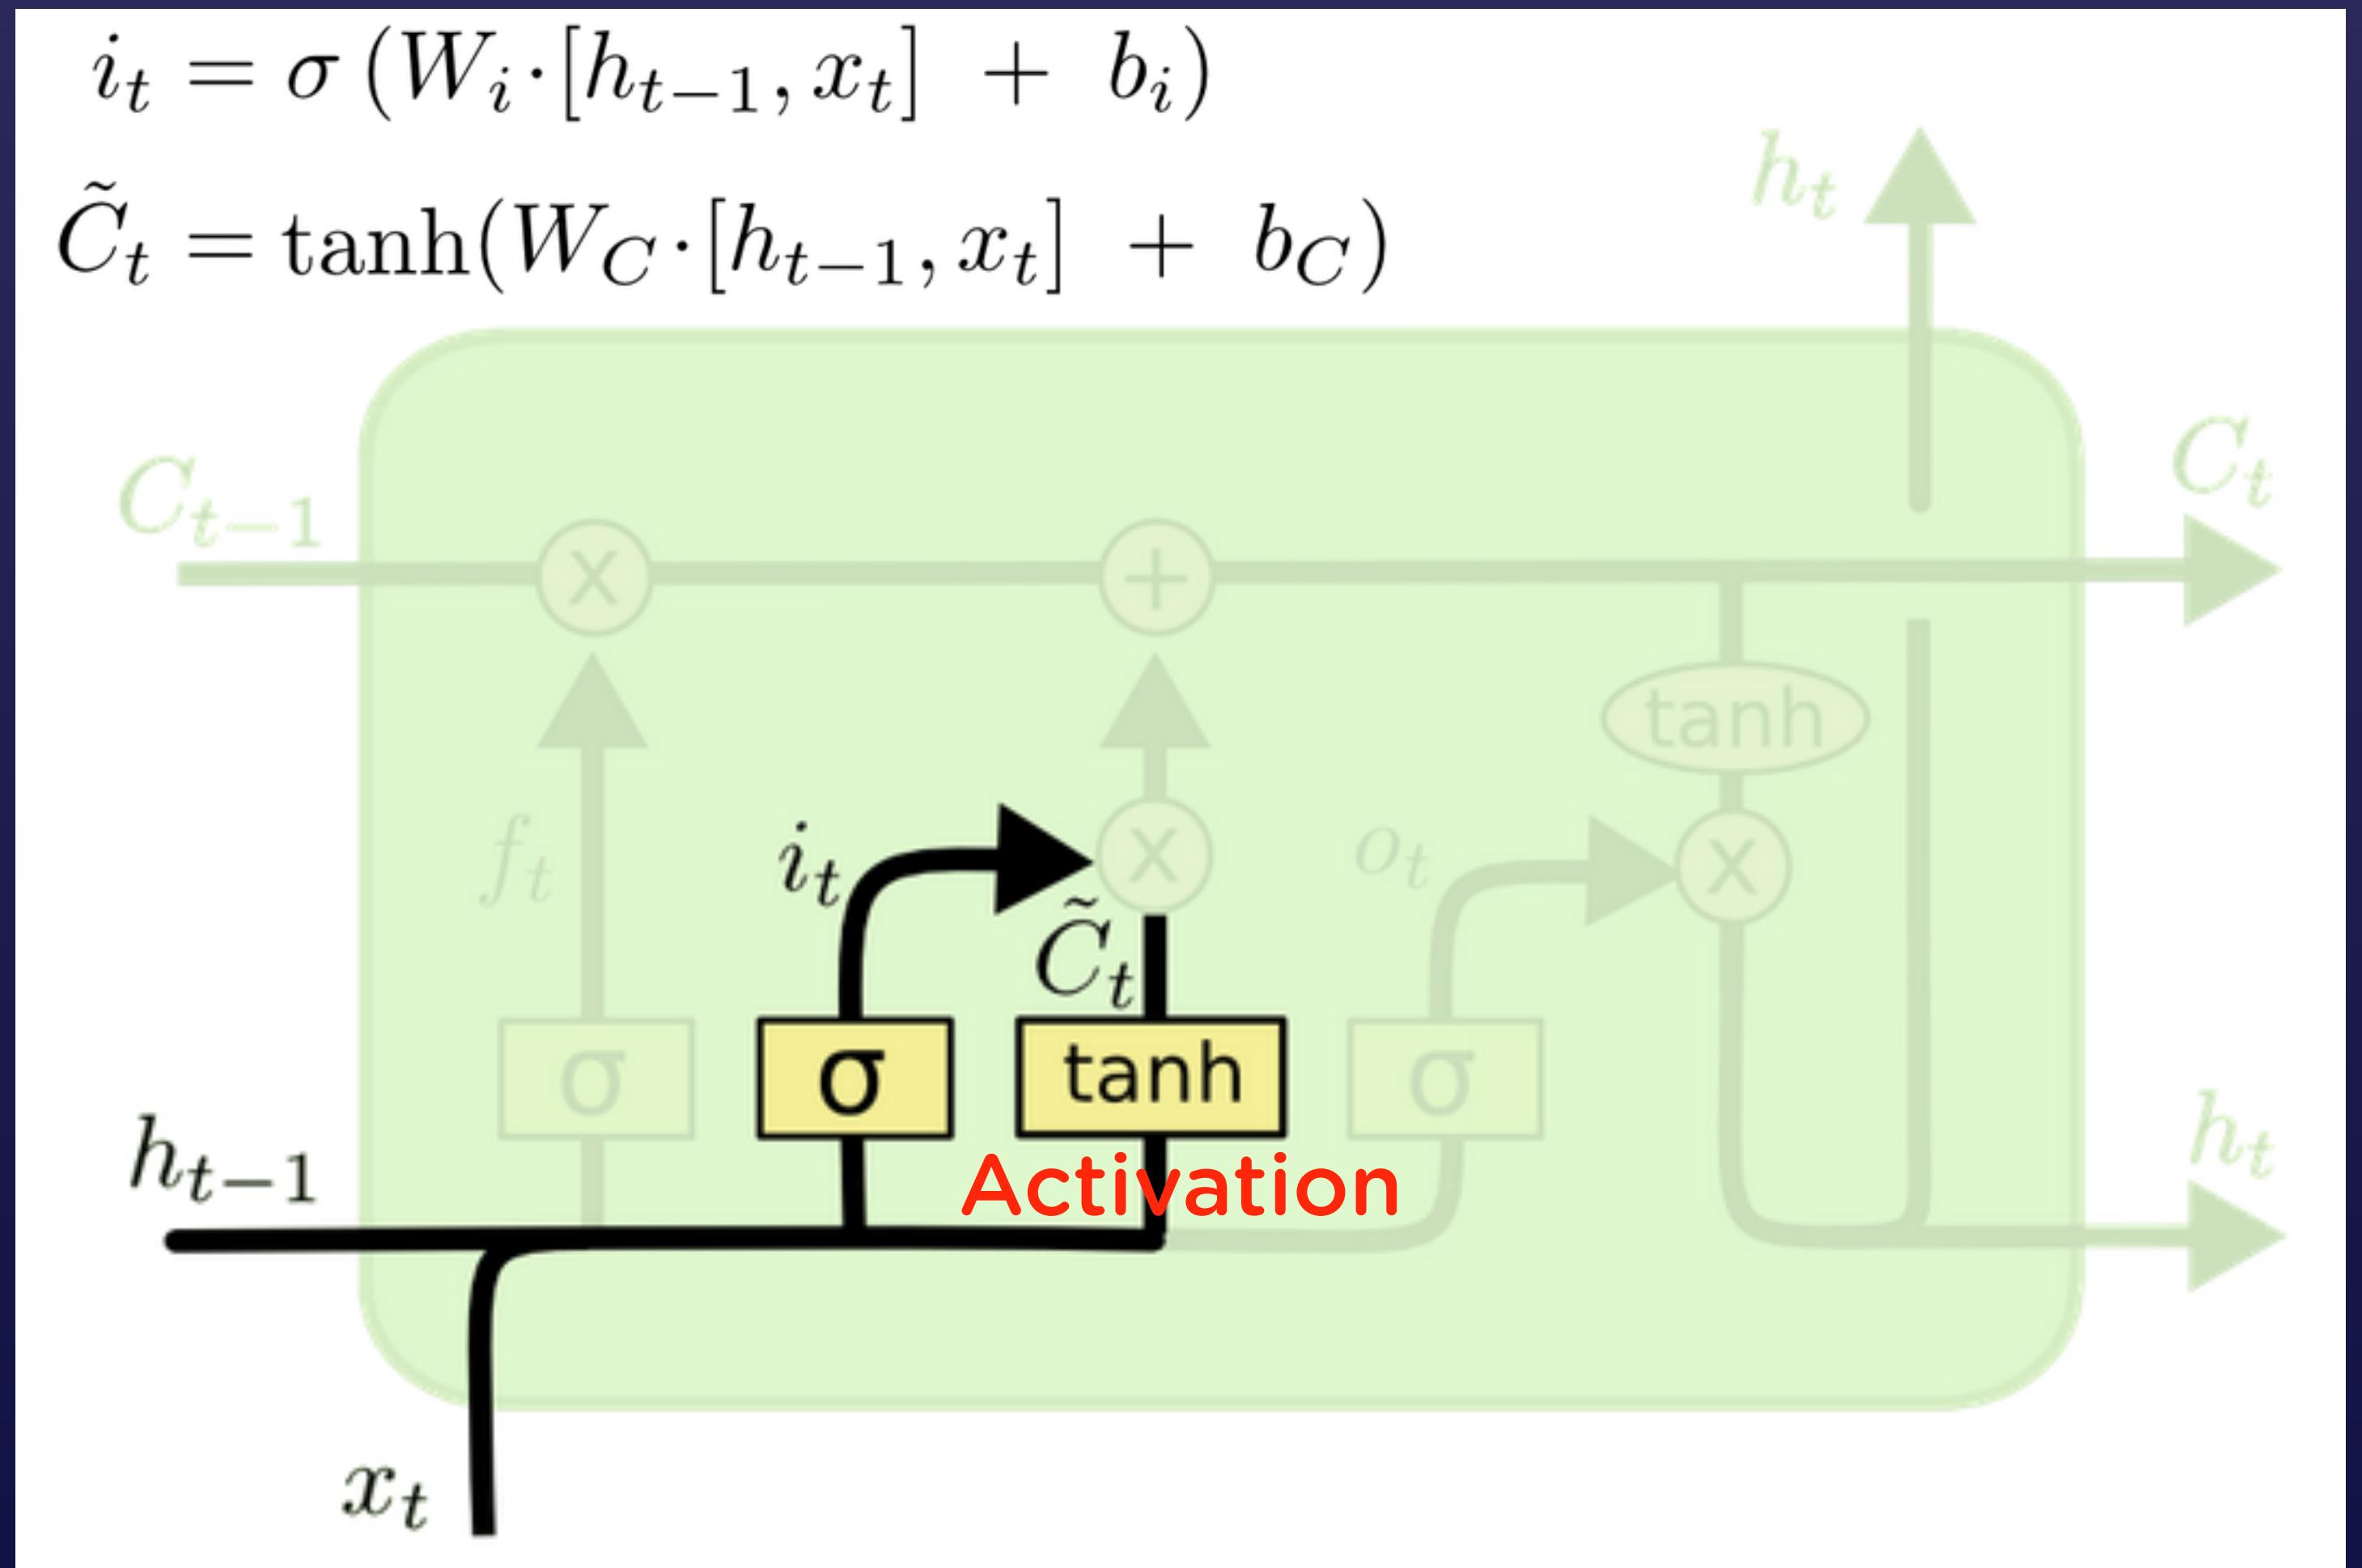
- Refer to input features of this step and the output of last step
- Decide how much memory to forget !
- $C_{t-1} * f_t$  : memory remained



# Long Short-Term Memory

## Input Step

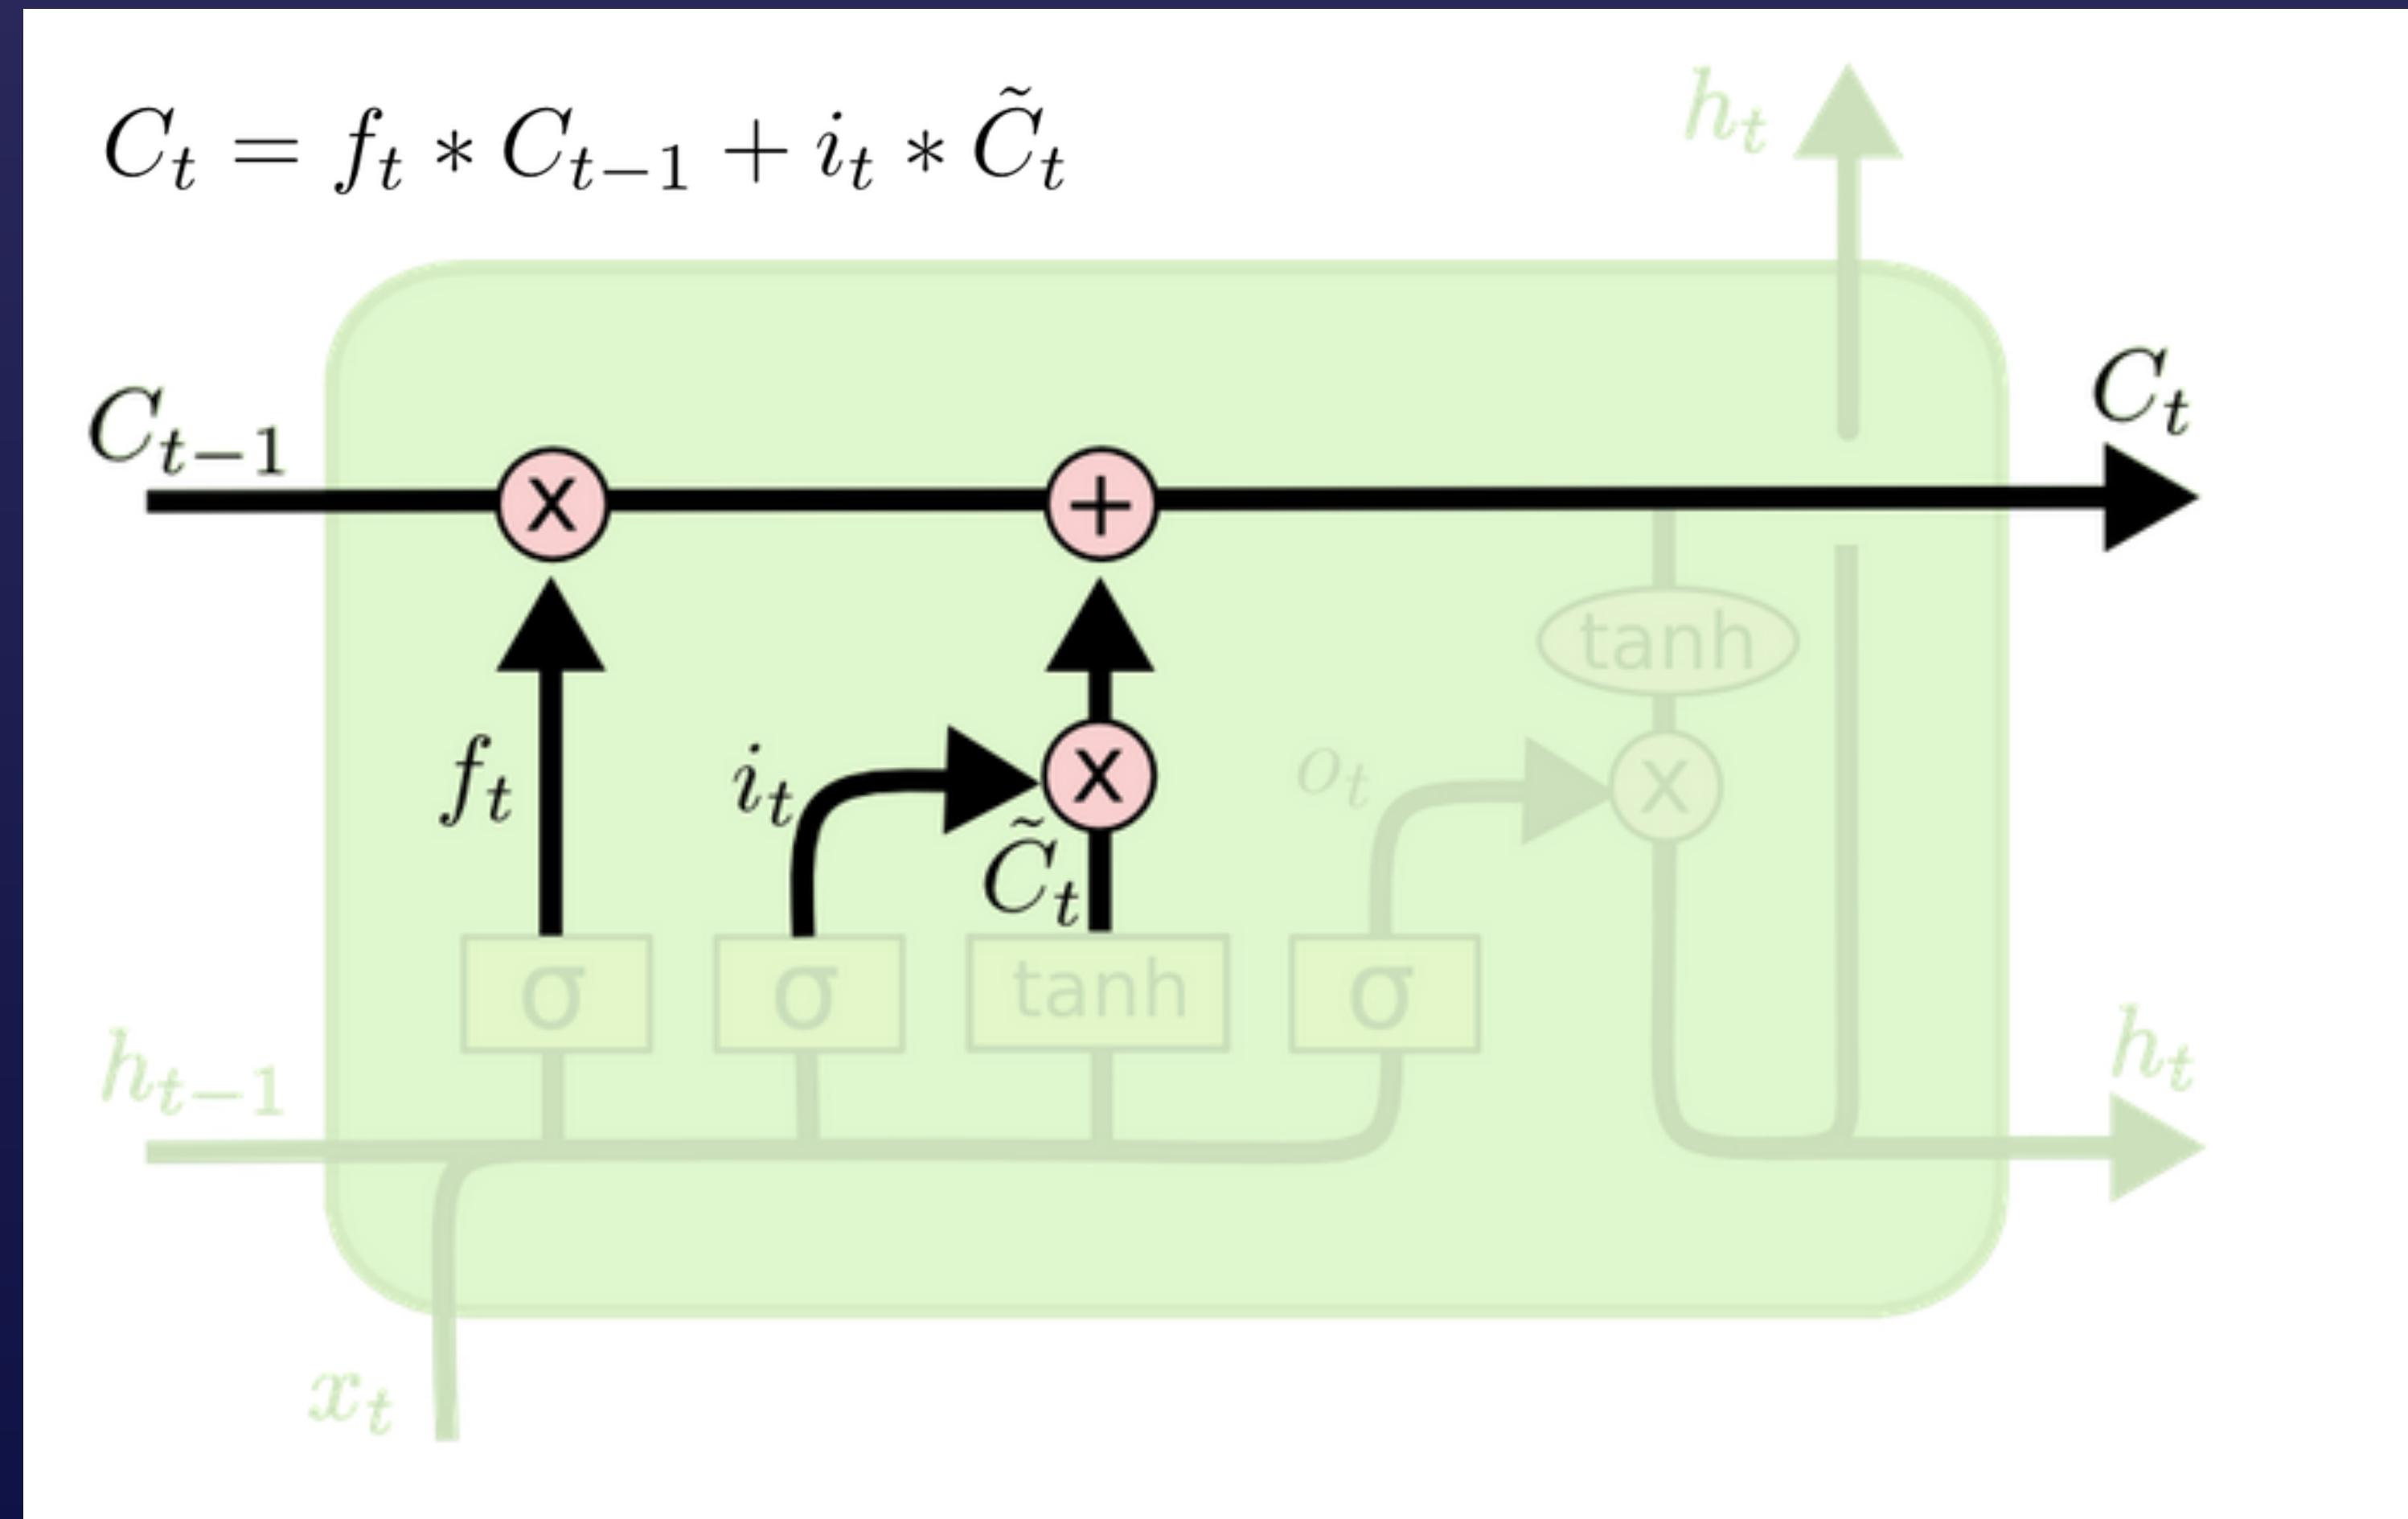
- Refer to  $x_t$  and  $h_{t-1}$
- Compute “new” info
- Decide how much new information to get in !
- $i_t * C'_t$  : new info to add to memory



# Long Short-Term Memory

## Update Step

- Update the memory
- Survived “old” memory plus passed “new” info



Reference : <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

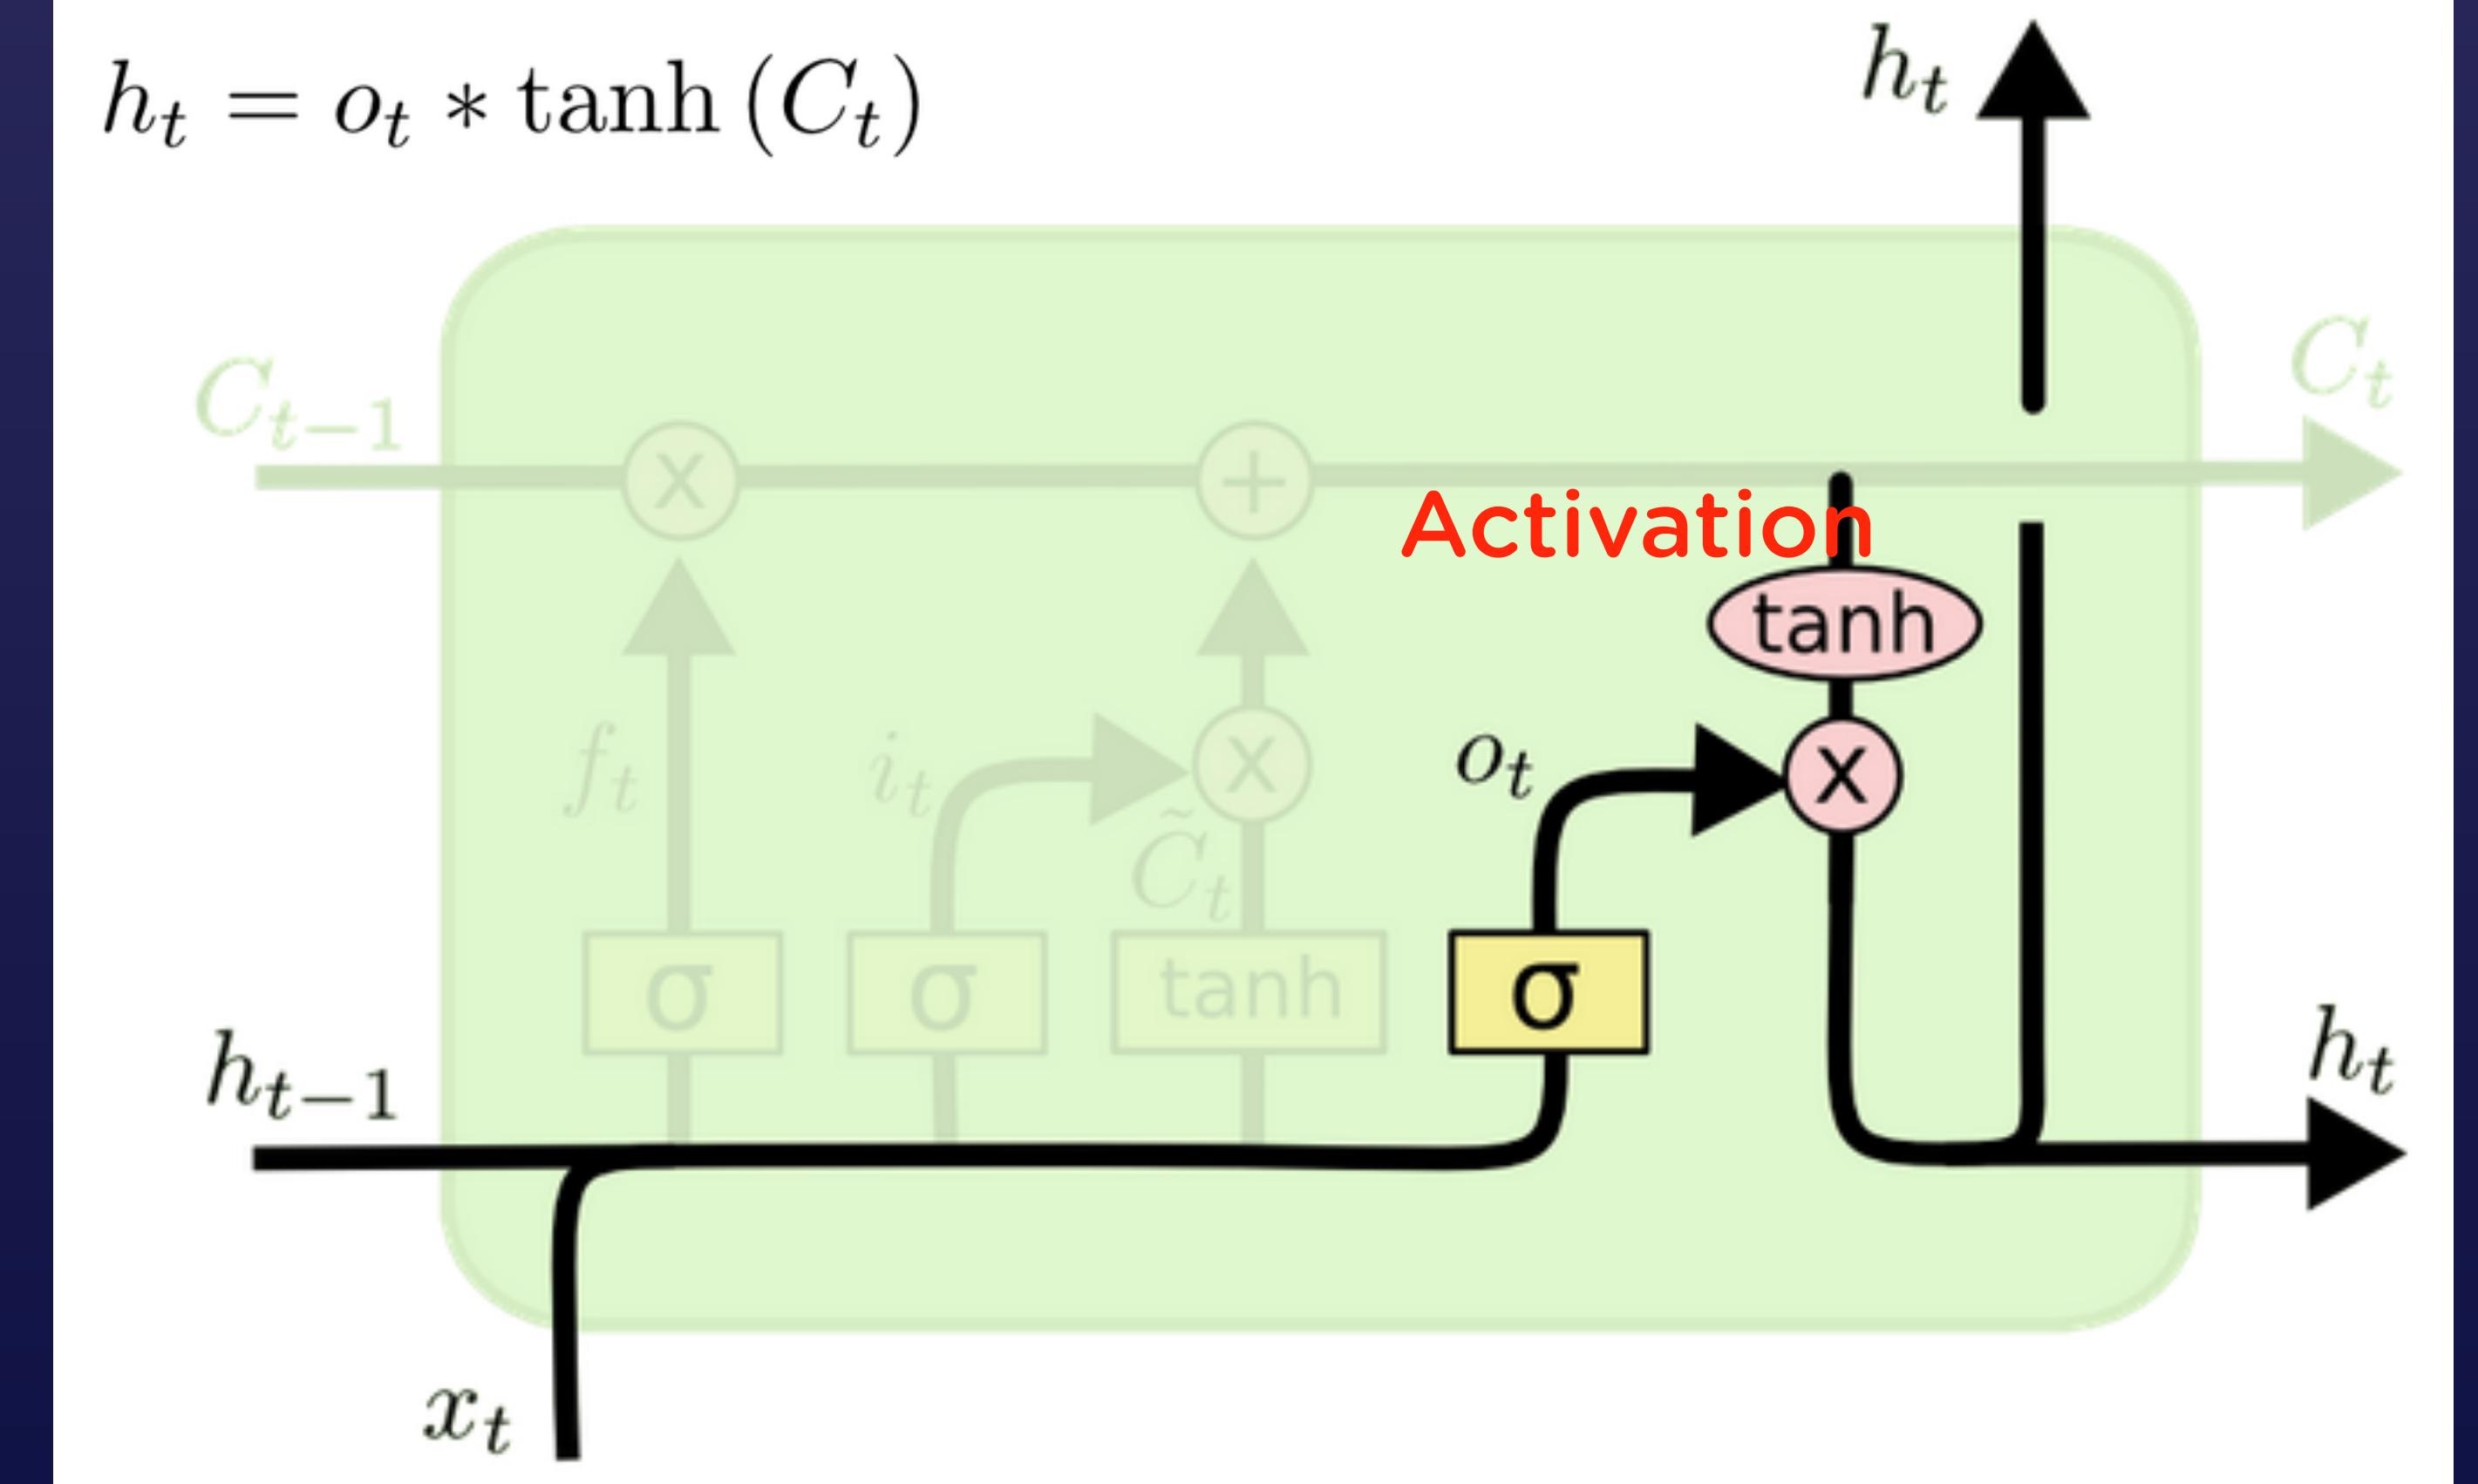
# Long Short-Term Memory

## Output Step

- Refer to  $x_t$  and  $h_{t-1}$
- memory + activation:  
ready to get out
- Decide how much memory  
info could get out !

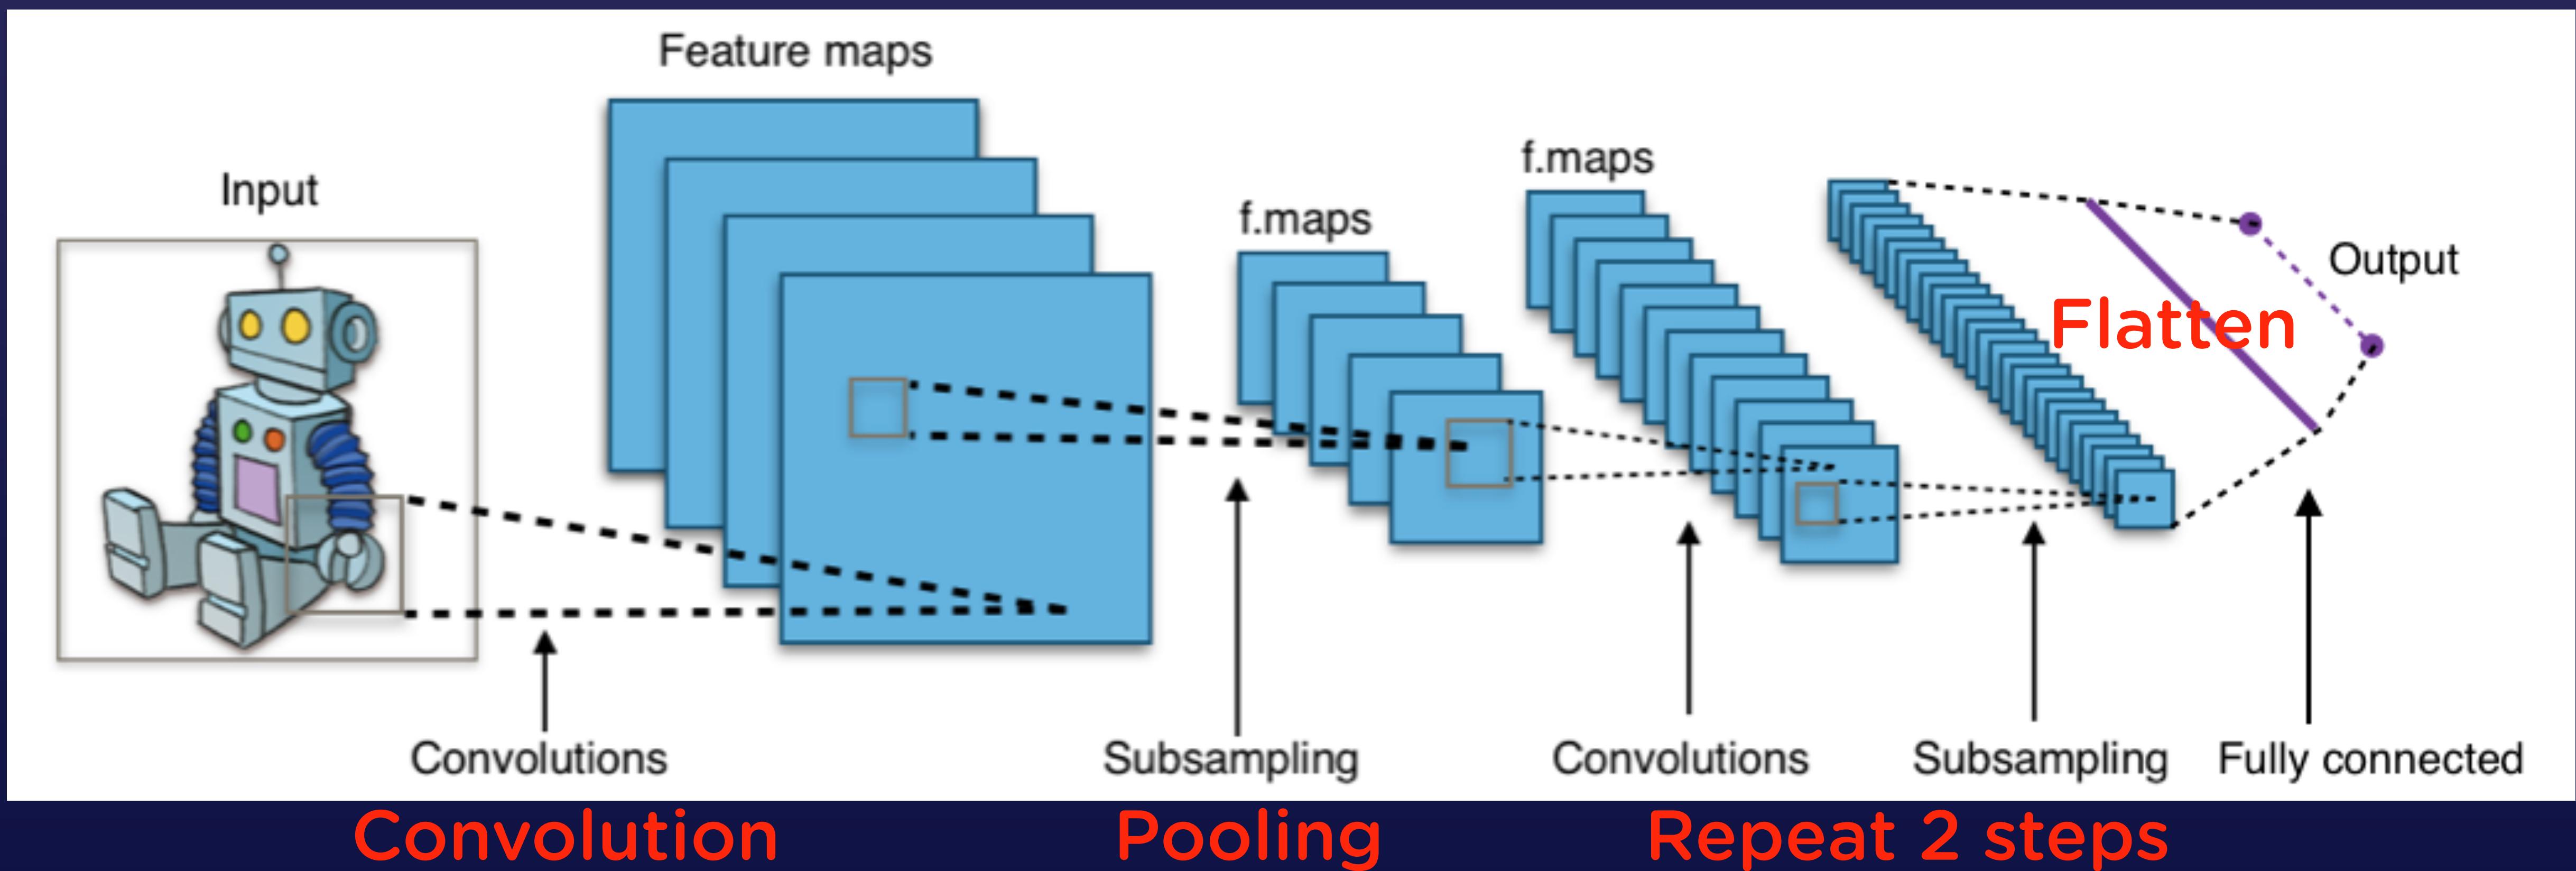
$$o_t = \sigma (W_o [ h_{t-1}, x_t ] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$



# Convolutional Neural Network

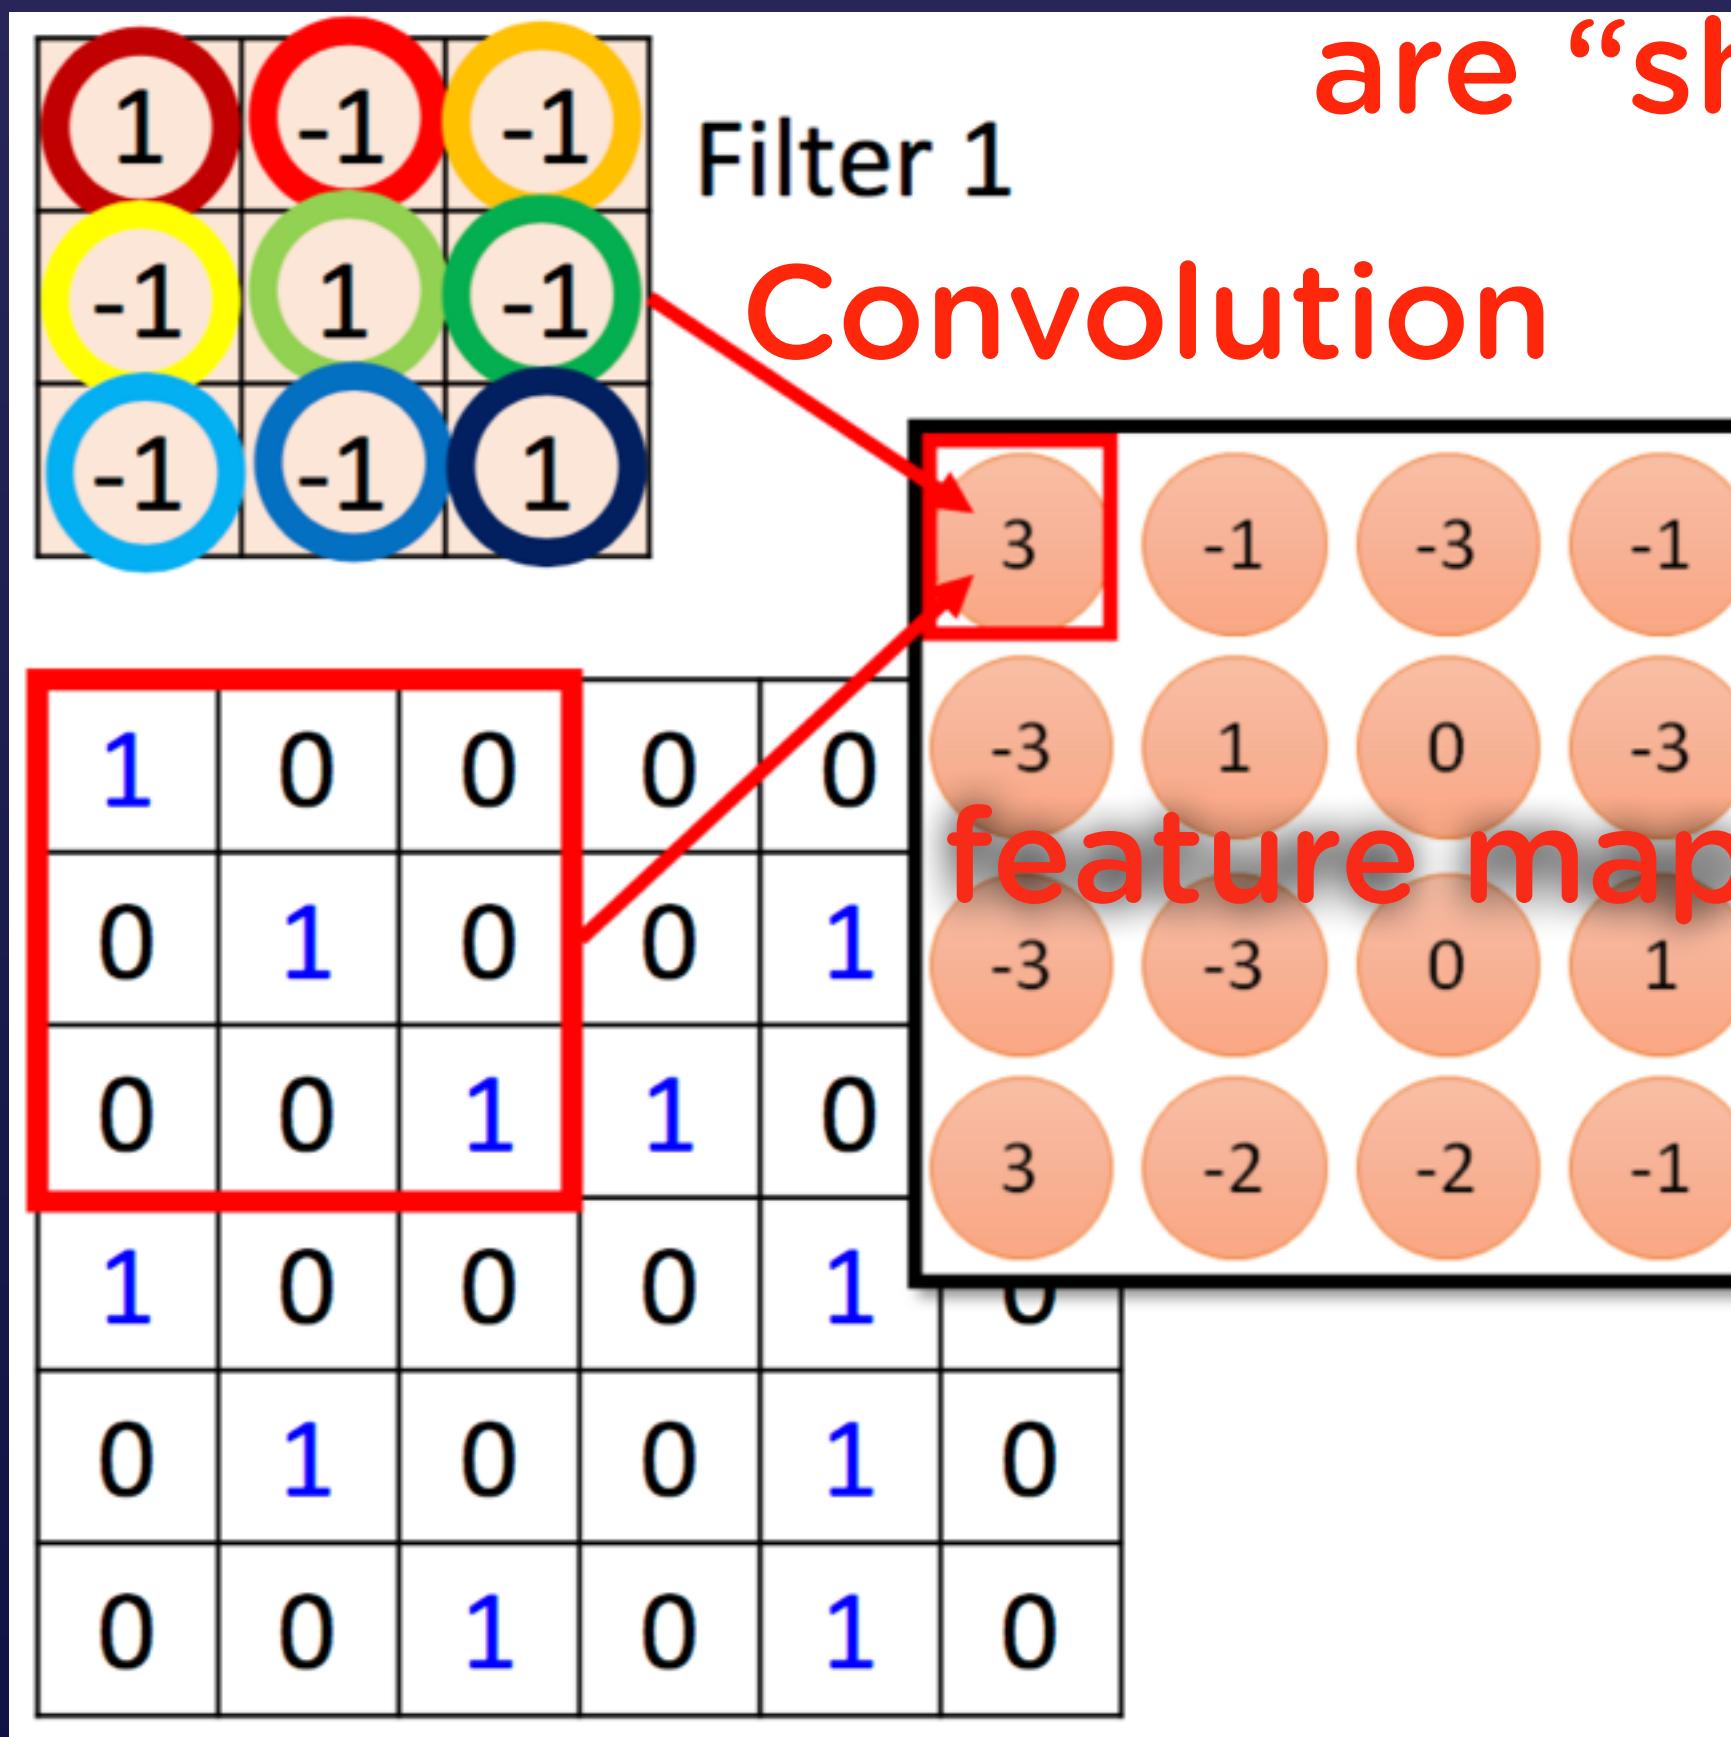
Images : local properties!



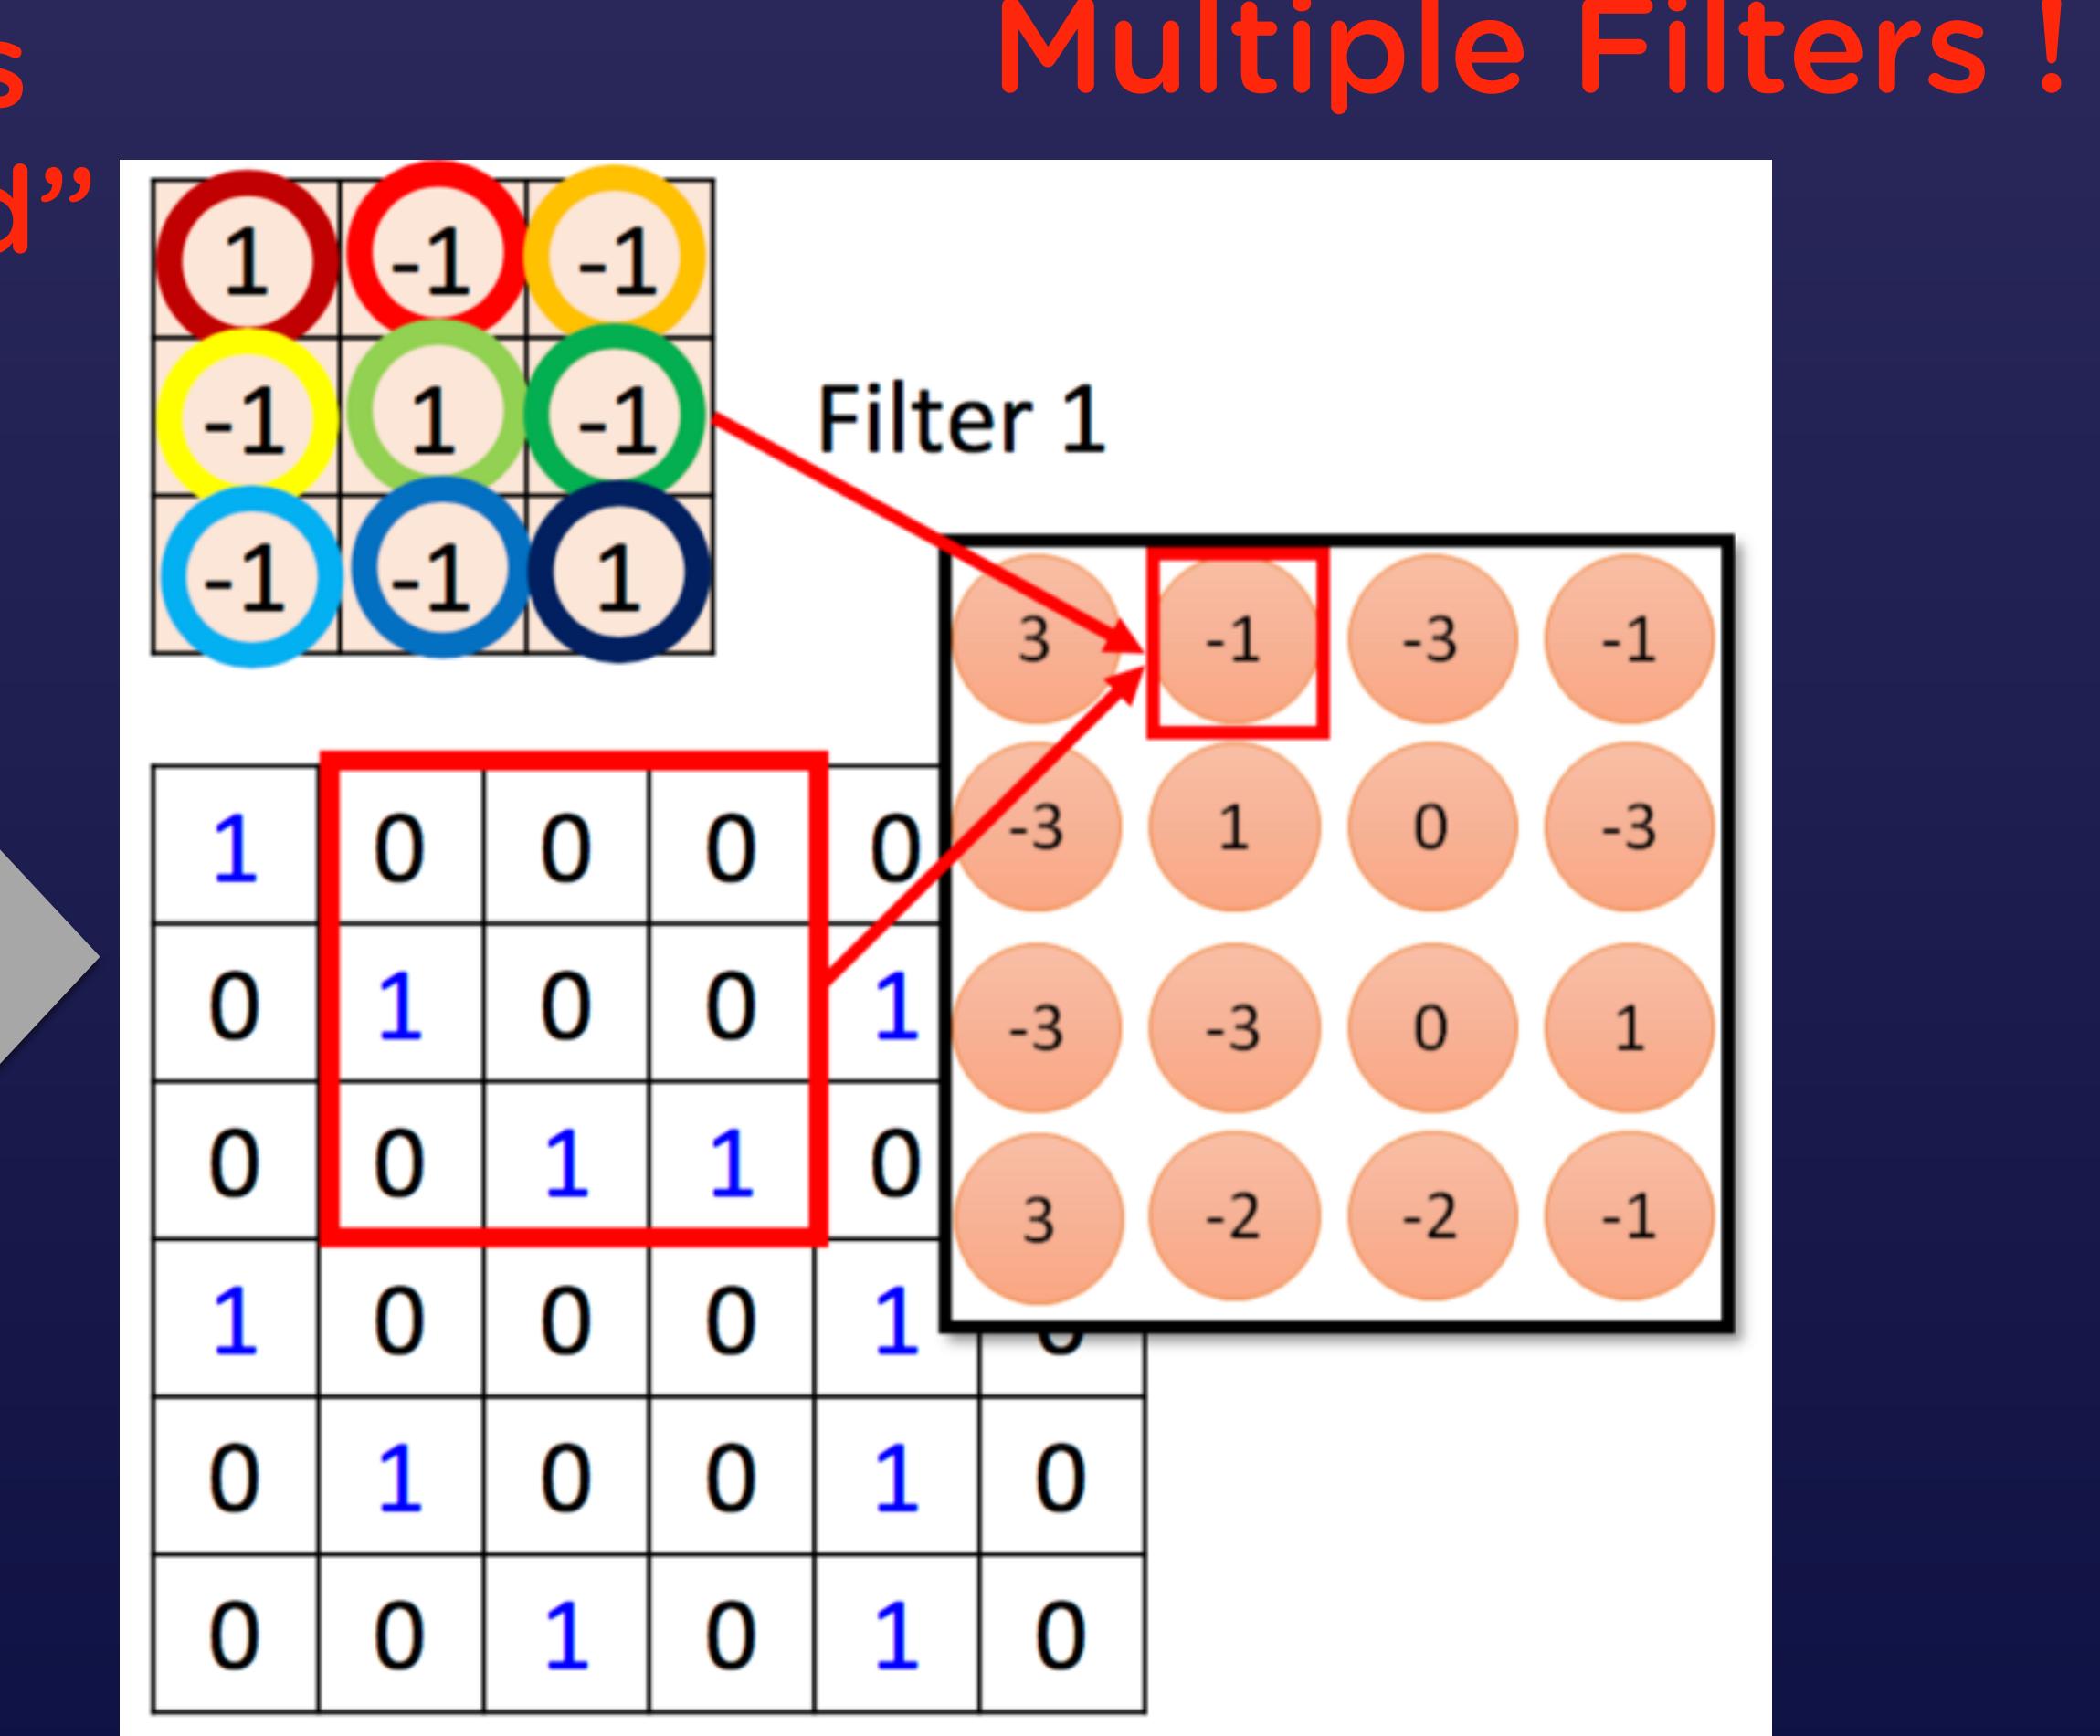
Reference : Wikipedia

# Convolutional Neural Network

## Convolution

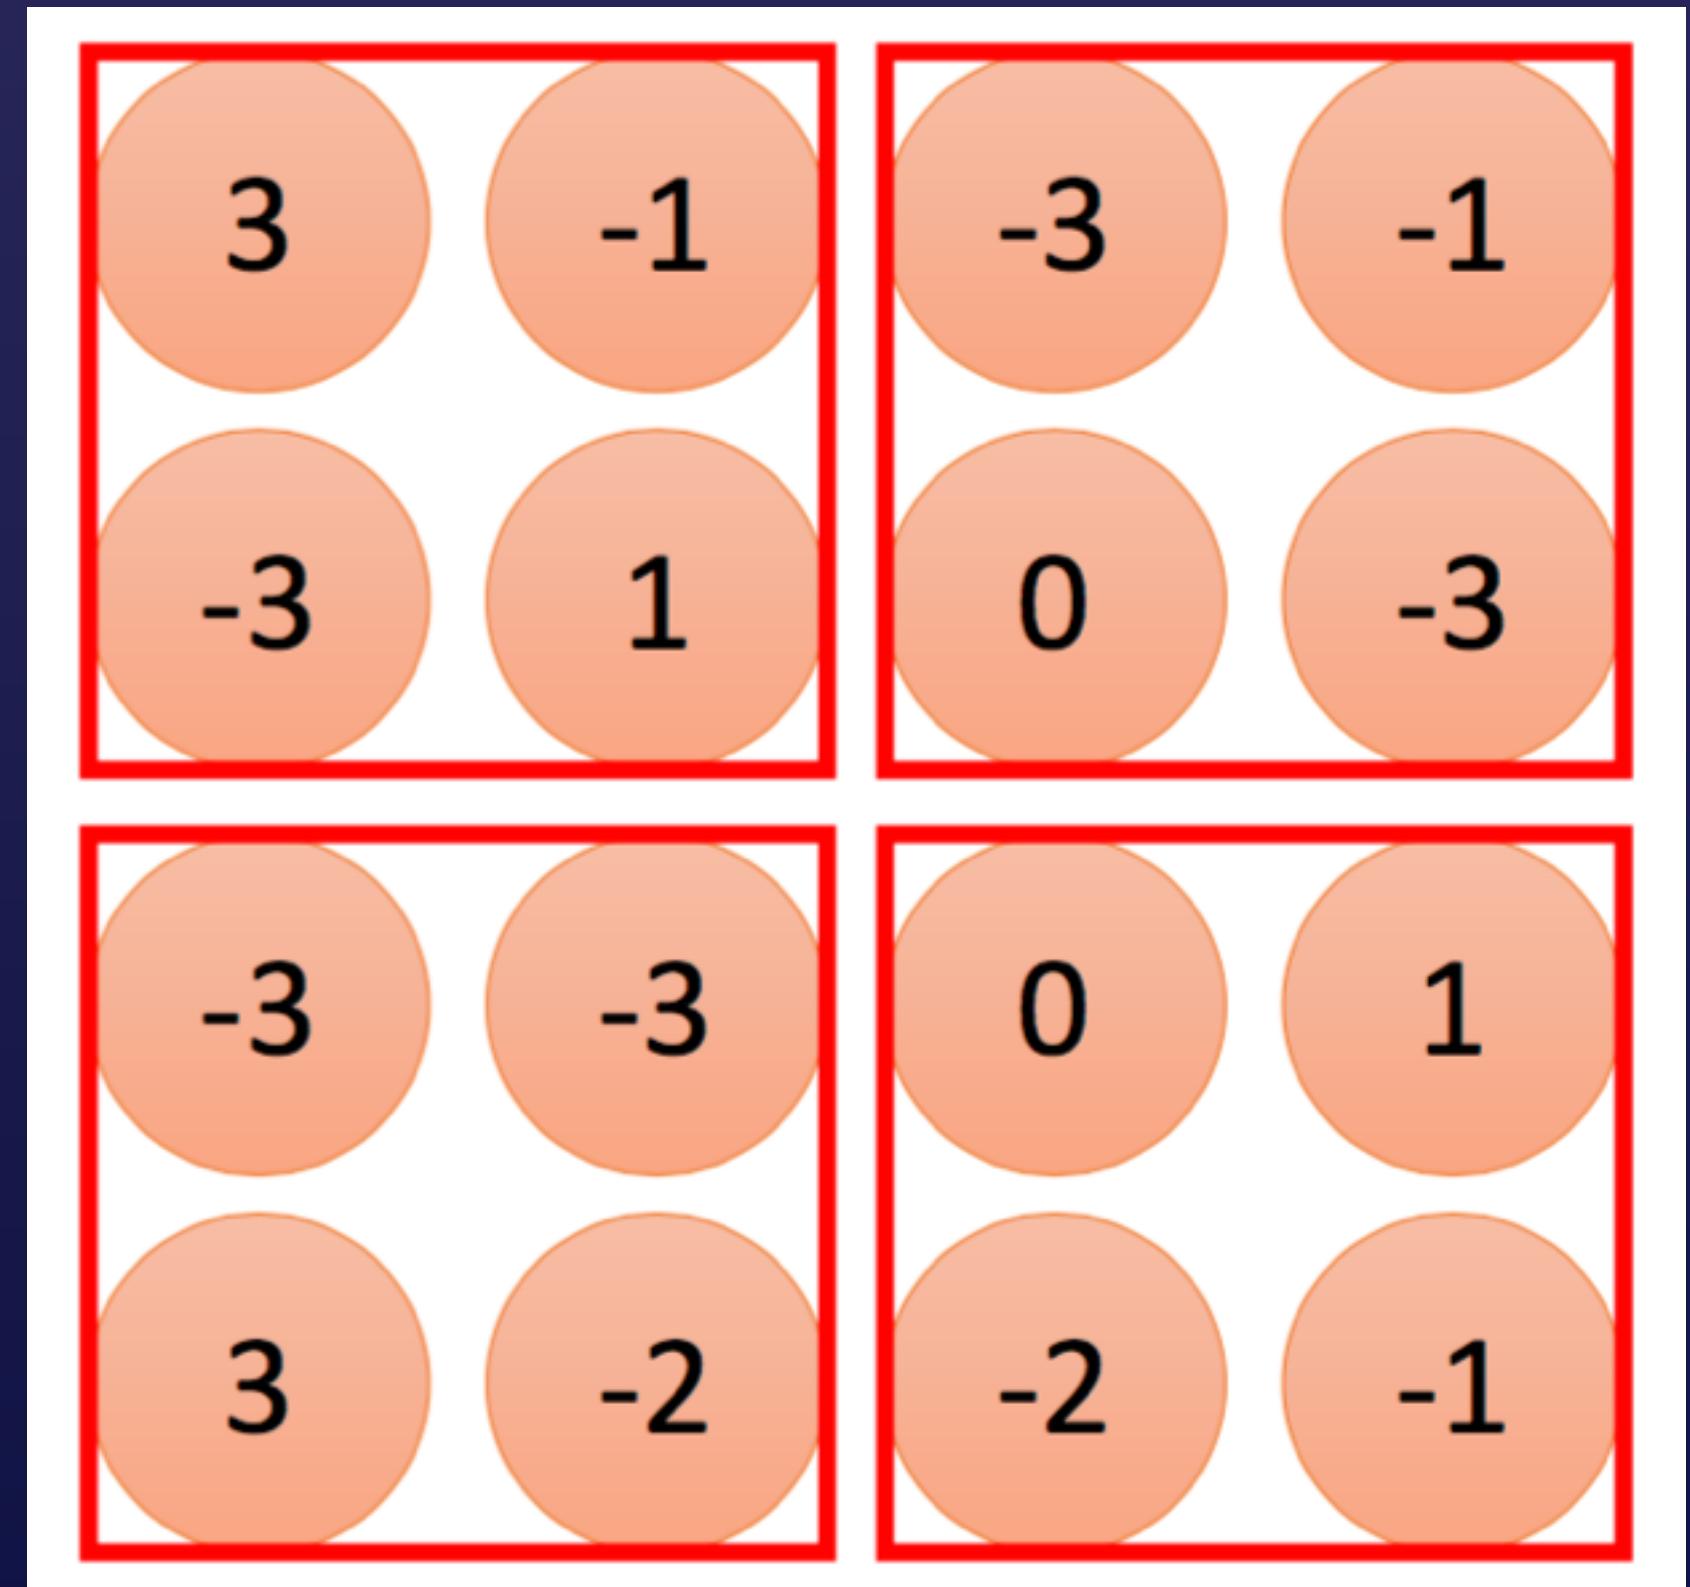
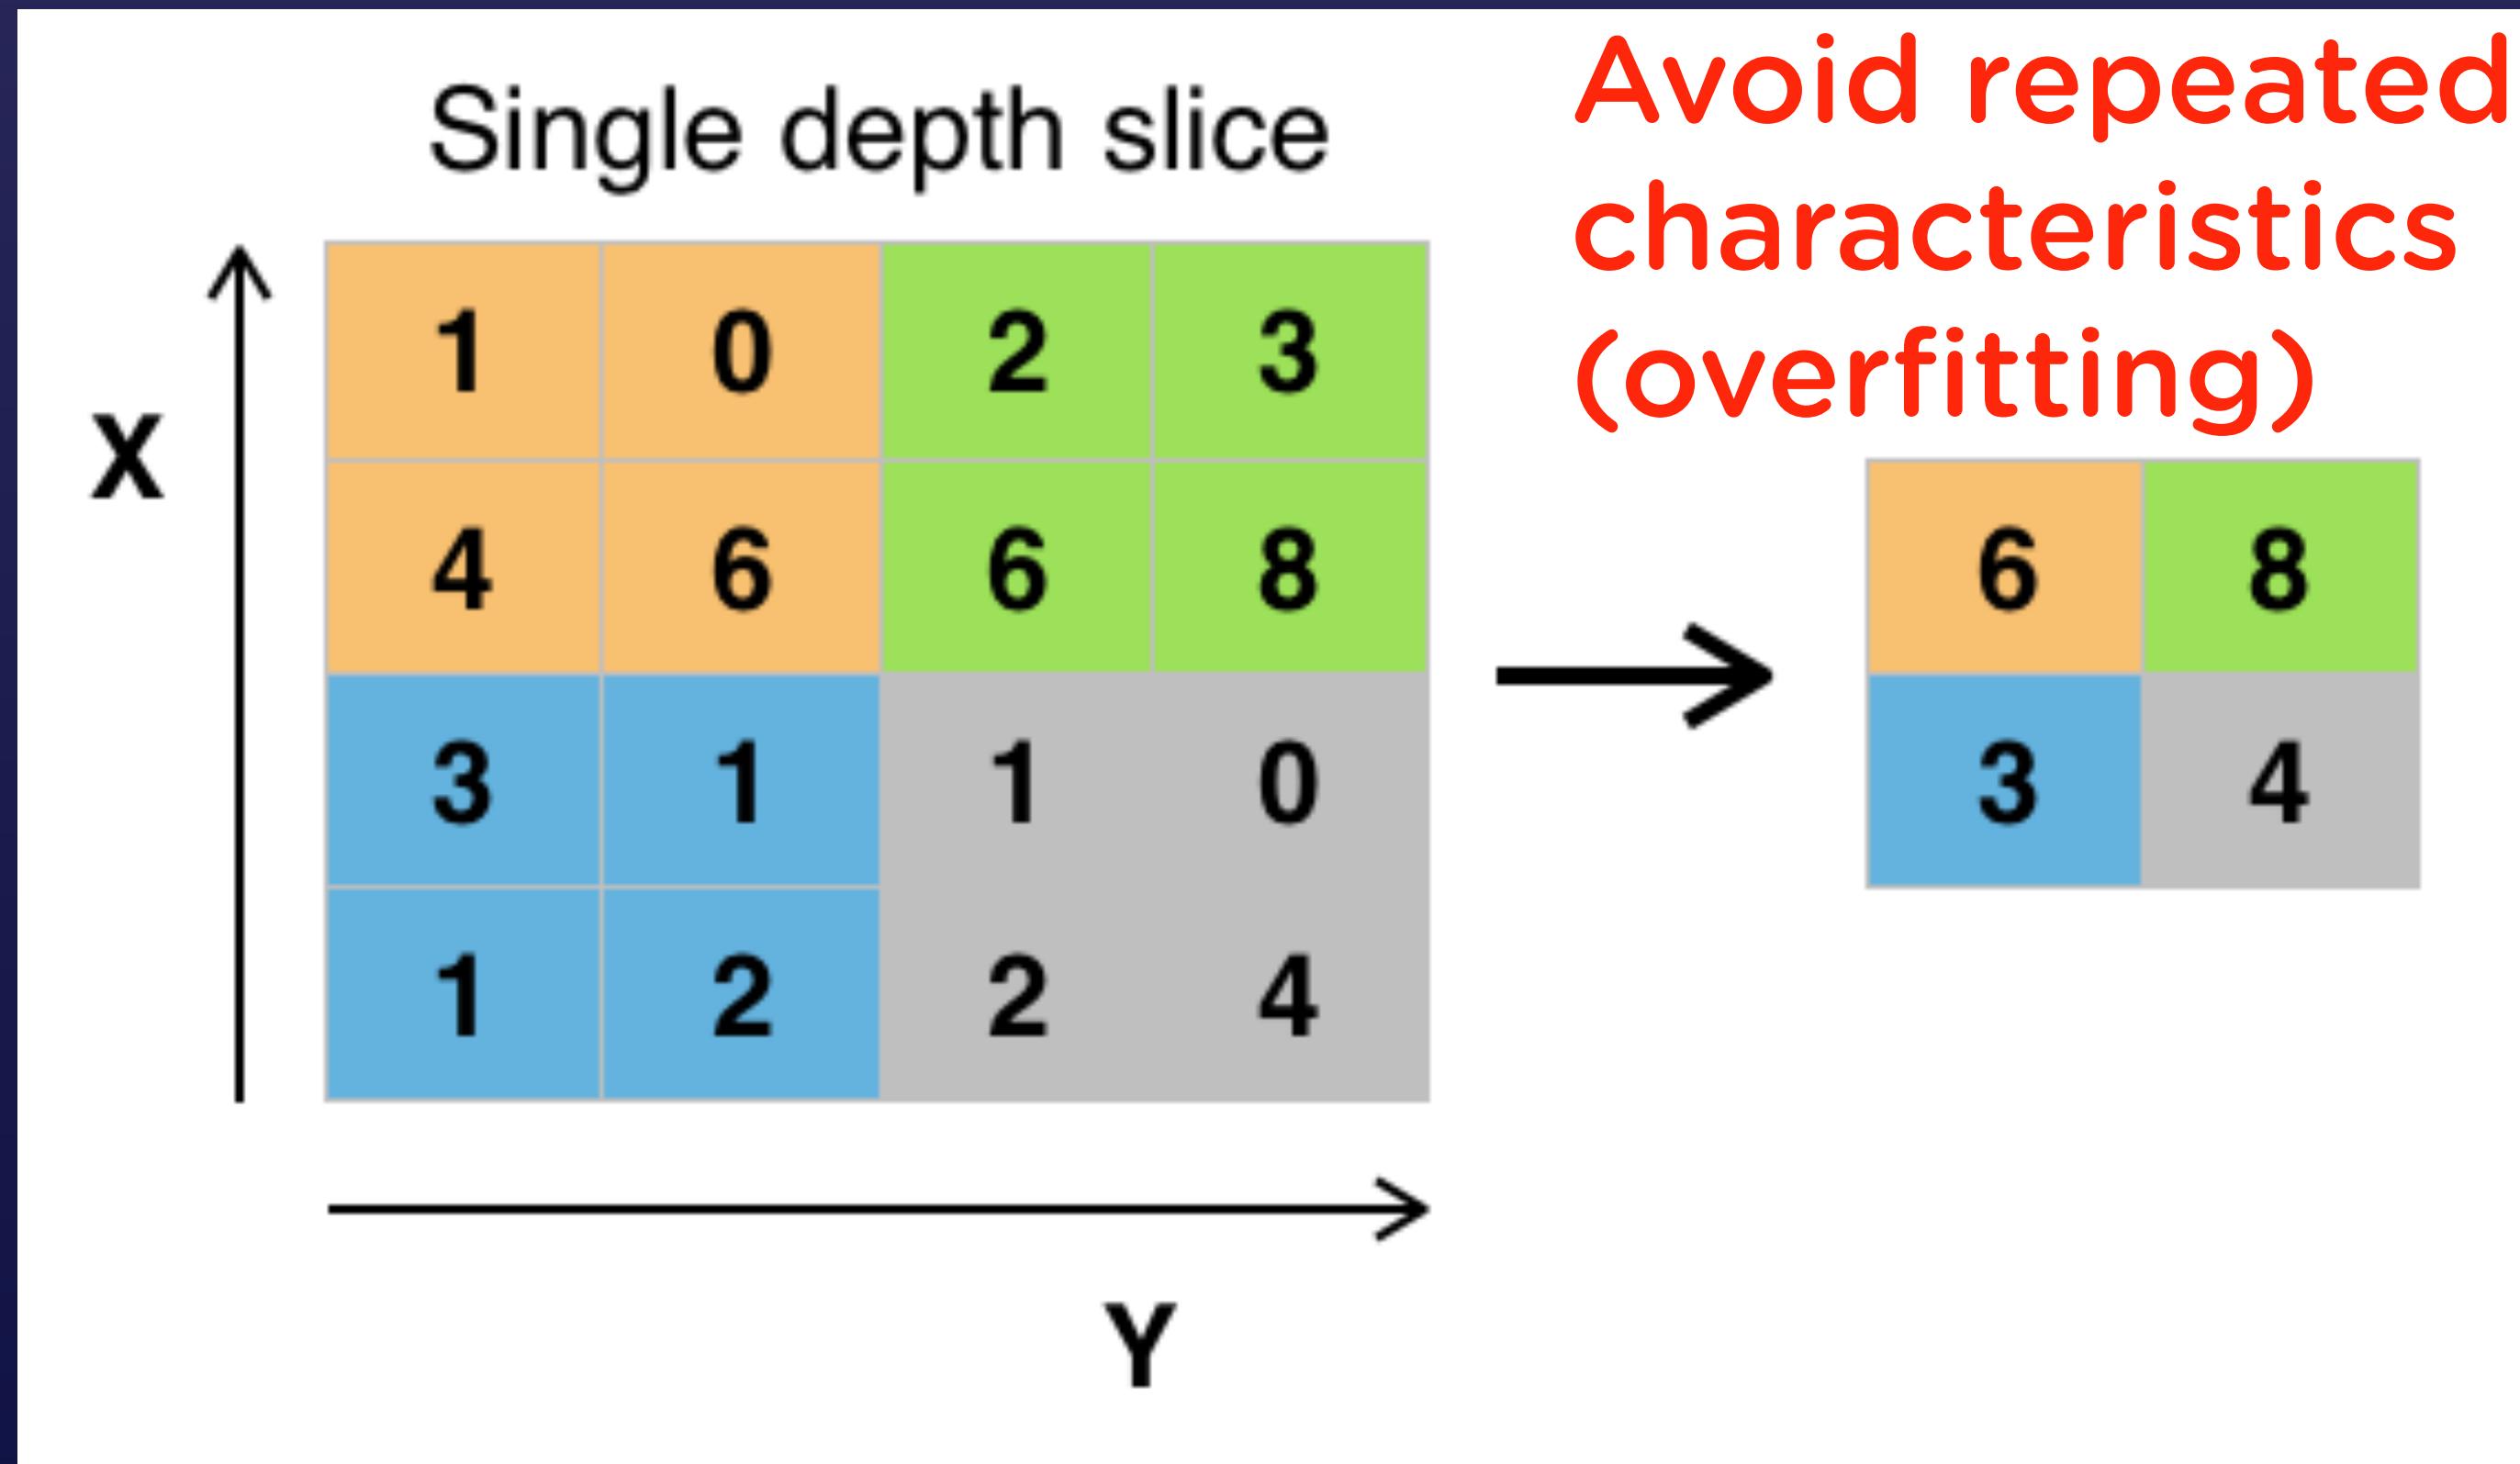


Parameters  
are “shared”



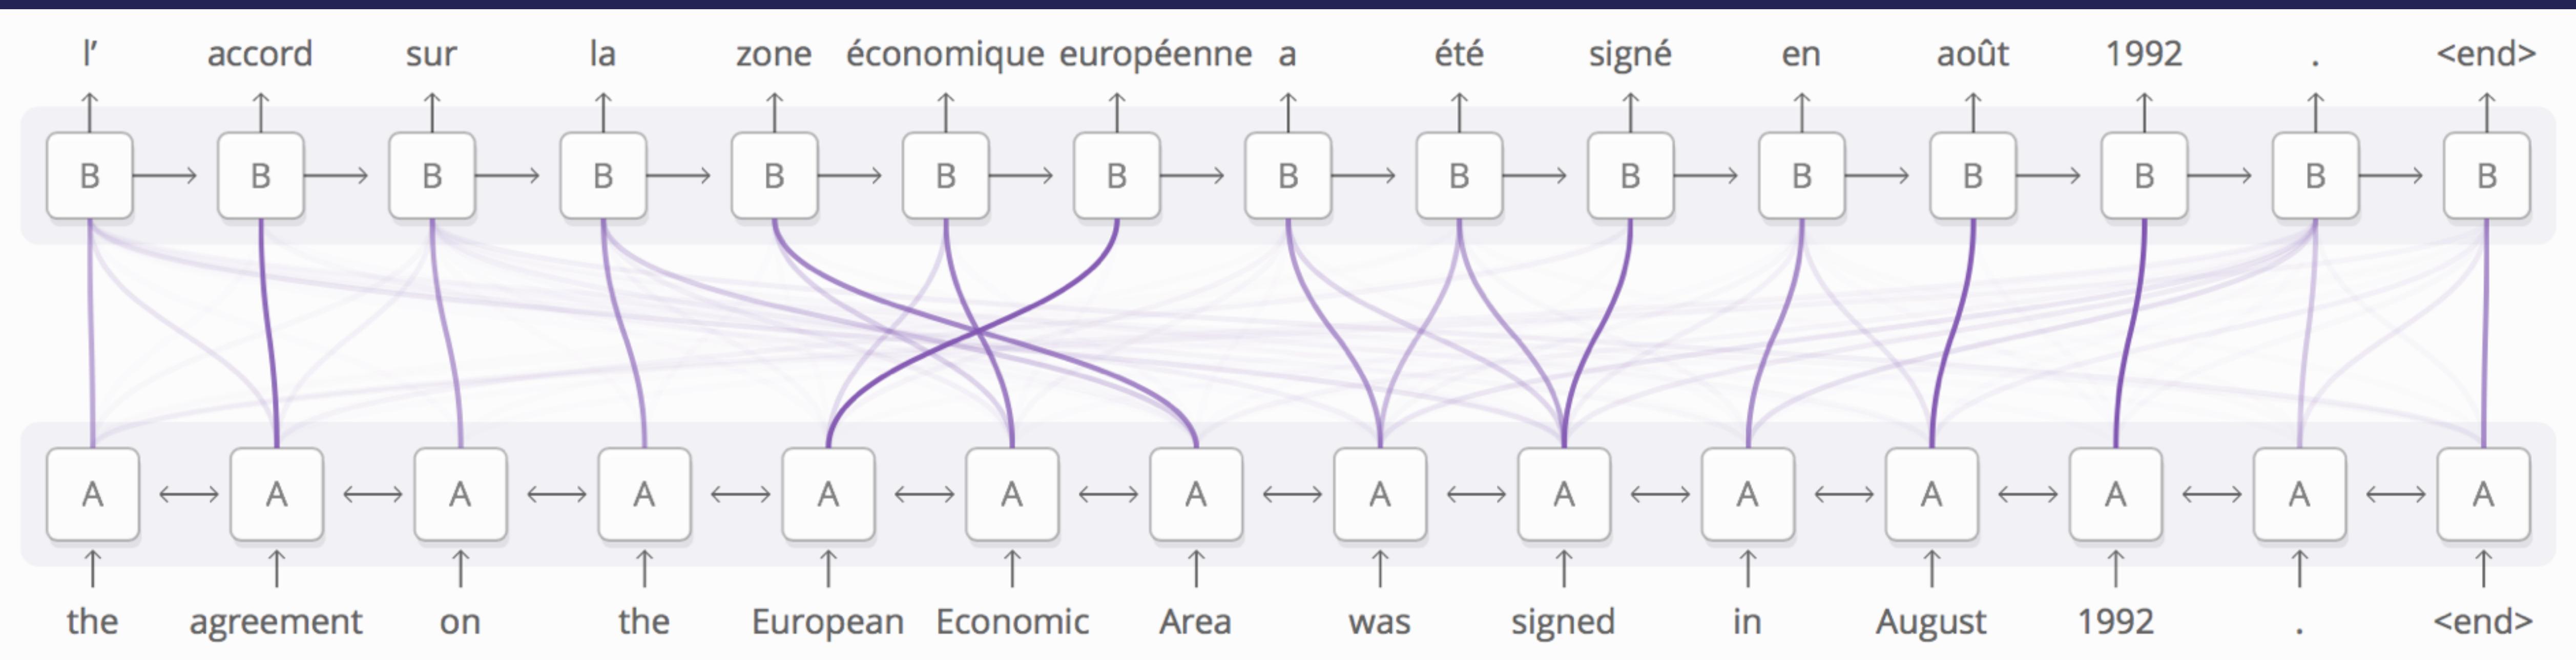
# Convolutional Neural Network

## Max Pooling



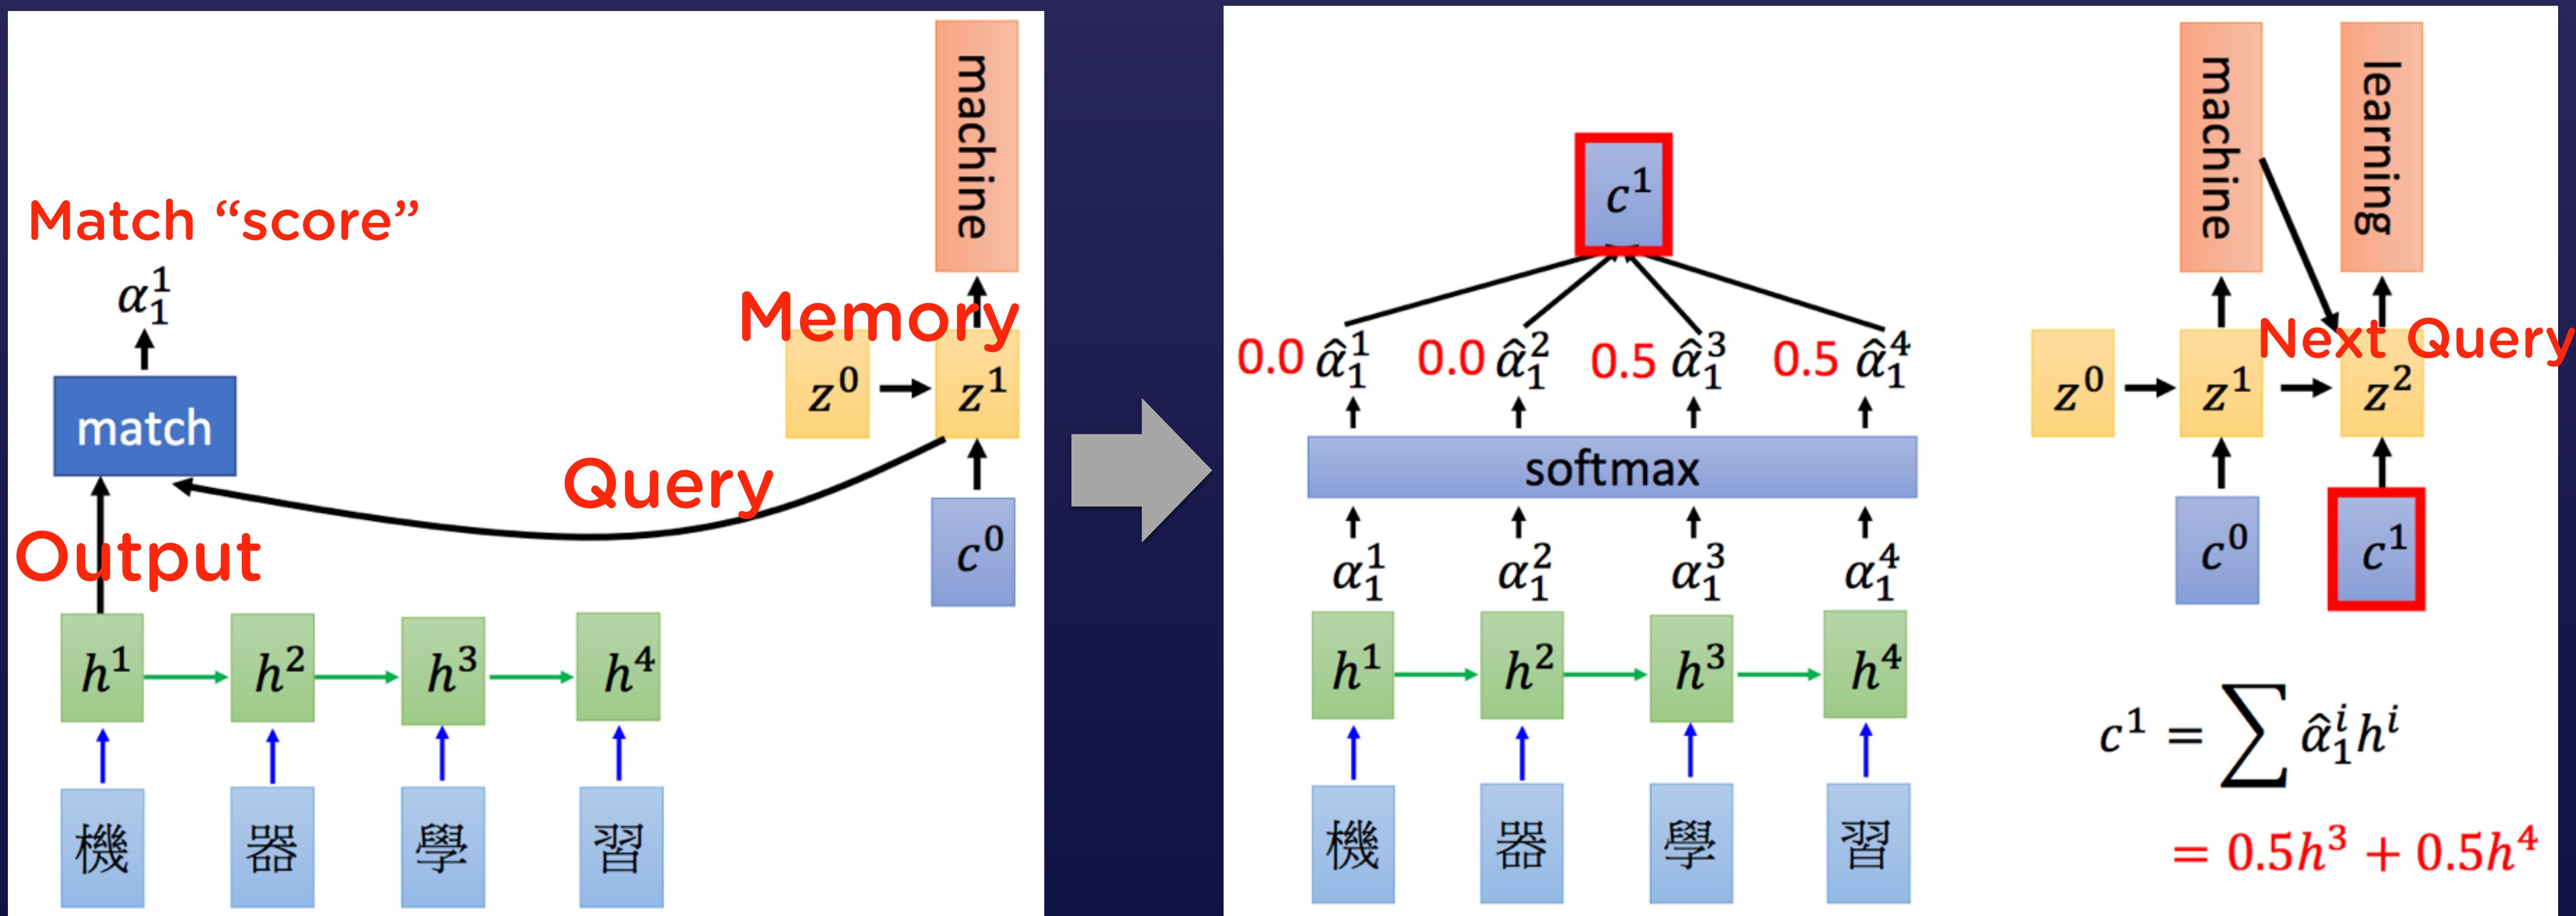
# Attention Mechanism

**Attention** : See the bigger picture  
and focus on somewhere “important”



# Attention Mechanism

Attention How? : “Query” and “Match”



# Attention Mechanism

## Match?

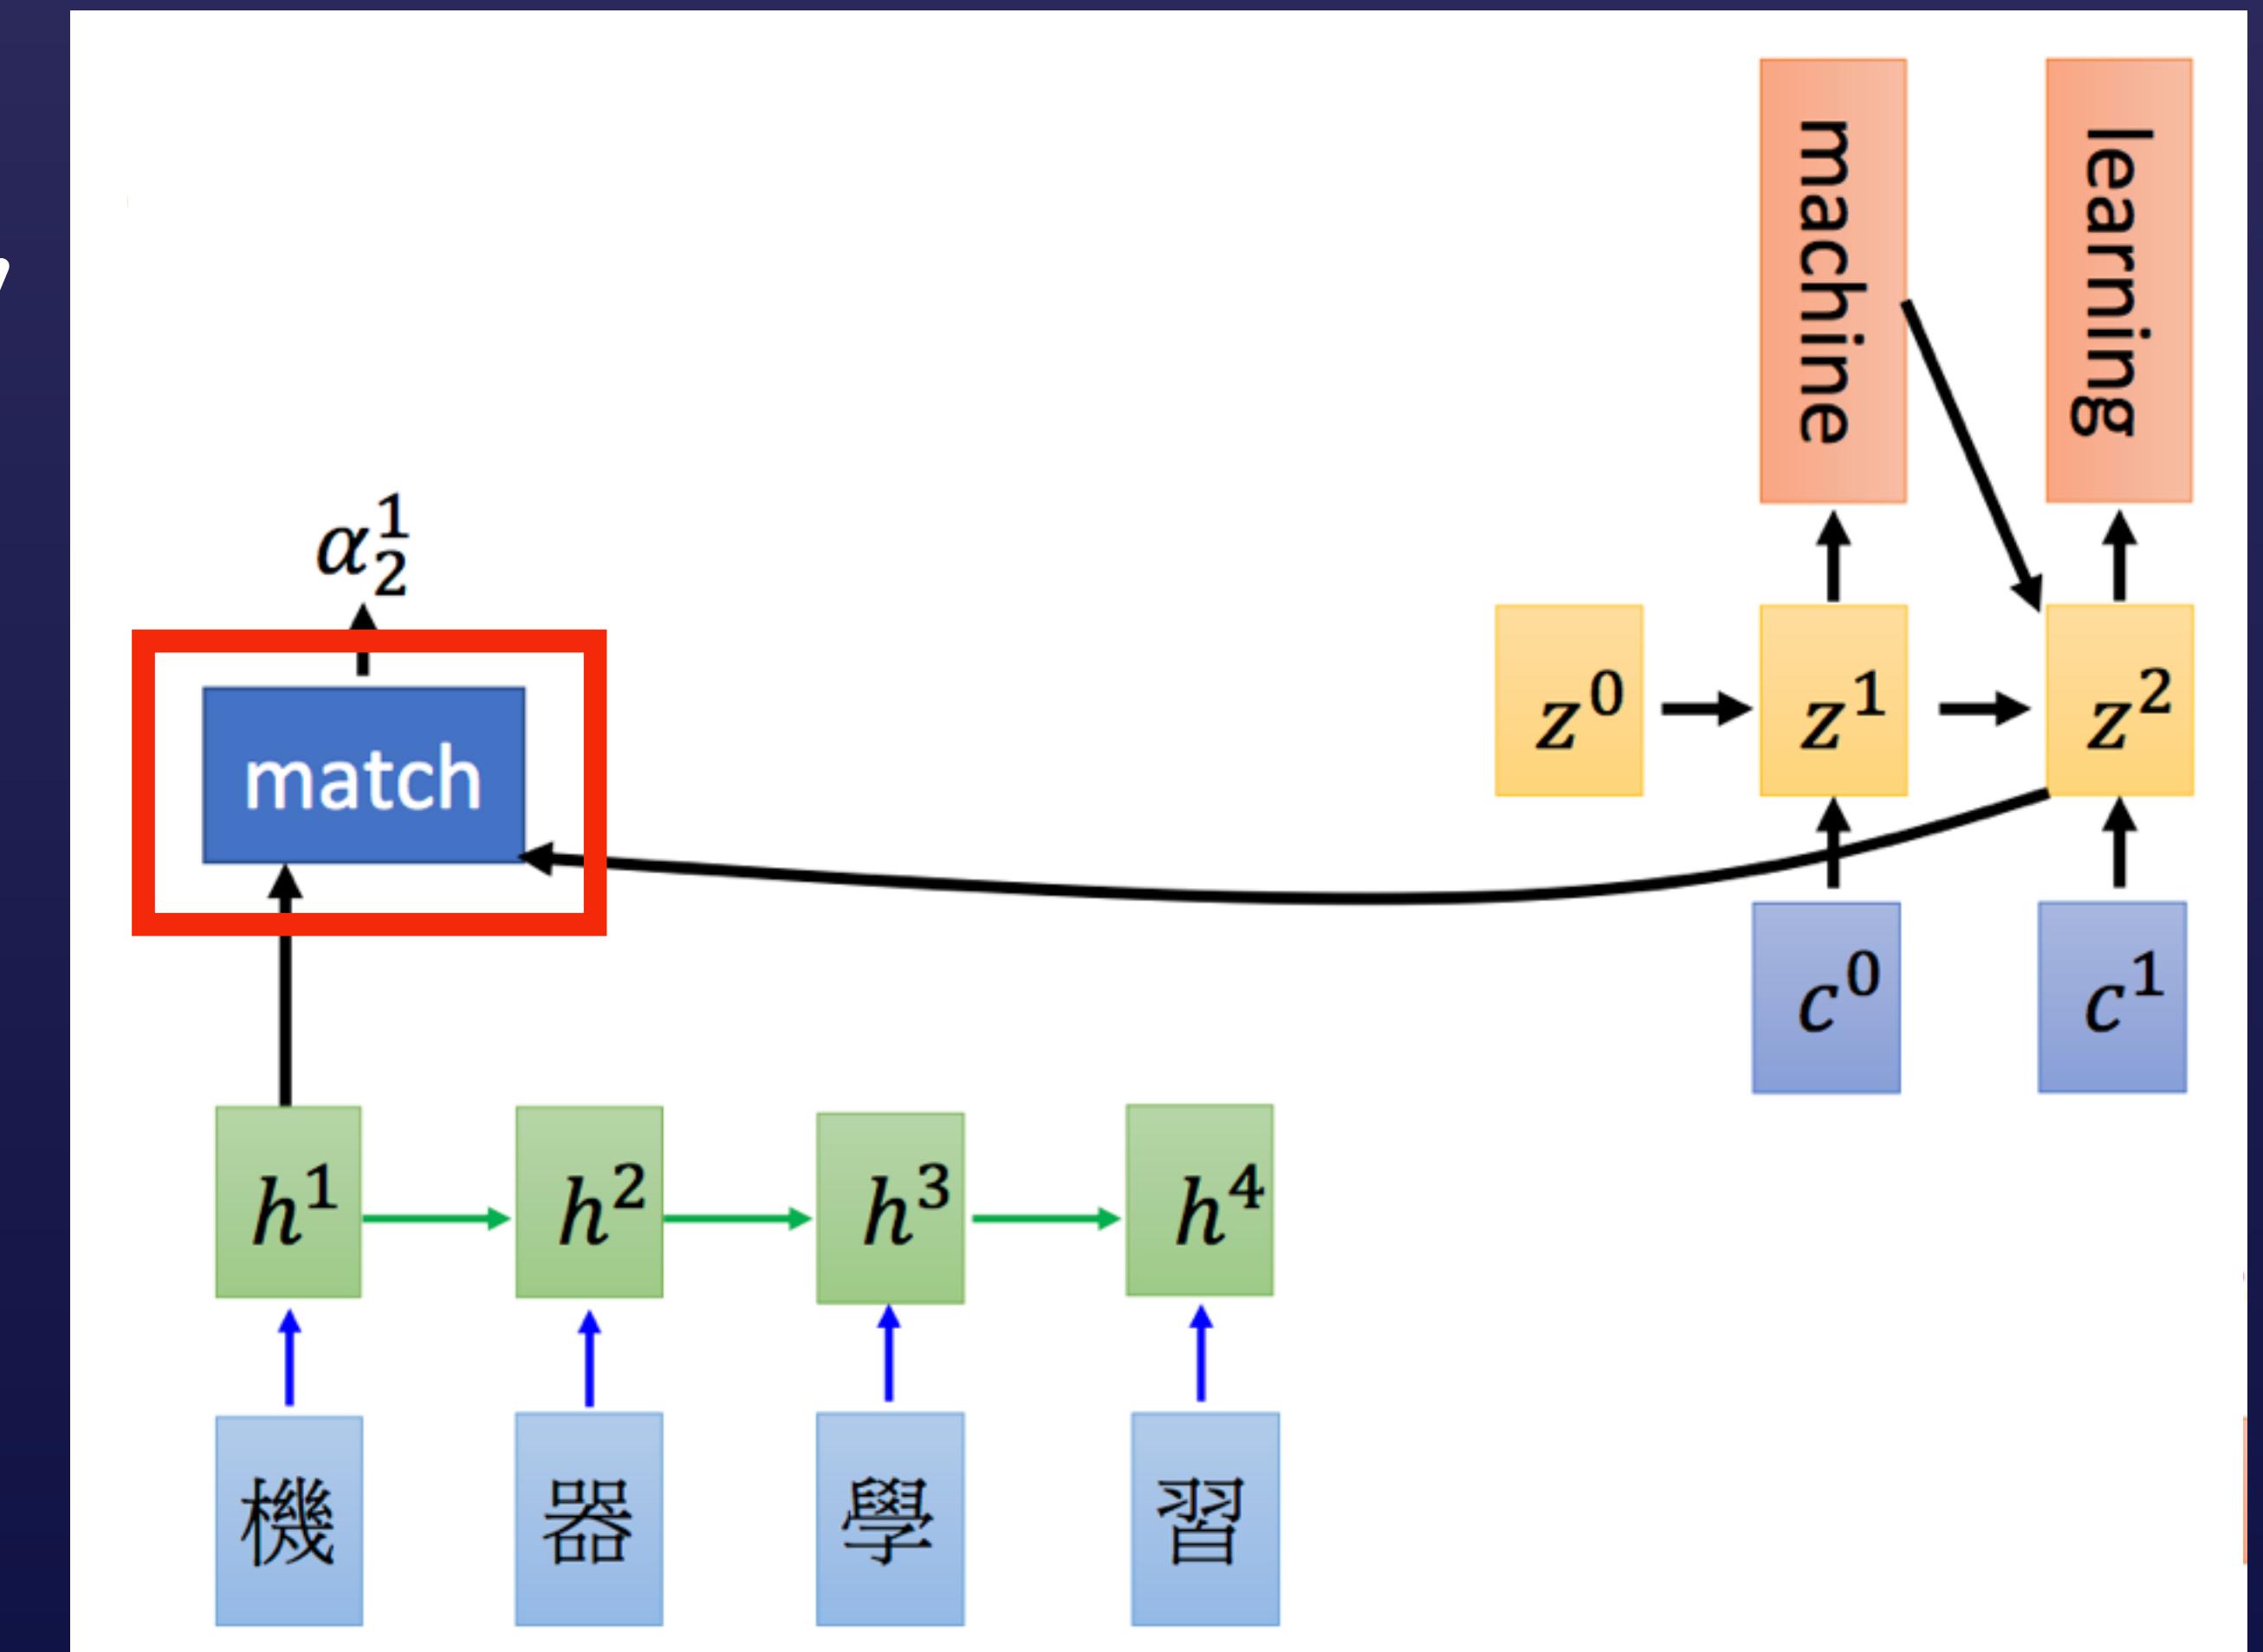
- Matchness of “**h**” and **memory**
- Some kind of “similarity”

Dot product

Cosine similarity

Score :  $h'Wz$

Small neural network



# Attention Mechanism

## Image Captioning



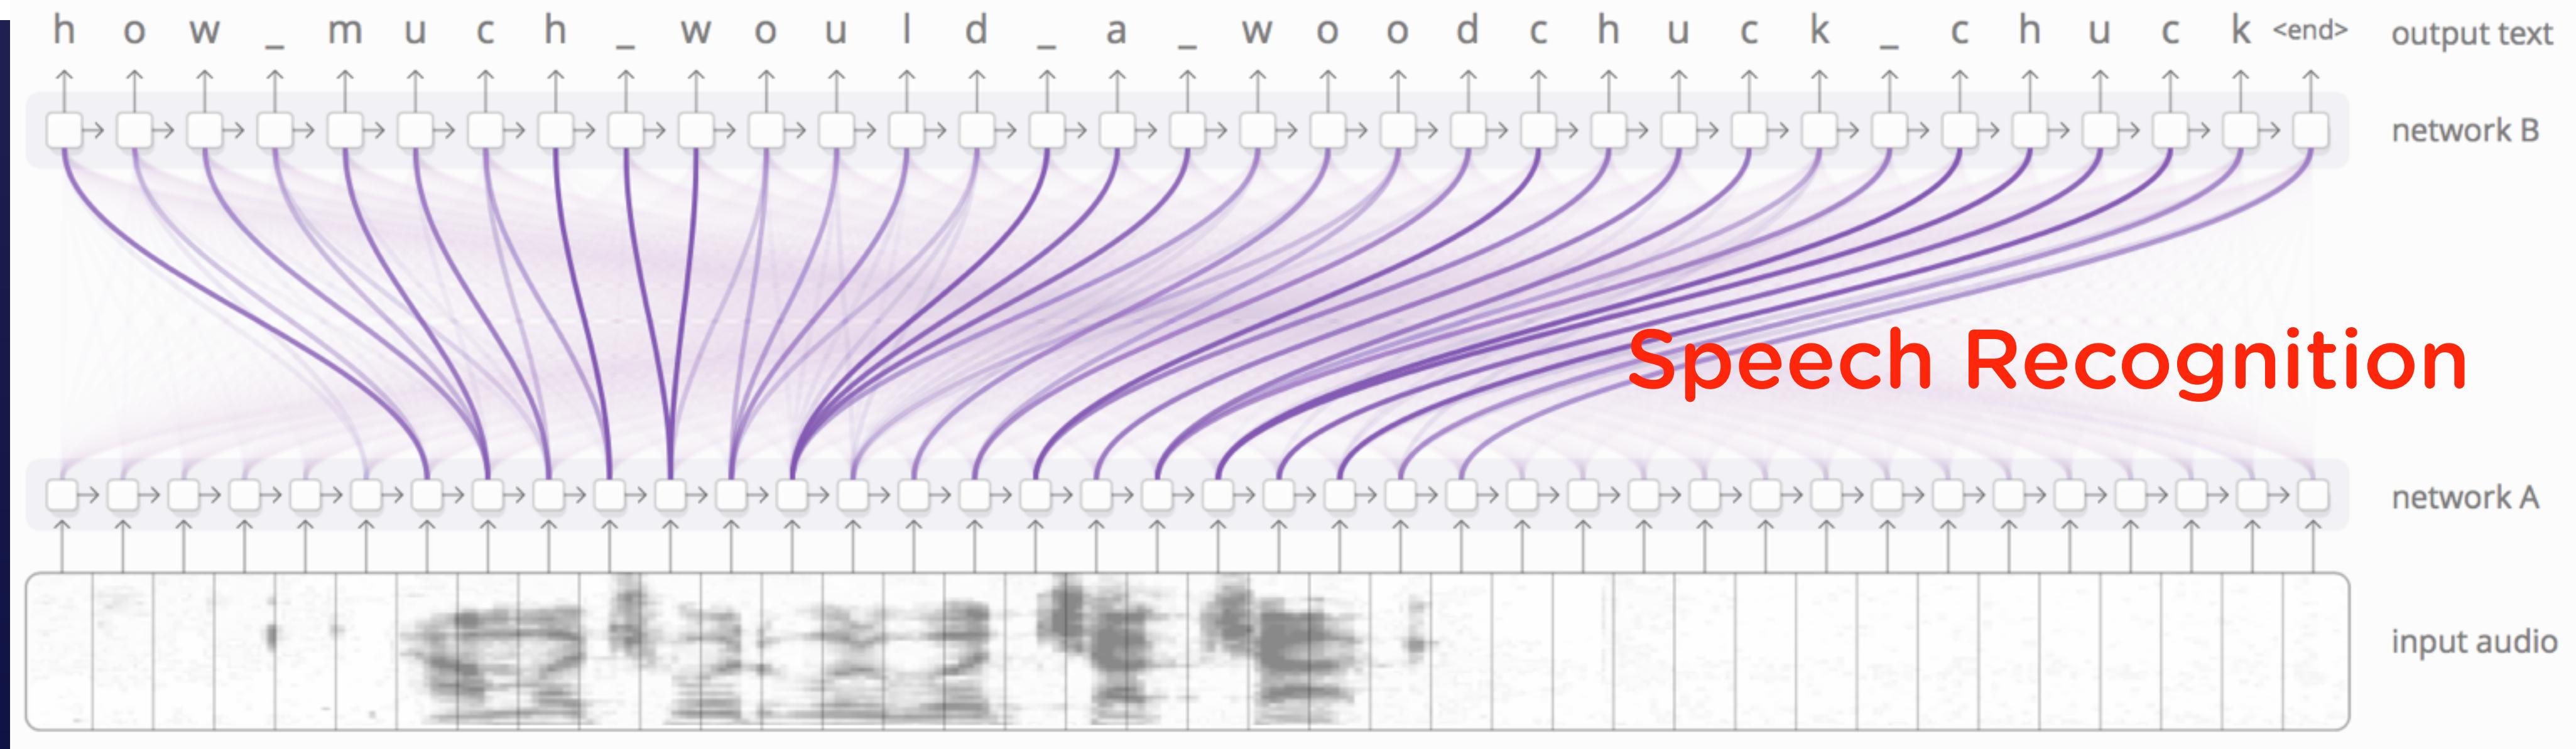
A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.

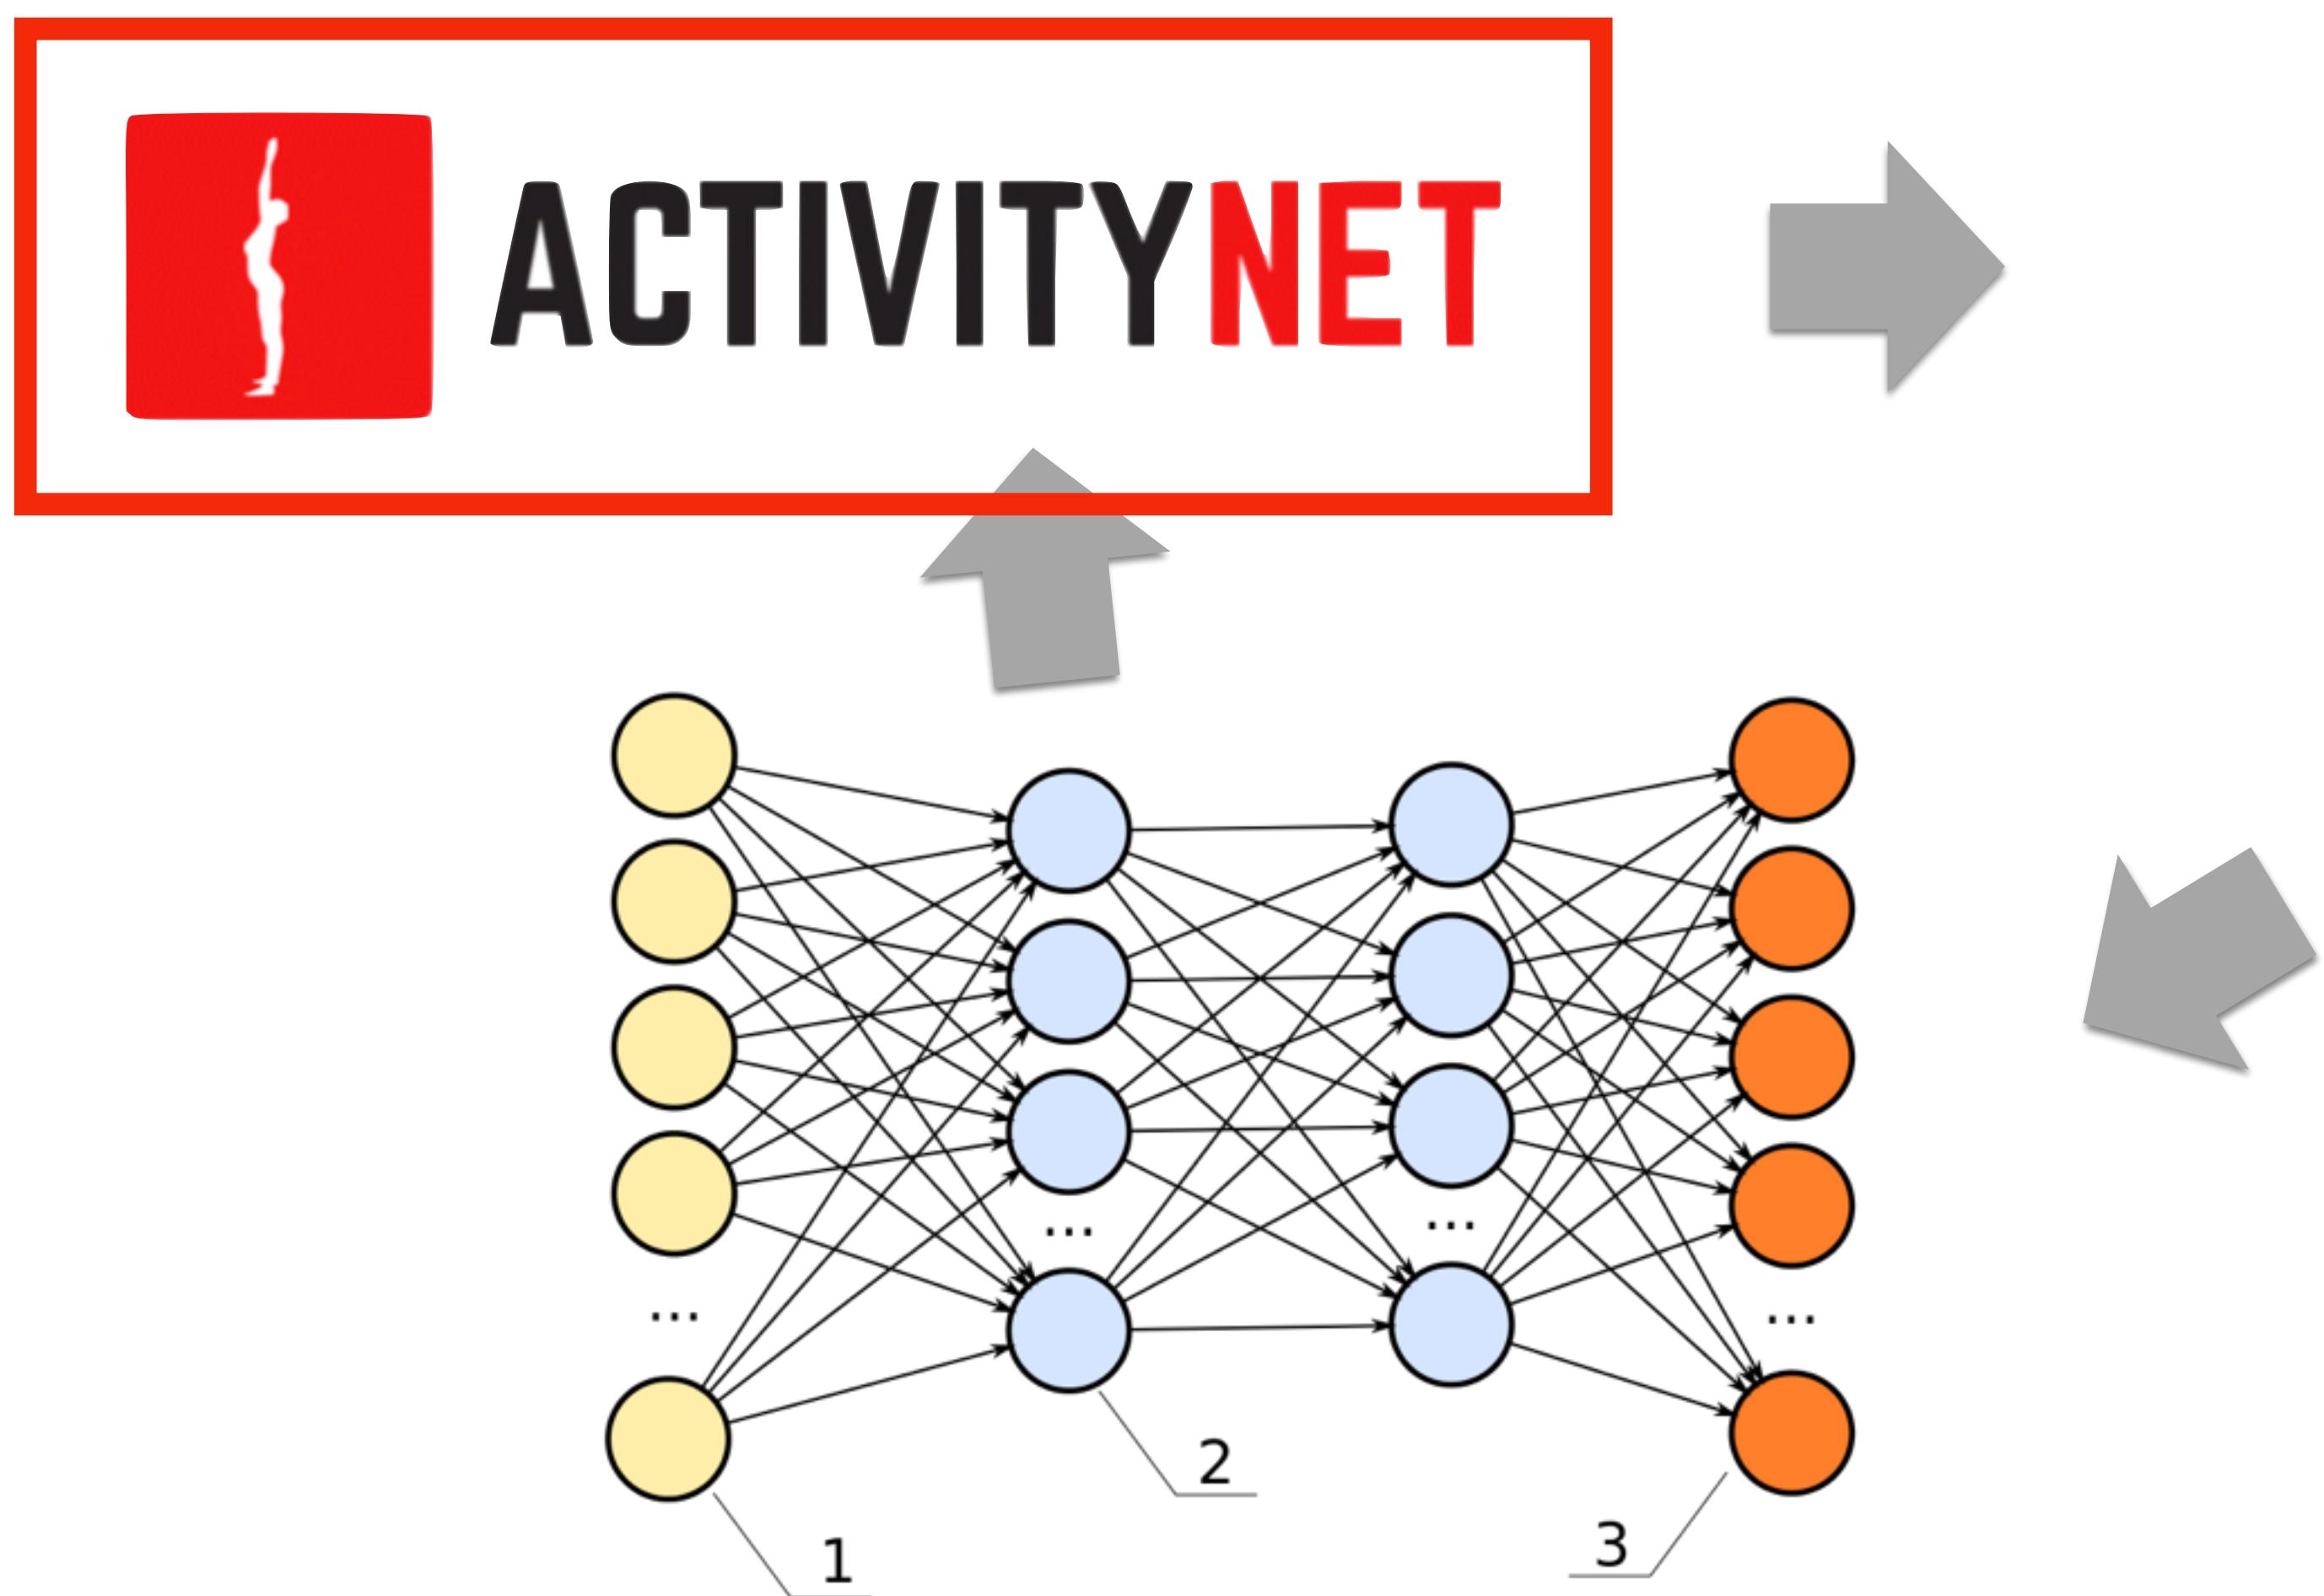


A stop sign is on a road with a mountain in the background.



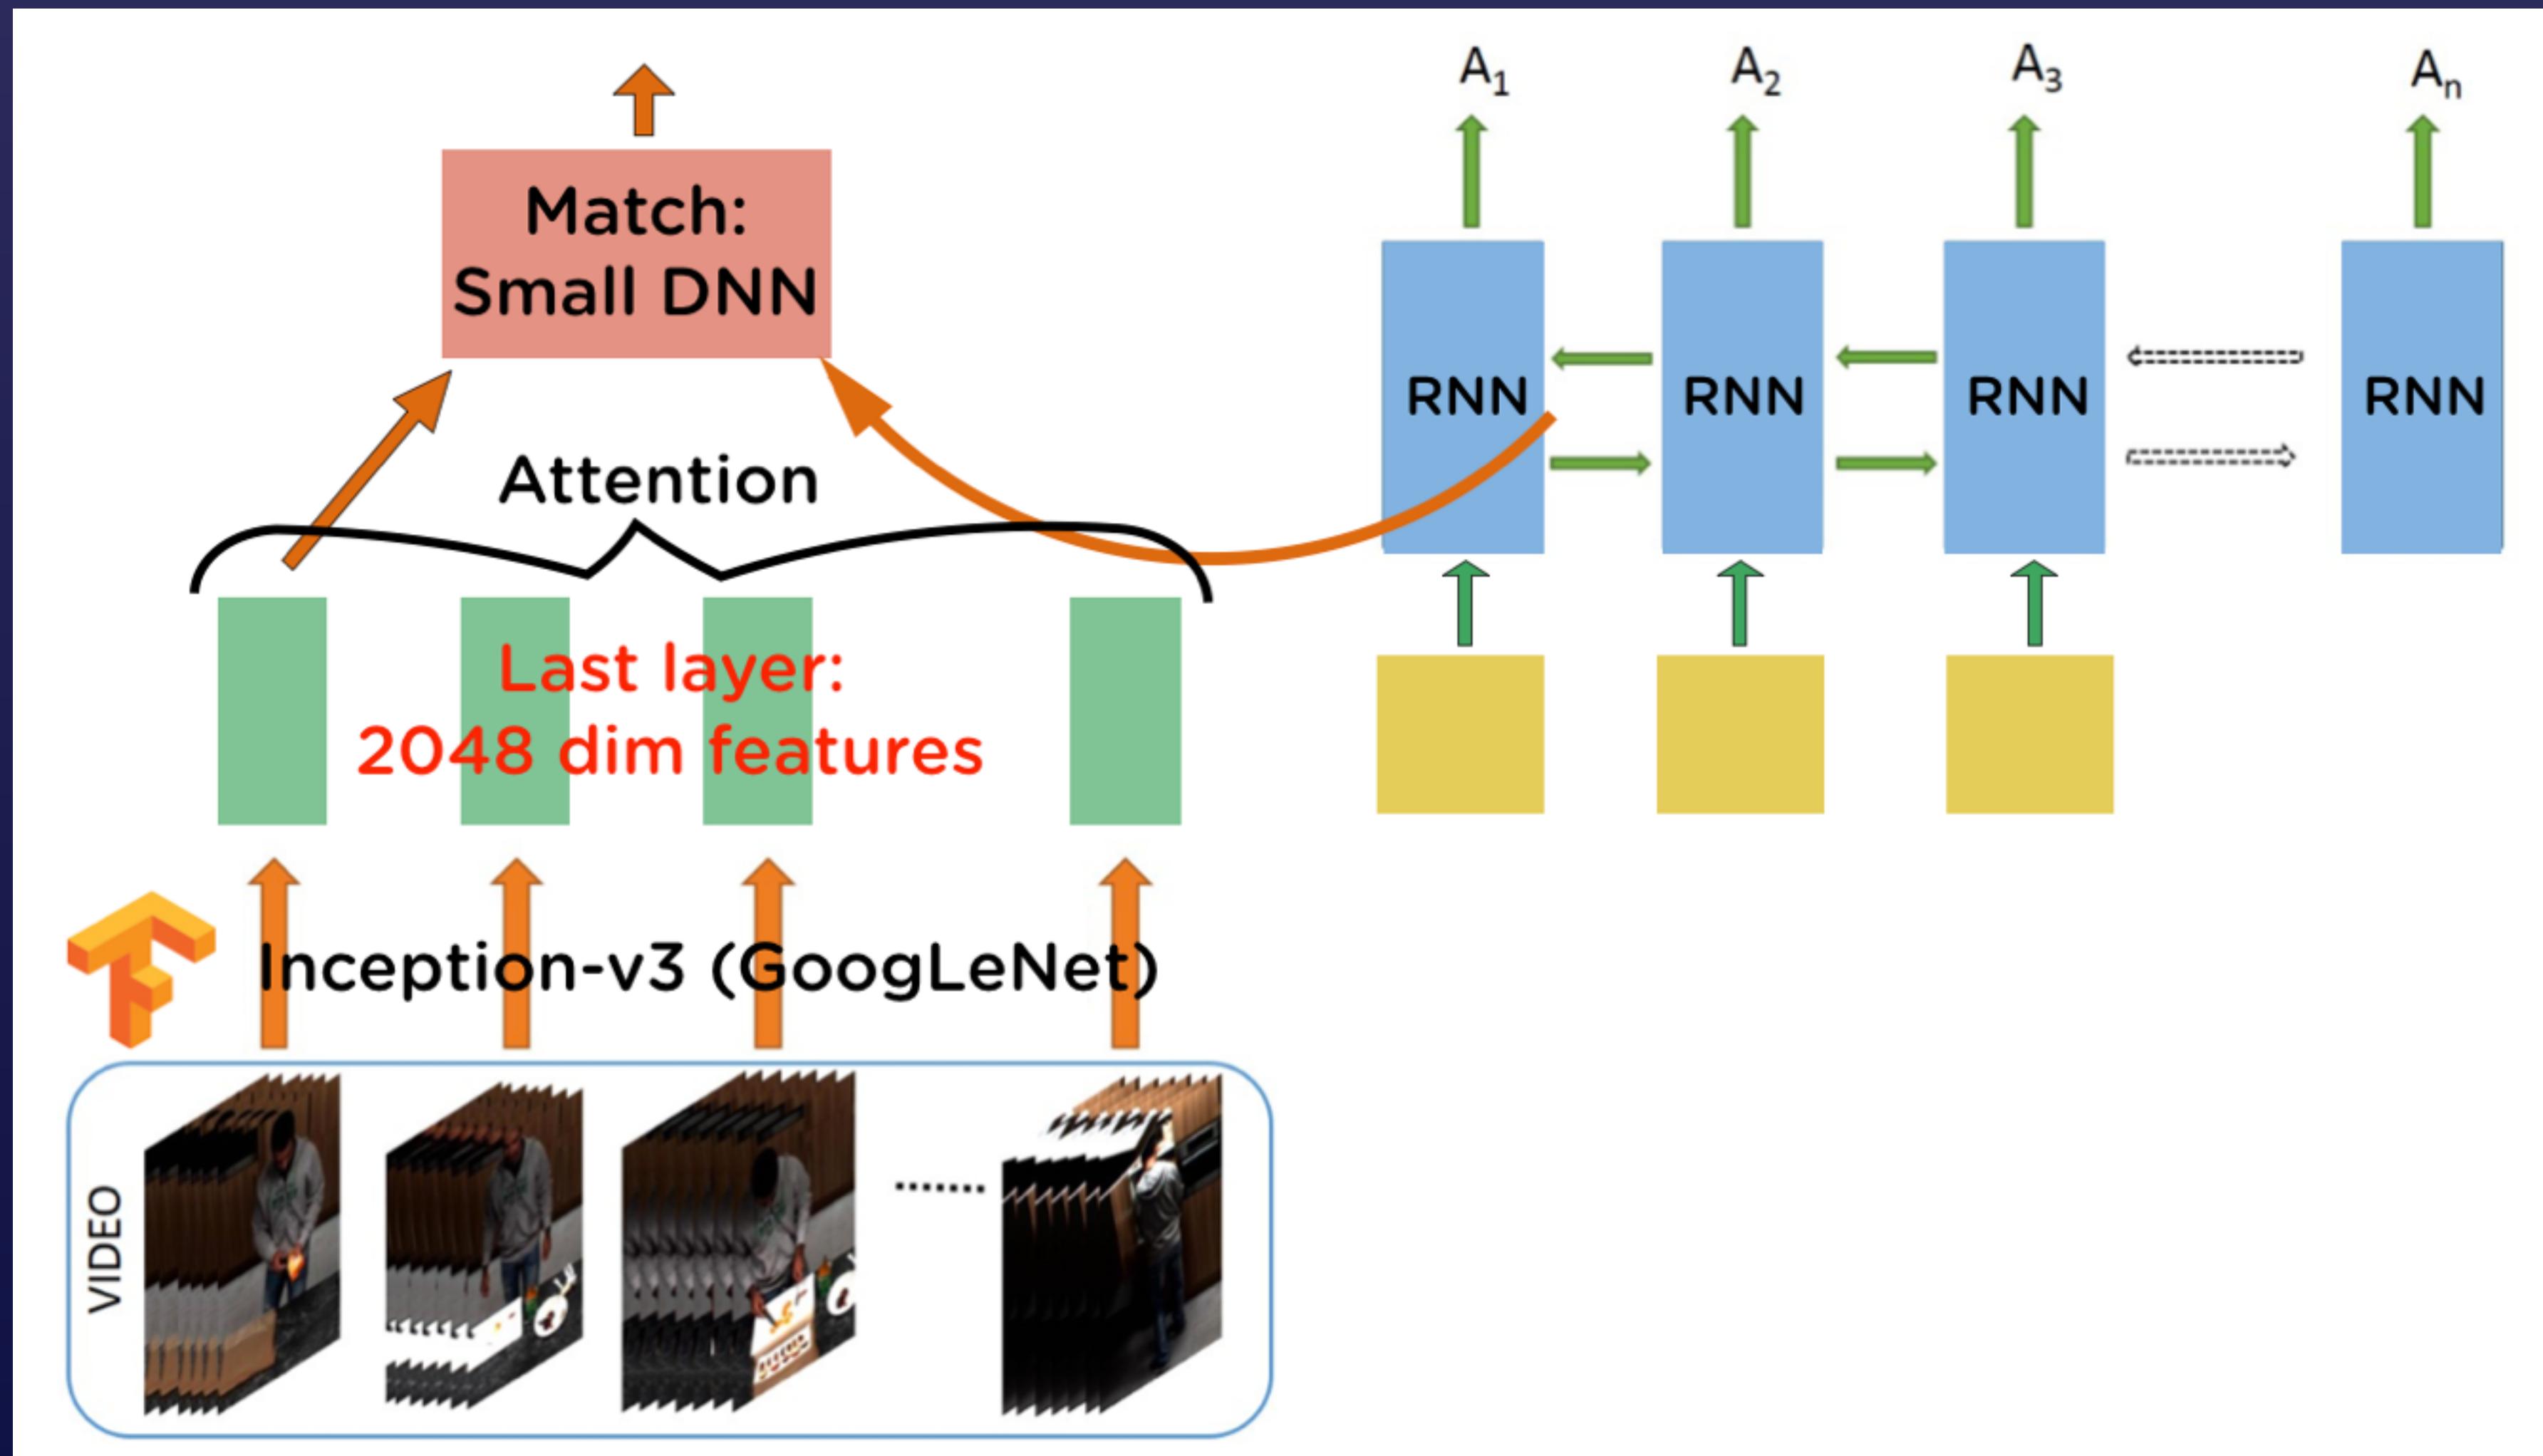
Reference : <http://distill.pub/2016/augmented-rnns/>

# Outline



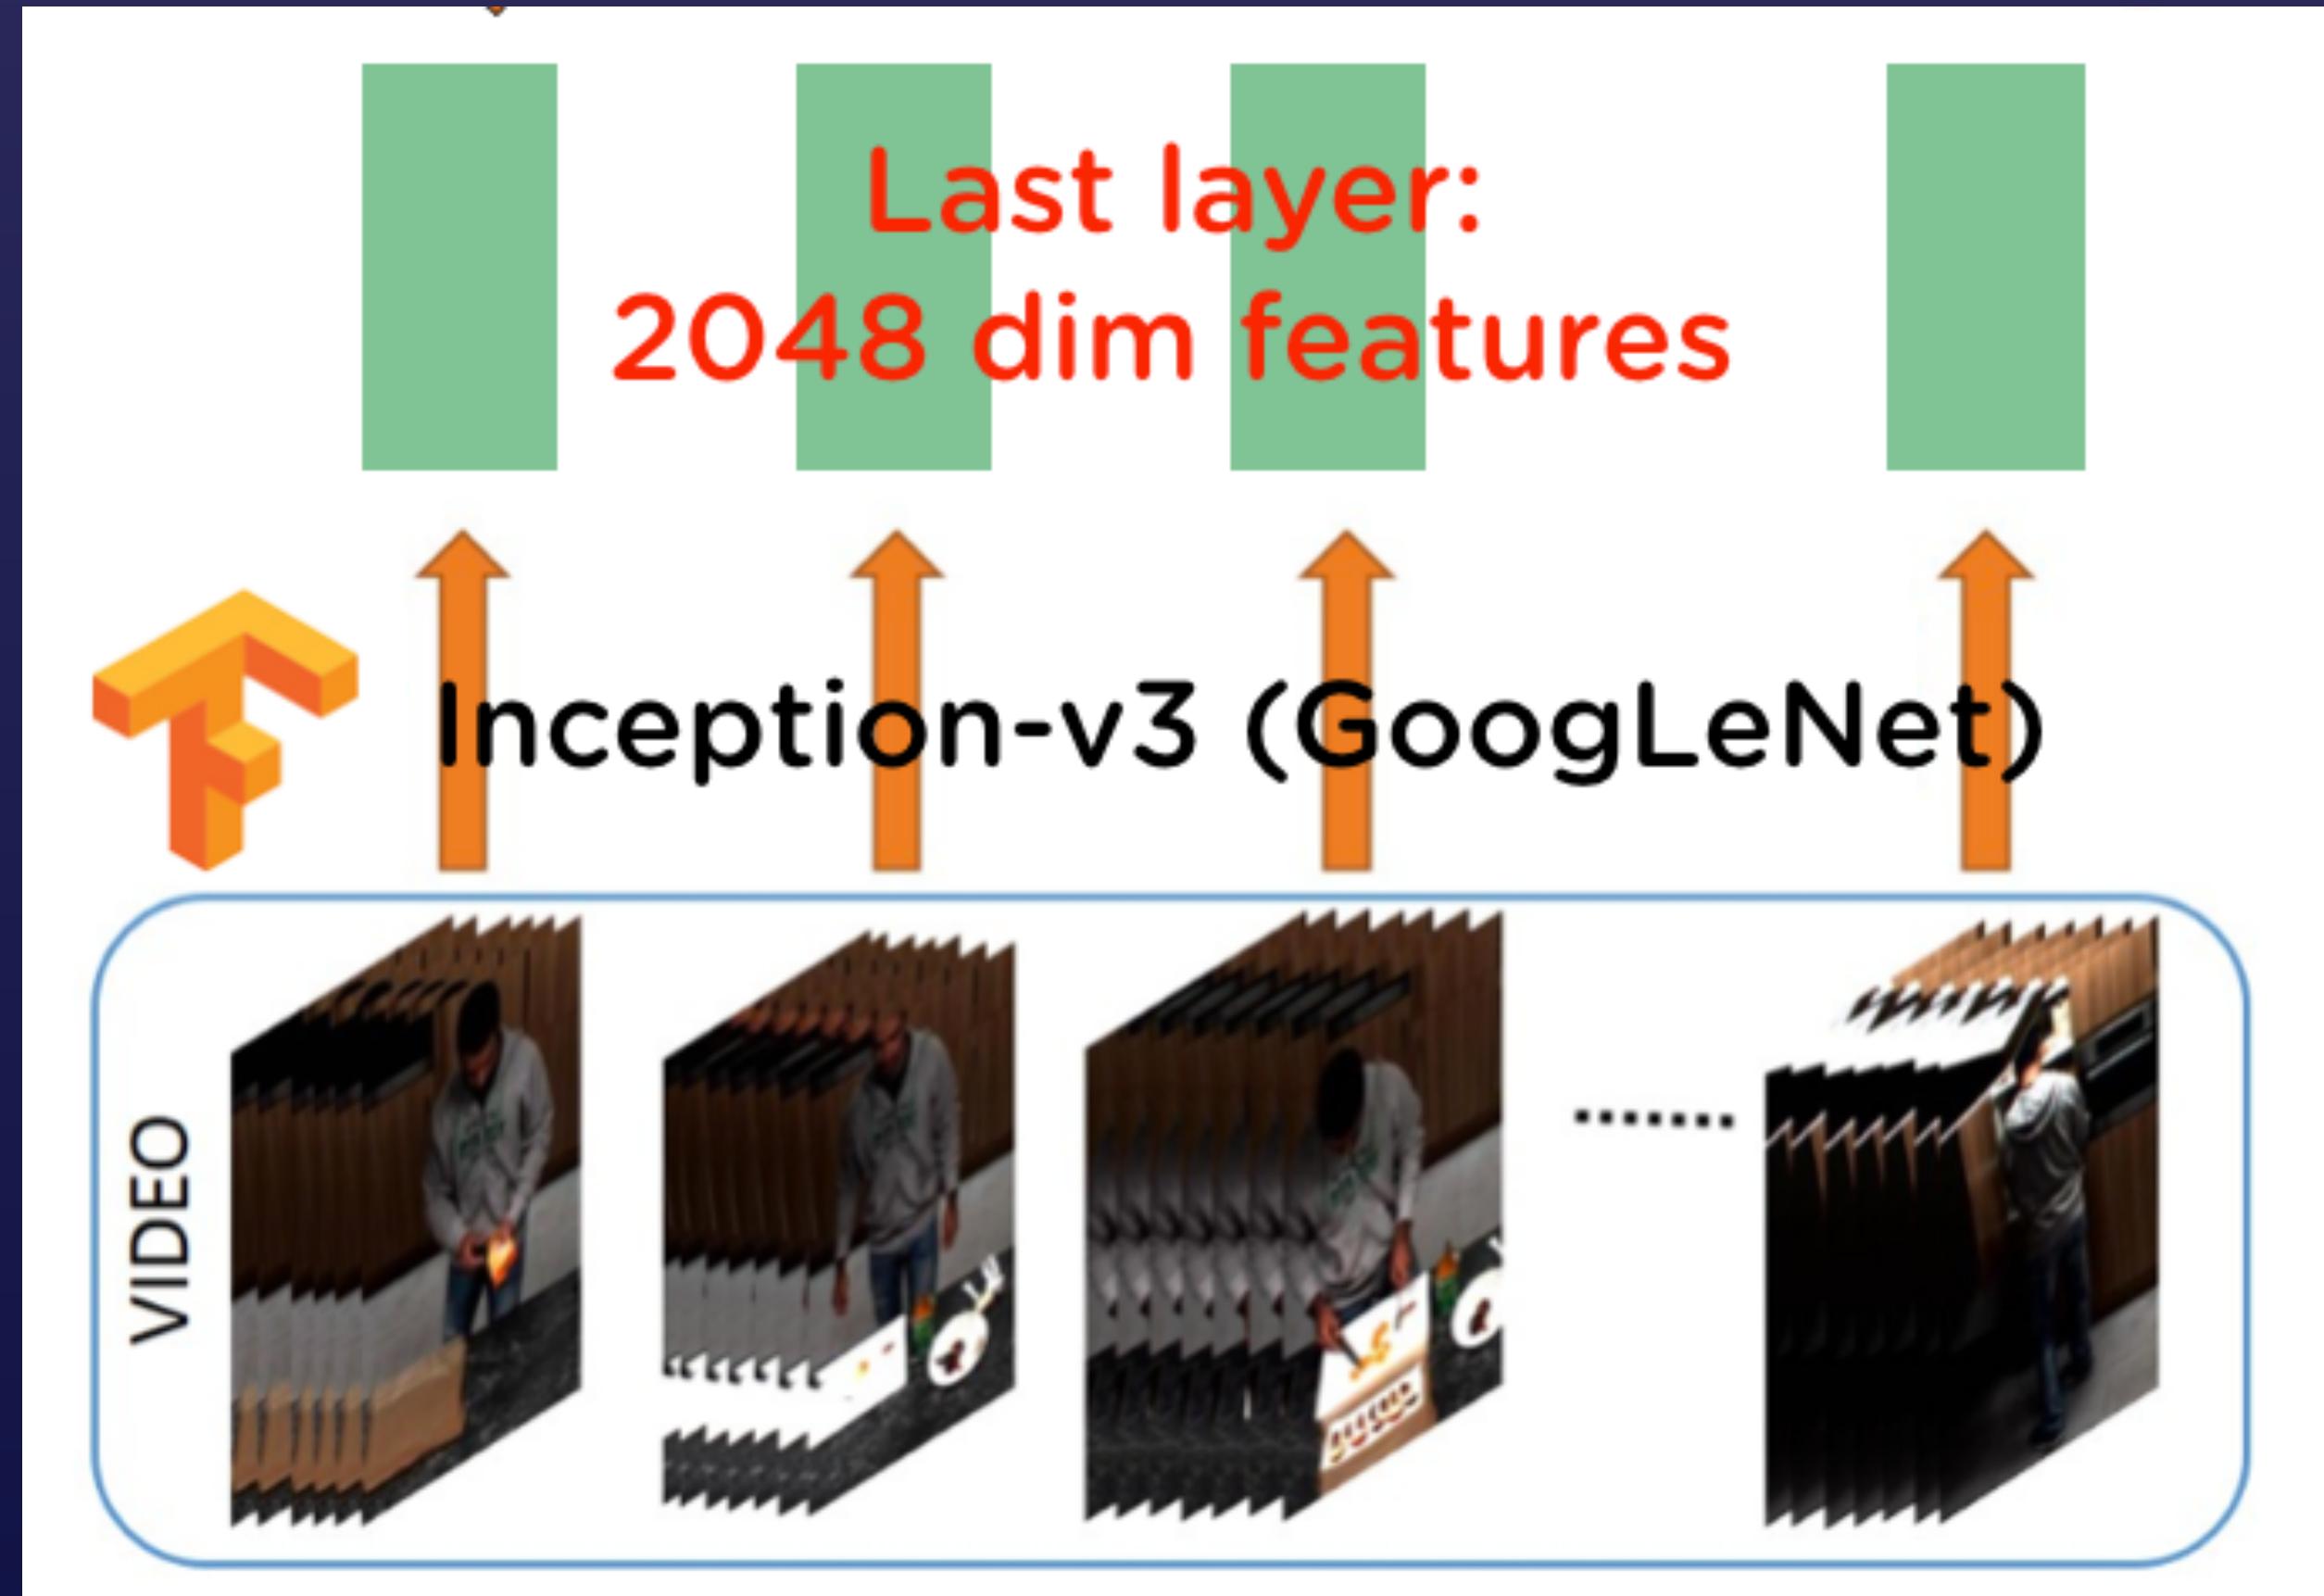
TensorFlow

# Putting All Things Together



Reference : <http://www.cs.umd.edu/~bharat/cvpr2016.pdf>

# Frame Features



## Convolution Neural Network

Pre-trained  
Inception-v3



Reference : [https://www.tensorflow.org/versions/r0.12/tutorials/image\\_recognition/index.html](https://www.tensorflow.org/versions/r0.12/tutorials/image_recognition/index.html)

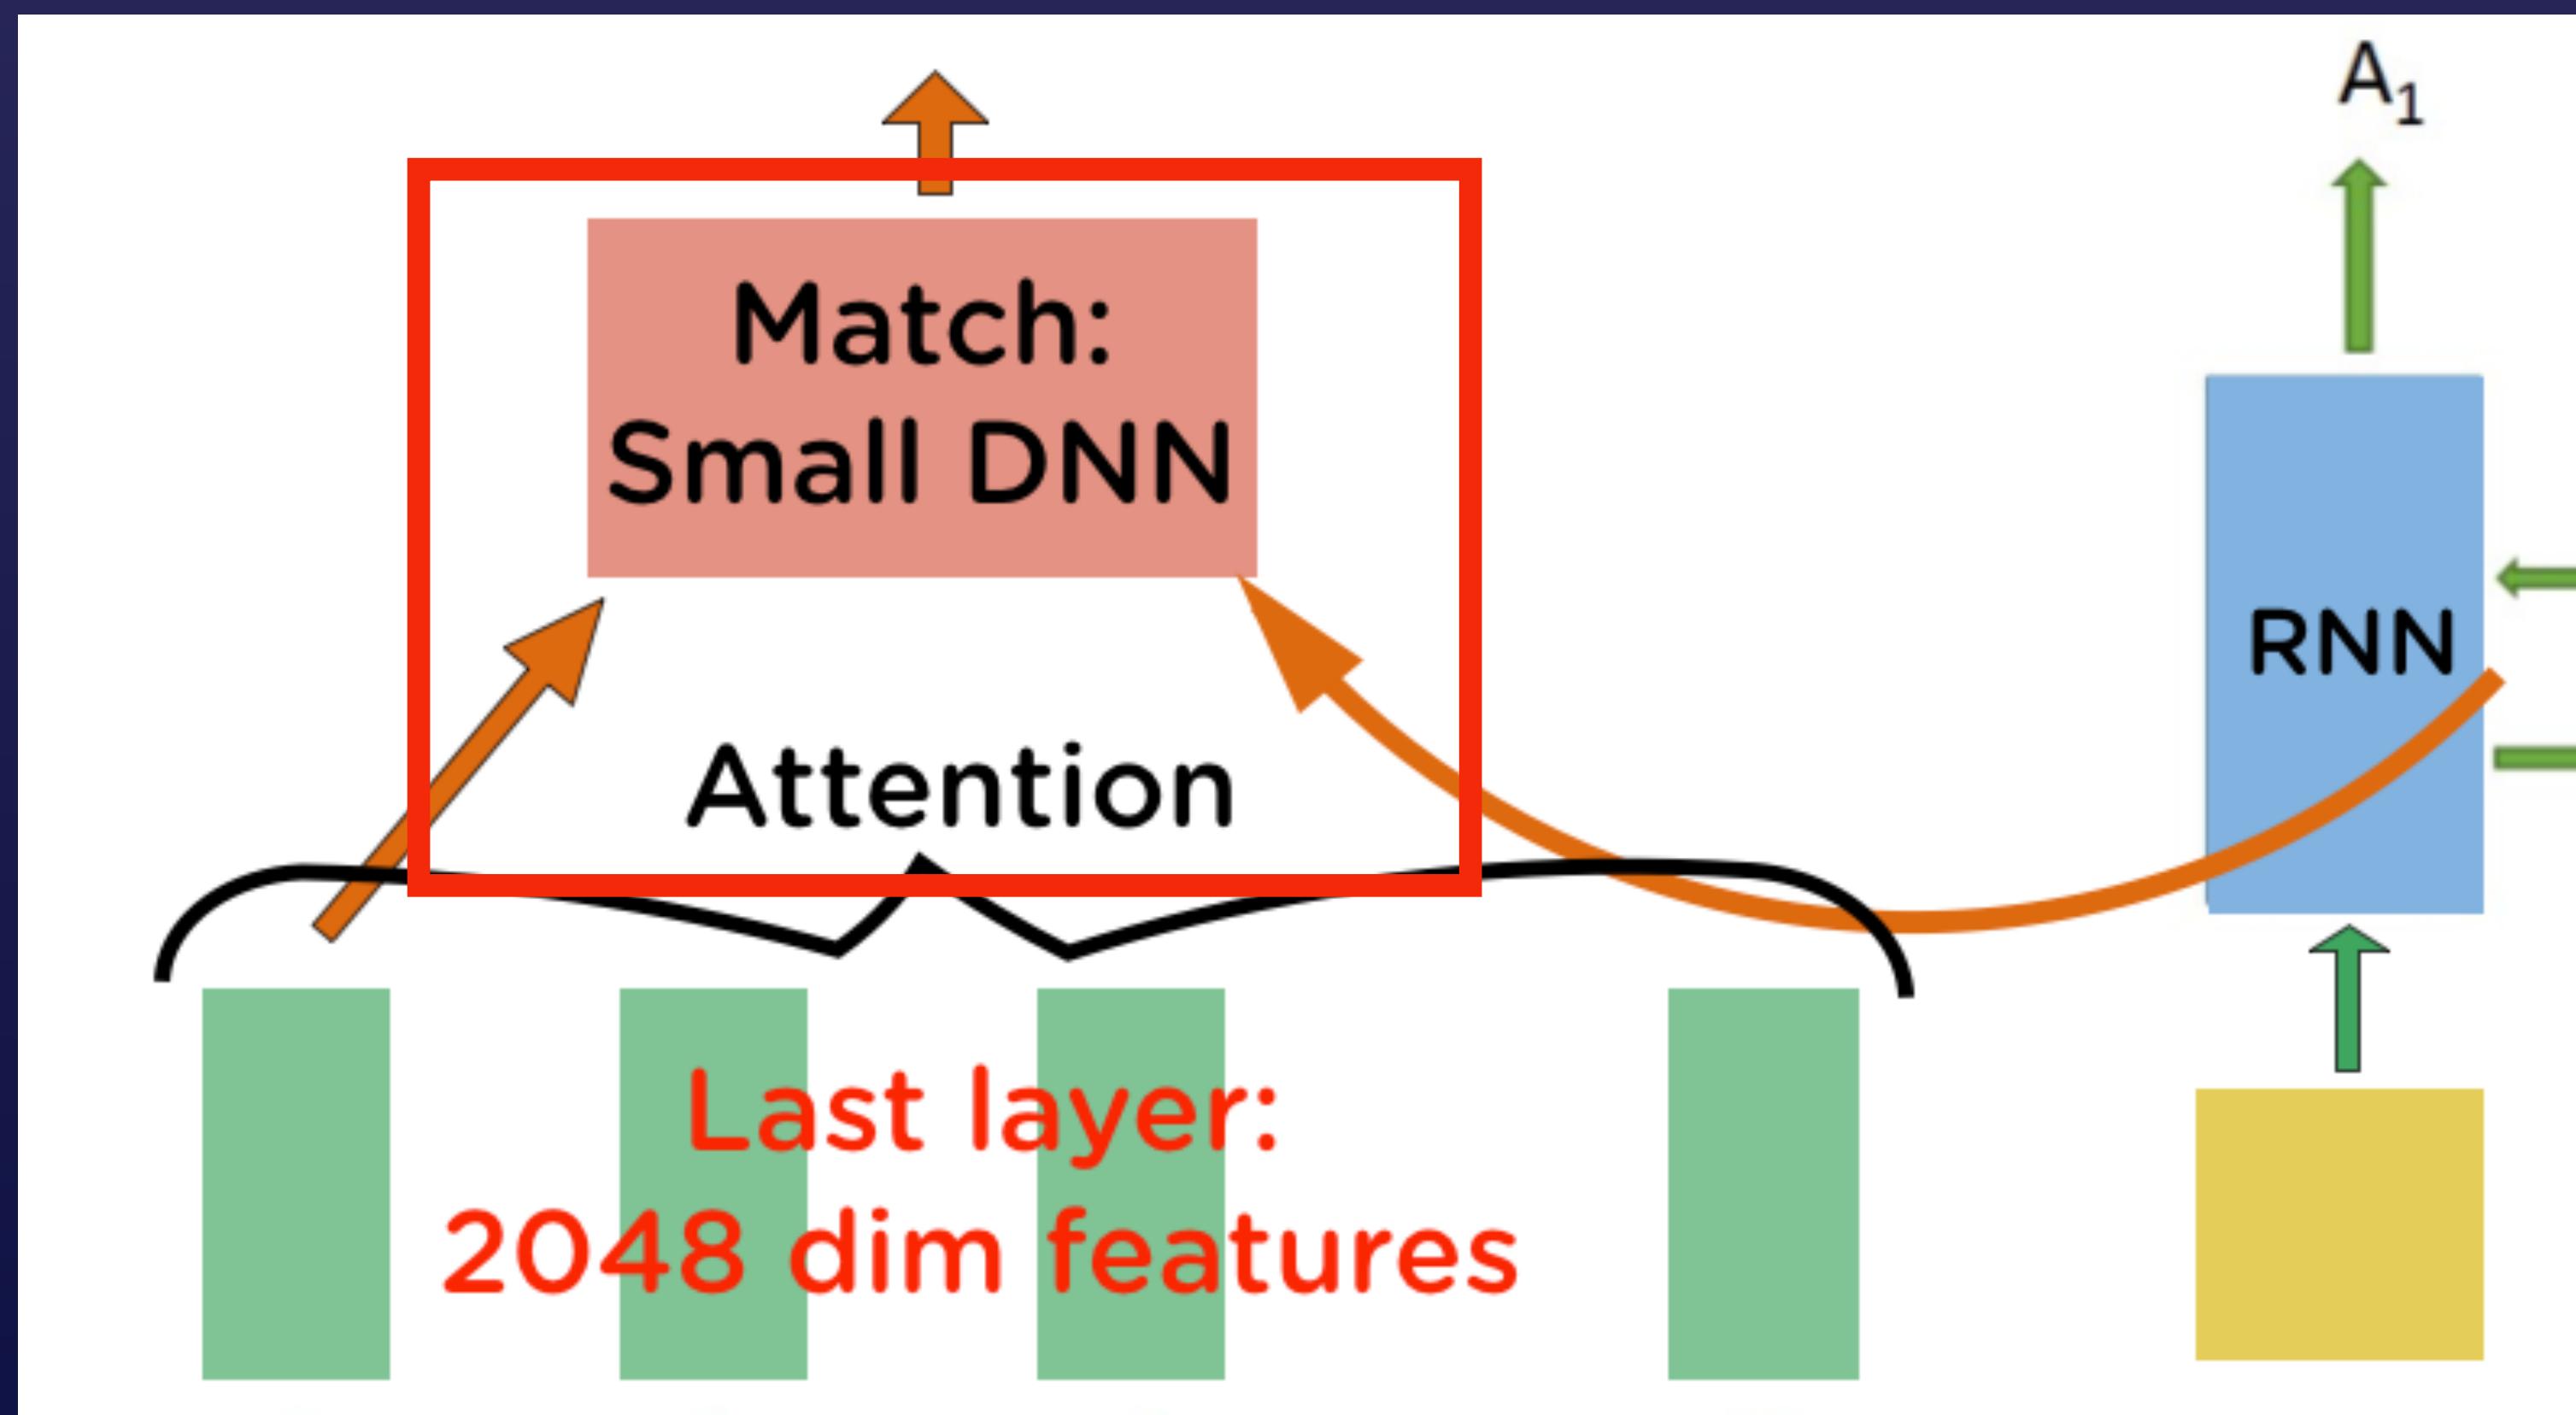
GDG DevFest  
TAIPEI 2016

GLIA CLOUD

# Attention and DNN

## Attention Mechanism

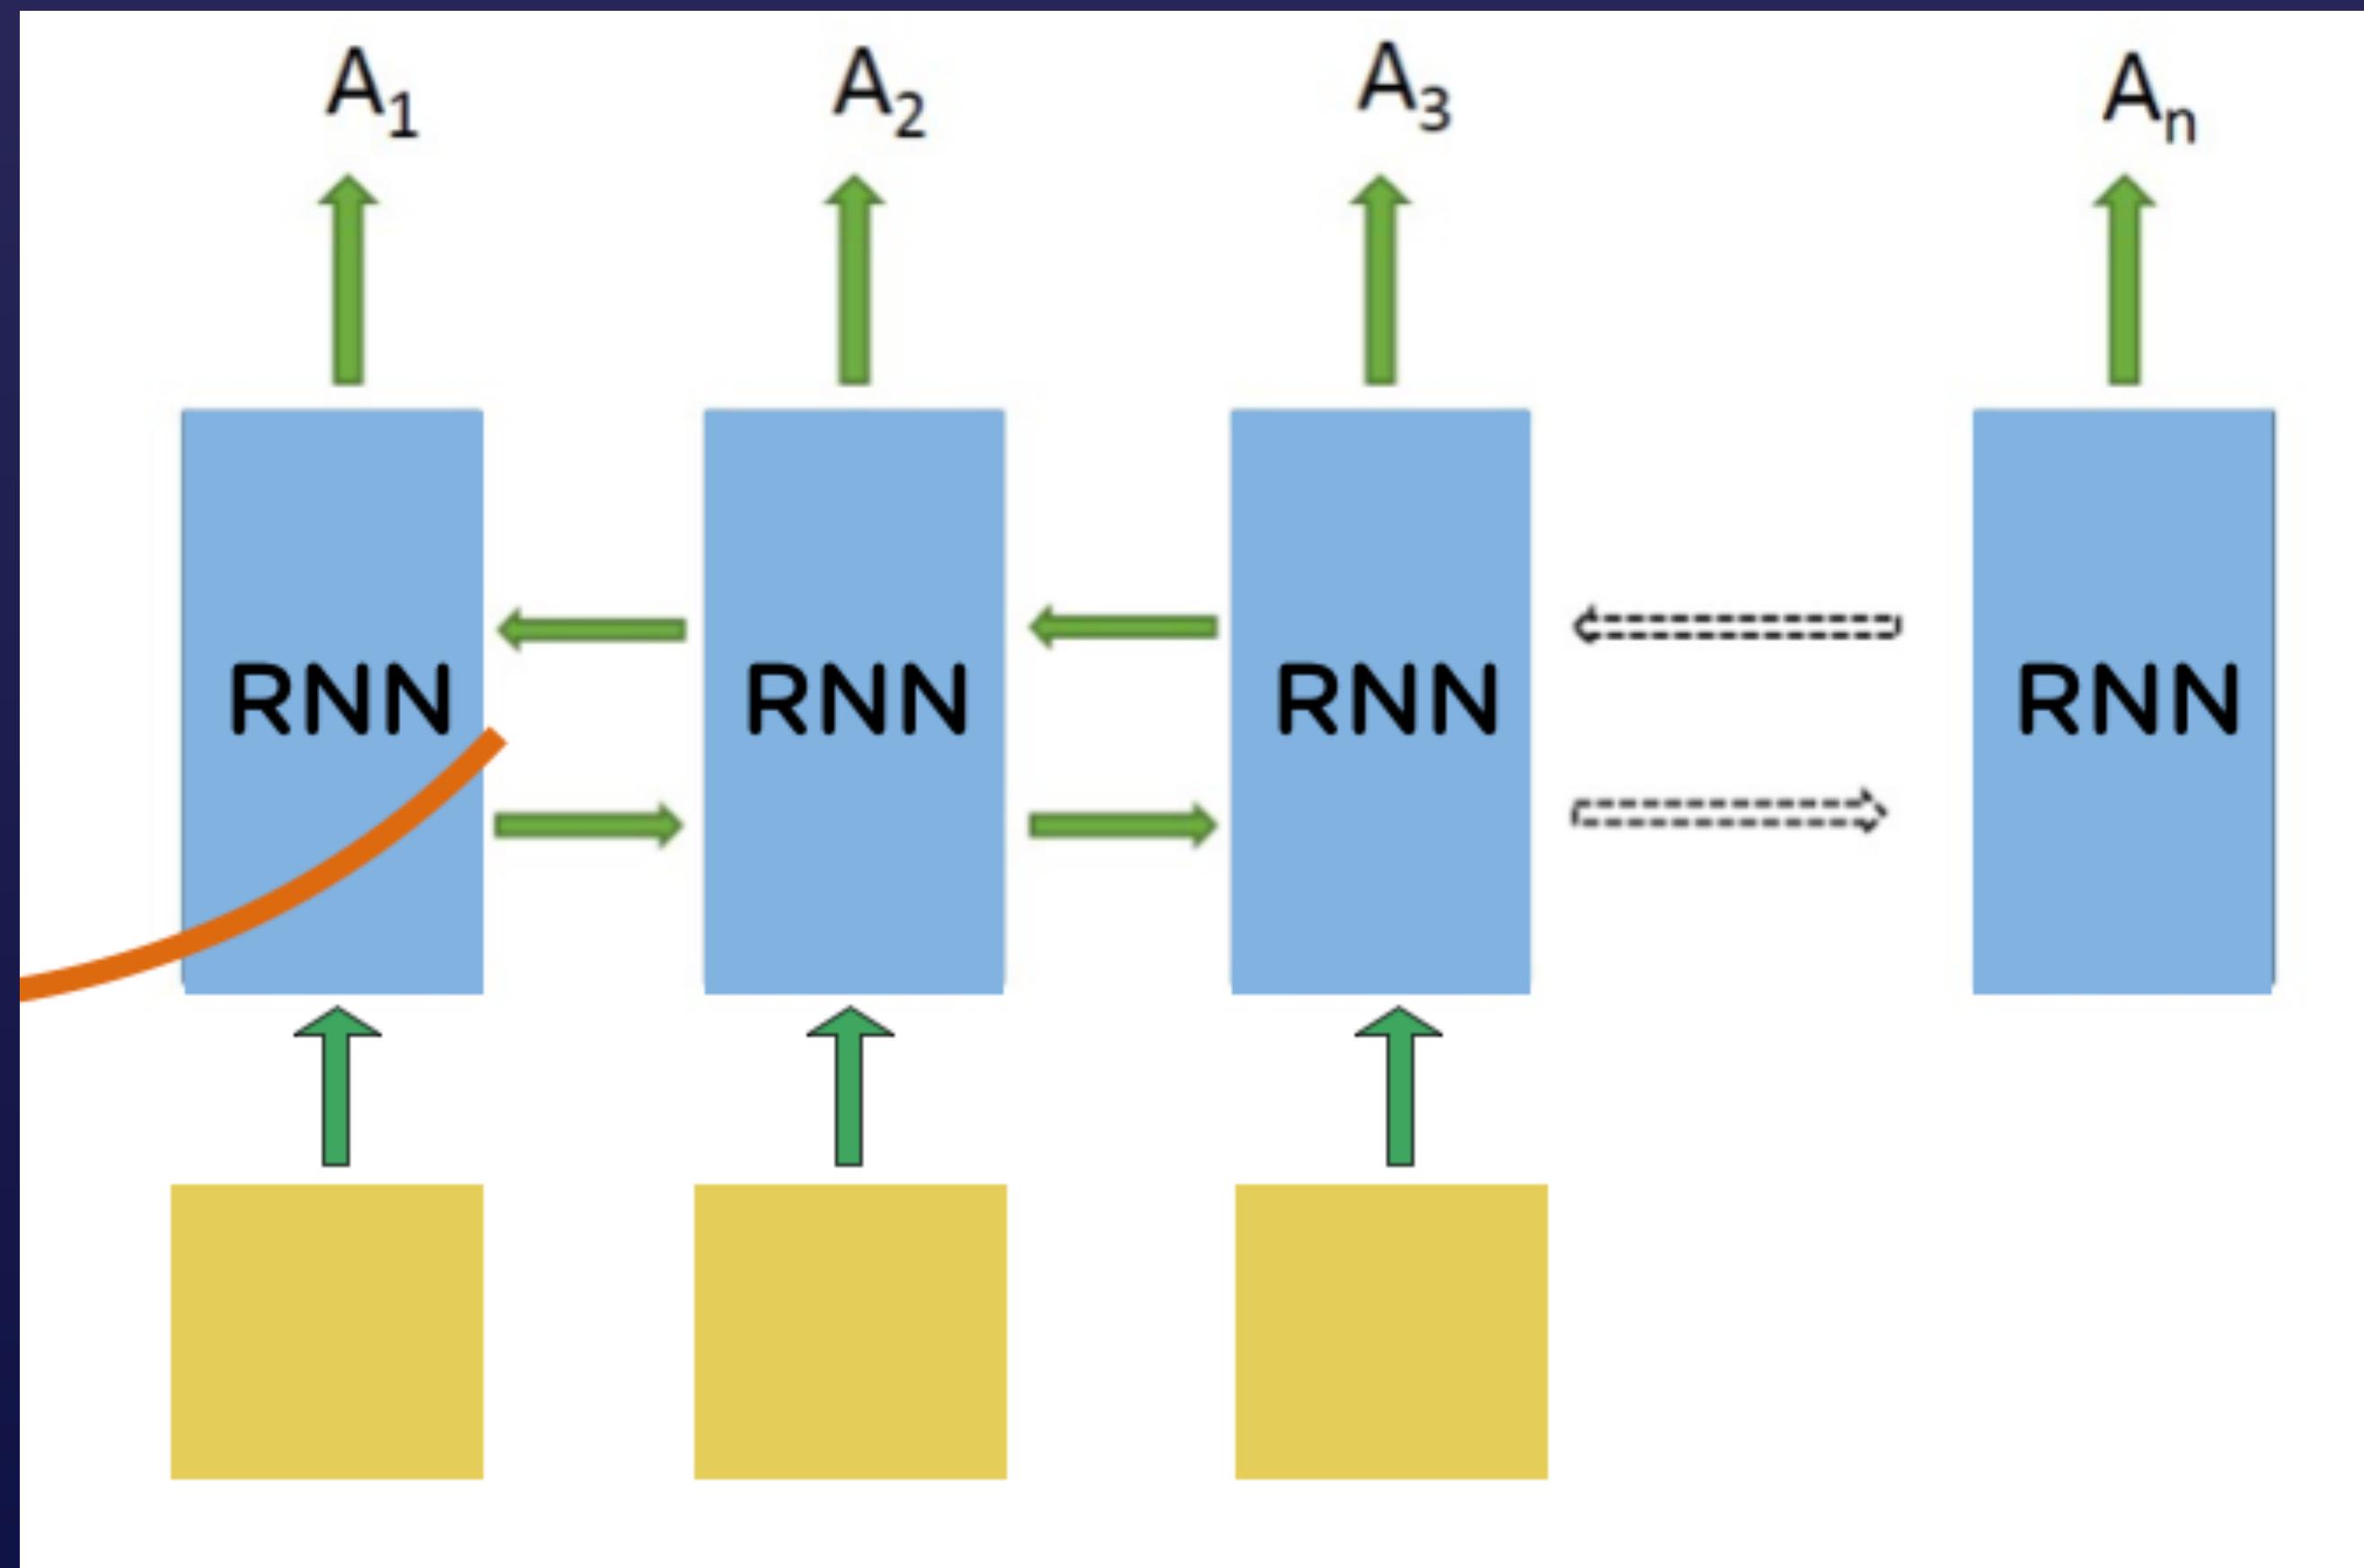
## Deep Neural Network



Reference : <http://www.cs.umd.edu/~bharat/cvpr2016.pdf>

# Sequence Labeling

## Recurrent Neural Network



Reference : <http://www.cs.umd.edu/~bharat/cvpr2016.pdf>

# Sample Code

## Initialize all variables

Attention  
DNN

RNN Cell

```
def RNN_attention(x, n_input, n_hidden, n_attn_hidden, n_classes, scope=None):
    with tf.variable_scope(scope or "RNN"):
        W_in = tf.get_variable("W_in_atten", [n_input + n_hidden, n_attn_hidden])
        b_in = tf.get_variable("b_in_atten", [n_attn_hidden])
        W2 = tf.get_variable("W2_atten", [n_attn_hidden, n_attn_hidden])
        b2 = tf.get_variable("b2_atten", [n_attn_hidden])
        W_out = tf.get_variable("W_out_atten", [n_attn_hidden, 1])
        b_out = tf.get_variable("b_out_atten", [1])
        d_0 = tf.get_variable("d_0", [1, n_input + n_hidden])

        W_h = tf.get_variable("W_mem", [n_hidden, n_hidden])
        W_i = tf.get_variable("W_in", [n_input, n_hidden])
        b_i = tf.get_variable("b_in", [n_hidden])
        W_o = tf.get_variable("W_out", [n_hidden, n_classes])
        b_o = tf.get_variable("b_out", [n_classes])
        state = tf.get_variable("state0", [1, n_hidden])
```

Output Scalar

# Sample Code

## Attention

```
# attention mechaism
def attention(query, x, scope):
    with tf.variable_scope(scope):
        tf.get_variable_scope().reuse_variables()

    d_0 = tf.get_variable("d_0")
    dnn_input = tf.scan(lambda _, b: tf.concat(1, [b, query]),
                        tf.transpose(x, [1, 0, 2]), initializer=d_0)
    dnn_input = tf.reshape(dnn_input, [-1, n_input + n_hidden])

    W_in = tf.get_variable("Win_atten")
    b_in = tf.get_variable("bin_atten")
    W2 = tf.get_variable("W2_atten")
    b2 = tf.get_variable("b2_atten")
    W_out = tf.get_variable("Wout_atten")
    b_out = tf.get_variable("bout_atten")

# attention "match" DNN
    a_attn = tf.nn.relu(tf.add(tf.matmul(dnn_input, W_in), b_in))
    a_attn2 = tf.nn.relu(tf.add(tf.matmul(a_attn, W2), b2))
    attn_score = tf.add(tf.matmul(a_attn2, W_out), b_out)
    attn = tf.nn.softmax(tf.reshape(attn_score, [-1]))

    return attn
```

DNN

Reference : [https://github.com/AaronYALai/tensorflow\\_activity\\_recognition](https://github.com/AaronYALai/tensorflow_activity_recognition)



GDG DevFest  
TAIPEI 2016



GLIA CLOUD

# Sample Code

## RNN Cell

```
# RNN cell
def cell(memory, _):
    tf.get_variable_scope().reuse_variables()
    W_h = tf.get_variable("W_mem")
    W_i = tf.get_variable("W_in")
    b_i = tf.get_variable("b_in")

    mem = tf.matmul(memory, W_h)
    attn = attention(memory, x, scope)

    x_t = tf.transpose(x[0], [1, 0])
    atten_x = tf.reshape(tf.reduce_sum(tf.mul(x_t, attn), 1), [1, n_input])

    memory = tf.nn.tanh(tf.add(tf.add(tf.matmul(atten_x, W_i), b_i), mem))
    return memory

out_seq = tf.scan(cell, tf.transpose(x, [1, 2, 0]), initializer=state)
out_seq = tf.reshape(out_seq, [-1, n_hidden])
outputs = tf.nn.softmax(tf.add(tf.matmul(out_seq, W_o), b_o))

return outputs
```

**One Step  
of RNN**

Reference : [https://github.com/AaronYALai/tensorflow\\_activity\\_recognition](https://github.com/AaronYALai/tensorflow_activity_recognition)

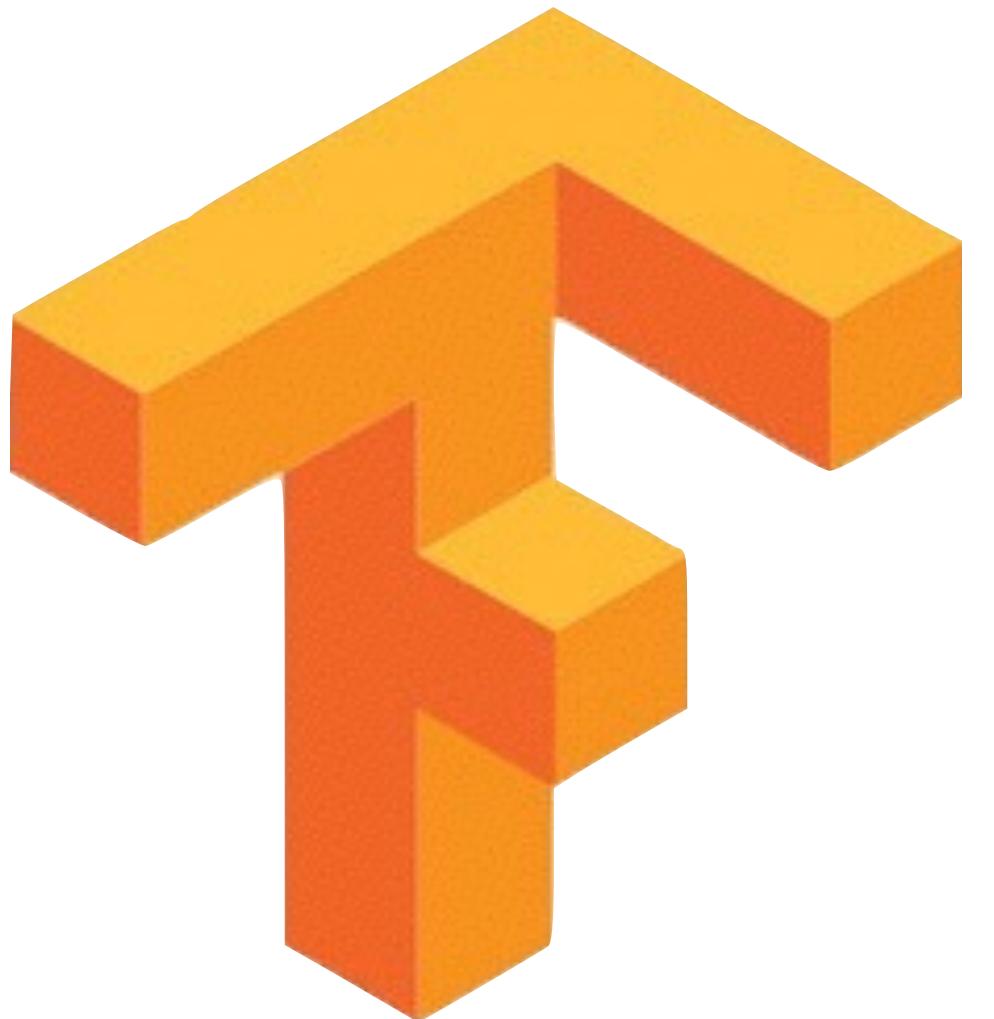


GDG DevFest  
TAIPEI 2016



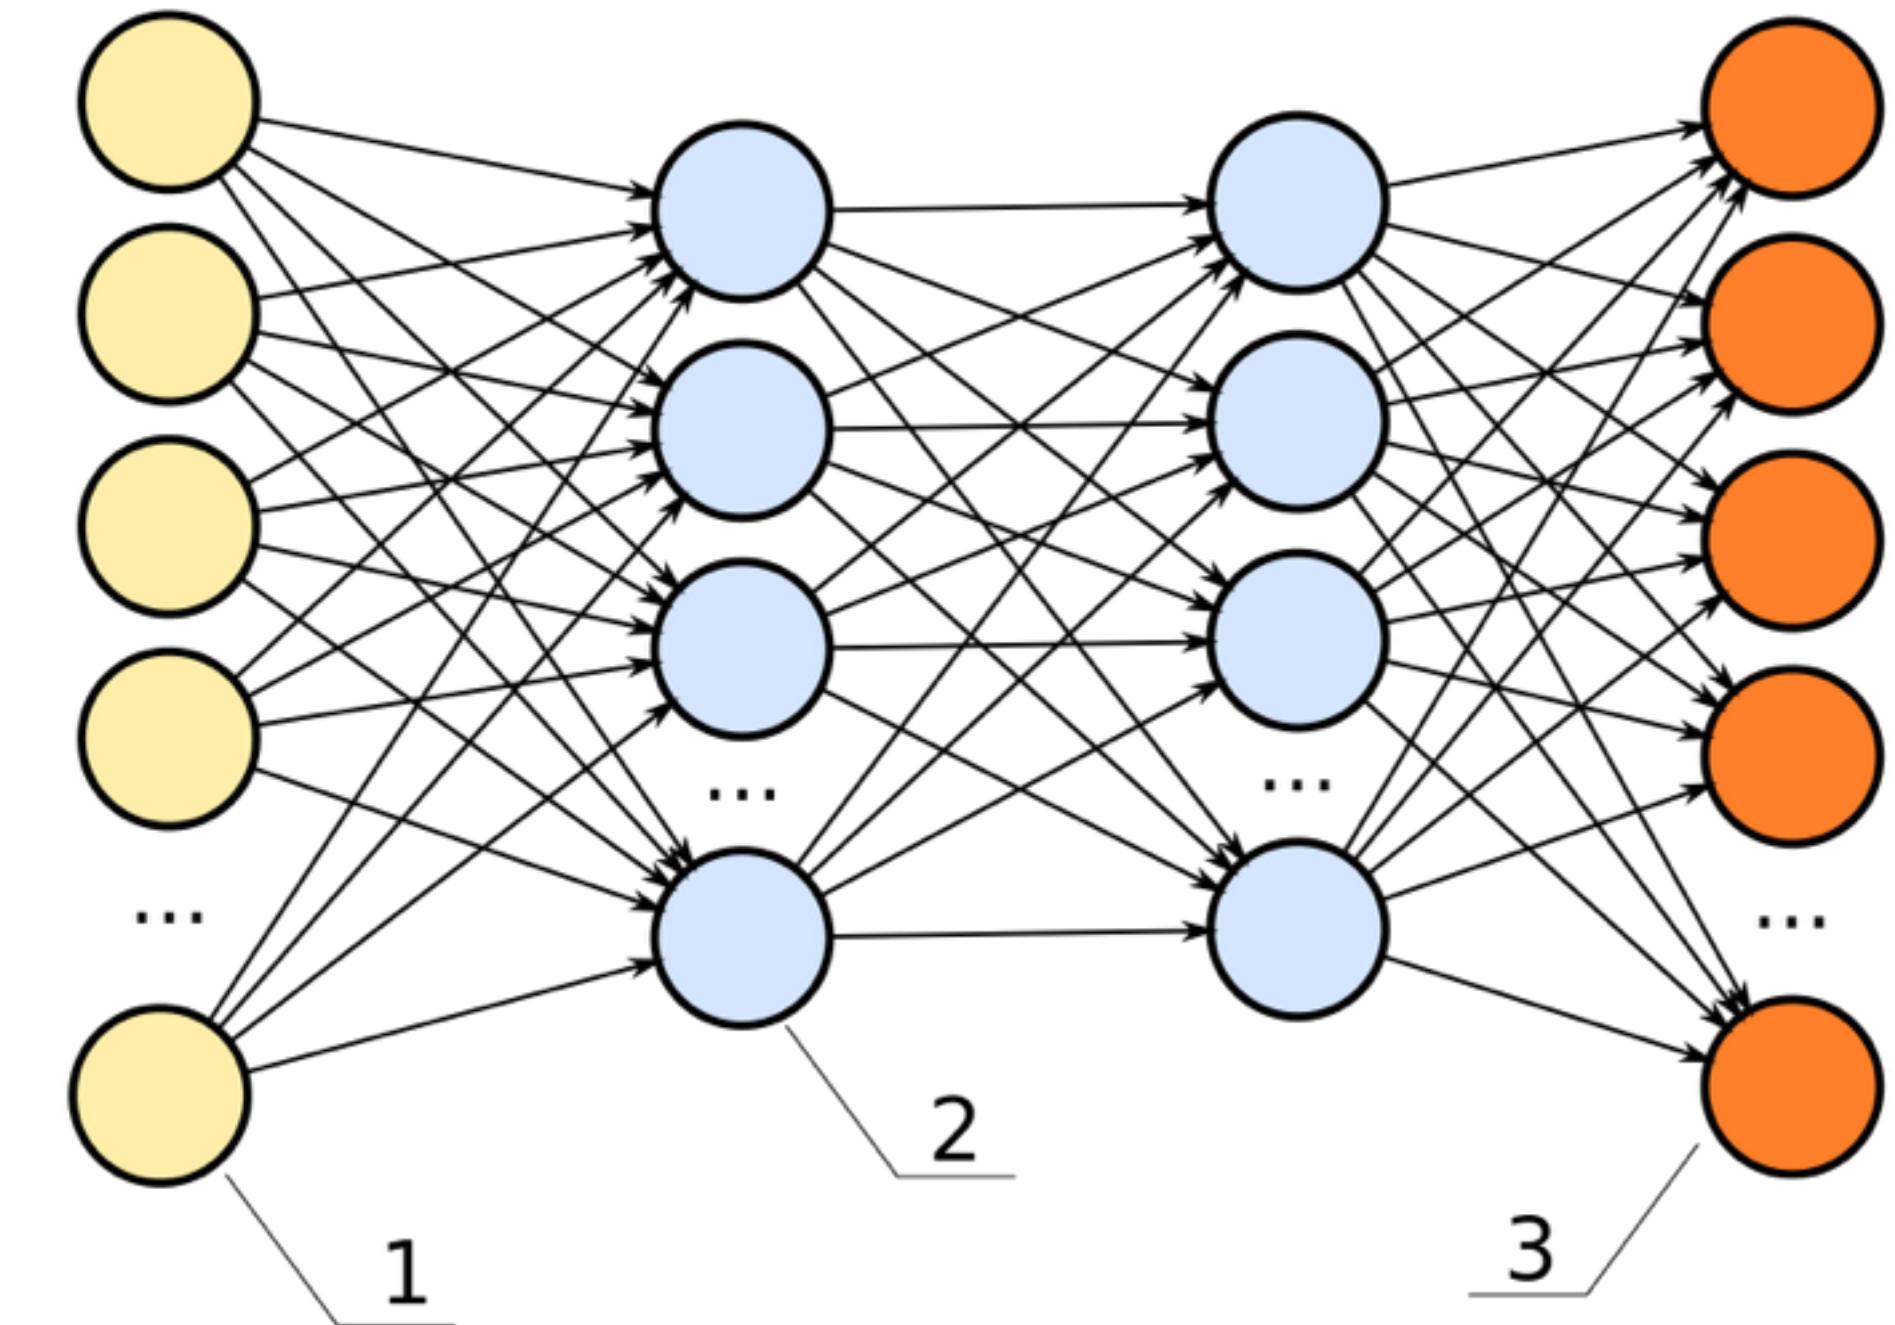
GLIA CLOUD

# Conclusions



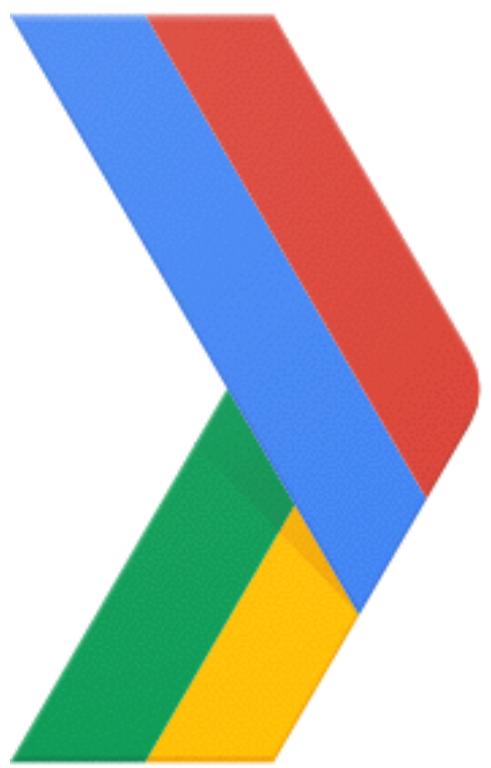
Still Young  
Many Tricks  
Be Powerful on  
Cloud Soon

More Examples  
[https://github.com/aymericdamien/  
TensorFlow-Examples](https://github.com/aymericdamien/TensorFlow-Examples)



# Resources

- Machine Learning and having it deep and structured(HungYi Lee):  
[http://speech.ee.ntu.edu.tw/~tlkagk/courses\\_MLSD15\\_2.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html)
- TensorFlow Deep Learning Examples:  
<https://github.com/aymericdamien/TensorFlow-Examples>
- Long Short-Term Memery:  
<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- “Deep Learning” by Ian Goodfellow, Yoshua Bengio and Aaron Courville:  
<http://www.deeplearningbook.org/>
- Gradient Descent Optimization Algorithms:  
<http://sebastianruder.com/optimizing-gradient-descent/index.html>



# GDG DevFest TAIPEI 2016



Thanks for Listening !



GLIA CLOUD